



# Advanced Time Series For Everyone

Bruno Gonçalves

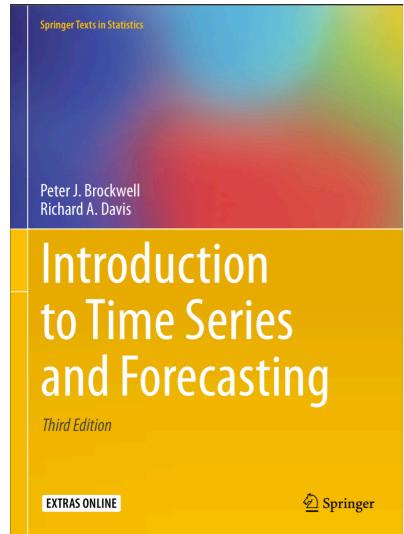
*www.data4sci.com/newsletter*

<https://github.com/DataForScience/AdvancedTimeseries>

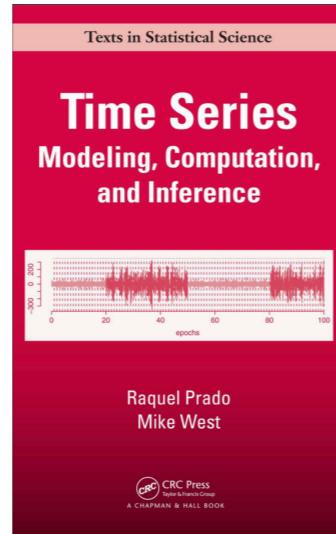


# References

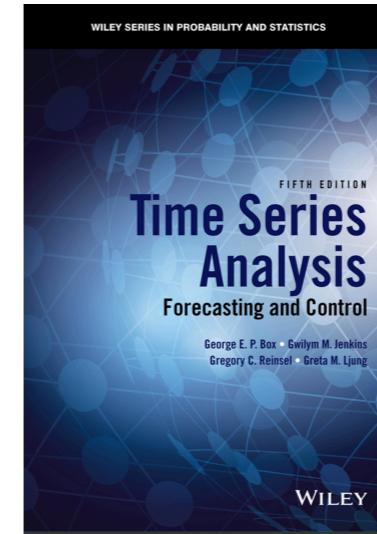
<https://github.com/DataForScience/AdvancedTimeseries>



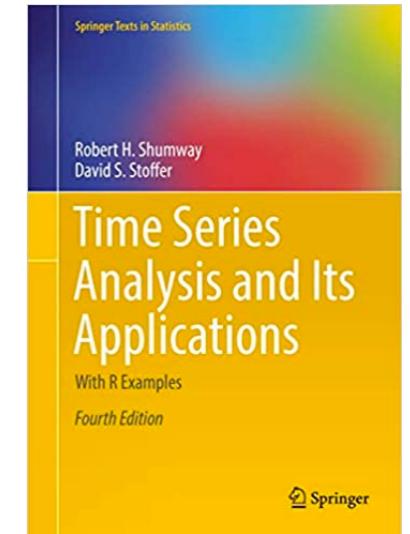
<https://amzn.to/30yEp8l>



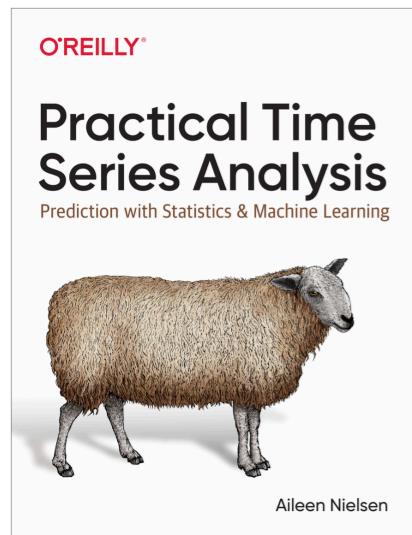
<https://amzn.to/2AhvCxt>



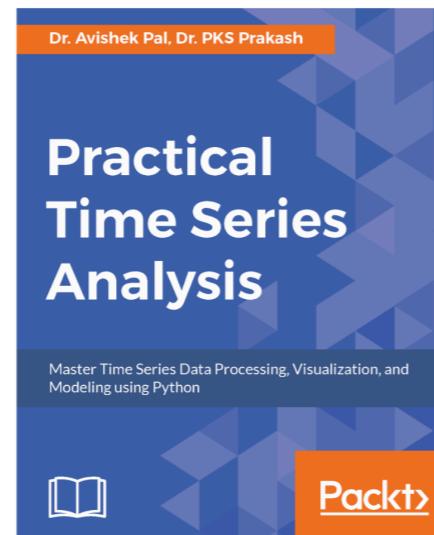
<https://amzn.to/30wY5db>



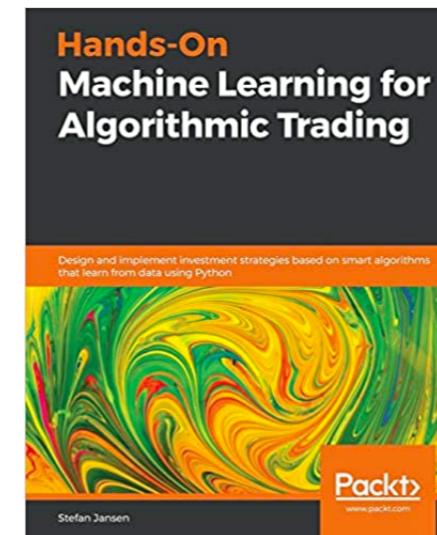
<https://amzn.to/3fOVLIM>



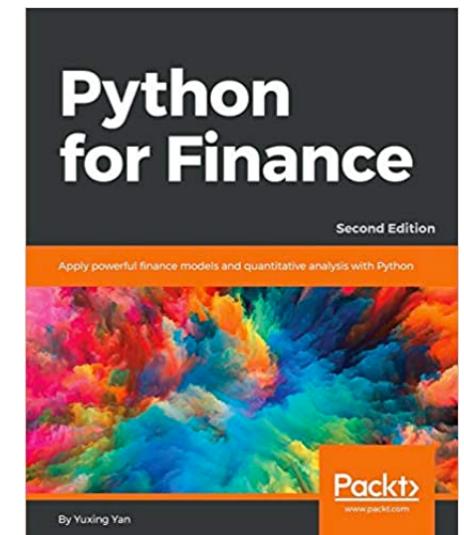
<https://amzn.to/2YoEKs6>



<https://amzn.to/30wY8Wp>



<https://amzn.to/2Ppvuj8>



<https://amzn.to/3fxzlp6>



## Lesson I:

# Pandas for Time Series

# Time Series

<https://github.com/DataForScience/AdvancedTimeseries>

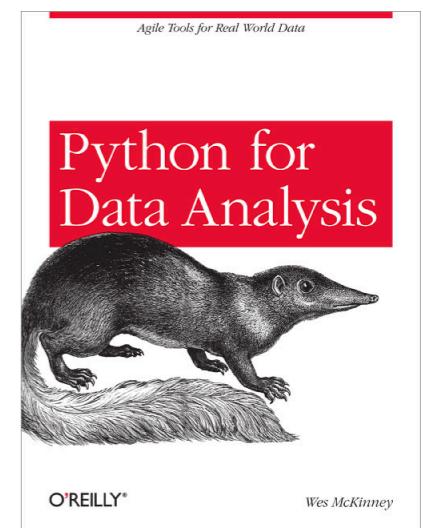
- A set of values measured **sequentially** in **time**
- Values are **typically** (but not always) measured at **equal intervals**,  $x_1, x_2, x_3, x_3$ , etc...
- Values can be:
  - **continuous**
  - **discrete** or **symbolic** (words).
- Associated with **empirical** observation of time varying phenomena:
  - Stock market prices (**day**, **hour**, minute, **tick**, etc...)
  - Temperatures (day, **minute**, second, etc...)
  - Number of patients (**week**, **month**, etc...)
  - GDP (**quarter**, **year**, etc...)
- **Forecasting** requires predicting **future** values based on **past** behavior

# pandas

[pandas.pydata.org](https://pandas.pydata.org)

- Created by [Wes McKinney](#) in 2008 while working for [AQR Capital Management](#), currently working at [Two Sigma](#) and [Ursa Labs](#)
- Pandas is specifically designed to handle time series and data frames
- High performance, flexible tool for quantitative analysis on financial data
- The functionality is built on top [numpy](#) and [matplotlib](#) (see next lecture)
- The fundamental data structures are [Series](#) (1-dimensional) and [DataFrame](#) (2-dimensional).
- Each column of a [DataFrame](#) can have a different [dtype](#) but all the elements in a column (or [Series](#)) must be of the same [dtype](#)
- Conventionally imported as

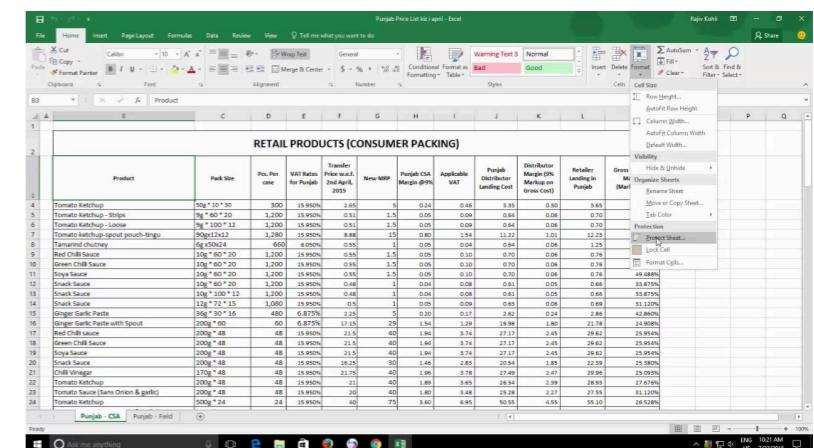
`import pandas as pd`



<https://amzn.to/2Pmg6UX>

# Series and Data Frames

- A **Series** is conceptually equivalent to a **numpy** array with some extra information (column and row names, etc)
  - A **DataFrame** can be thought of as the union of several **Series**, with names associated.
  - Similarly, you can think of a **DataFrame** as an Excel **sheet** and of a **Series** as an individual **column**
  - A minimal **DataFrame** implementation would be a dict where each element is a list



# Series

	id
0	23
1	42
2	12
3	86

+

## Series

	Name
0	“Bob”
1	“Karen”
2	“Kate”
3	“Bill”

2

# DataFrame

	<b>id</b>	<b>Name</b>
0	23	“Bob”
1	42	“Karen”
2	12	“Kate”
3	86	“Bill”

# DataFrame Fields

---

- **DataFrame**s contain a great deal of extra information in addition to just the data values. In particular, they include:
  - **columns** - column names
  - **index** - row names
  - **dtypes** - dtype associated with each column
  - **shape** - number of rows and columns
  - **ndim** - the number of dimensions
- **DataFrame**s make it easy to add new columns.
- **DataFrame** elements can be accessed and modified in various ways

# DataFrame Exploration

---

- We get information about the contents of a **DataFrame** using several simple methods:
  - **head()** - display the top 5 rows
  - **tail()** - display the bottom 5 rows
  - **info()** - Print a concise summary of a DataFrame
  - **describe()** - Generate descriptive statistics that summarize the data

# Indexing and Slicing

---

- pandas supports various ways of indexing the contents of a **DataFrame**
- [**<column name>**] - select a given column by it's name. column names can also be used as field names
- **.loc[<row name>]** - select a given row by it's (Index) name
- **.iloc[<position>]** - select a given row by it's position in the Index, starting at 0
- **.loc[<row name>, <column name>]** - select an individual element by row and column name
- **.iloc[<row position>, <row position>]** - select an individual element by row and column position (starting at 0)
- **.iloc** also supports ranges and slices similarly to **python** lists or **numpy** arrays

# Importing and Exporting Data

---

- pandas has powerful methods to read and write data from multiple sources
  - `pd.read_csv()` - Read a comma-separated values file
    - `sep=' '` - Define the separator to use. ',' is the default
    - `header=0` - Row number to use as column names
  - `pd.read_excel()` - Read an Excel file
    - `sheet_name` - The sheet name to load
  - `pd.read_html()` - Read tabular data from a URL (or local html file)
  - `pd.read_pickle()` - Read a Pickle file
- Each of these functions accepts a large number of options and parameters controlling its behavior. Use `help(<function name>)` to explore further.
- Each `read_*` function has a complementary `to_*` function to write out a **DataFrame** to disk. The `to_*` functions are members of the **DataFrame** object

# Time Series

---

- pandas was originally developed to handle financial data
- Temporal sequences (time series) are a common type of data encountered in financial applications, so, naturally, pandas has good support for the most common time series operations.
- **pd.to\_datetime** - converts a Series or value into a date time timestamp.
  - Format is inferred by looking at the entire Series data
  - Format can also be manually specified using specific arguments:
    - **format** - detailed string specifying the full format
    - **dayfirst=True** - parse 10/11/12 as Nov 10, 2012
    - **yearfirst=True** - parse 10/11/12 as Nov 12, 2010
- **pd.to\_timedelta** - converts a Series into an absolute difference in time
- Dates can also be parse automatically at read time by passing a list of the columns containing dates to the **parse\_dates** argument of **pd.read\_csv**

# Time Series

---

- Sequences of dates/times can be created using `pd.date_range()`, similar to `np.arange()`
  - **start/end** - specify the limits
  - **freq** - specify frequency (step)
    - B - business day frequency
    - D - calendar day frequency
    - W - weekly frequency
    - M - month end frequency
    - Q - quarter end frequency
    - A - year end frequency
    - H - hourly frequency
    - T - minutely frequency
    - S - secondly frequency
- By setting the index to be a date time, we turn the **DataFrame** into a **Time Series**.
- Pandas has several methods that are specialized to this case
- `index.day/index.month/index.year` - return the day, month and year of the time series index value

# Transformed values

---

- We often need to apply some simple transformation to the values in our original **Series**.  
Pandas offers several methods to accomplish this goal:
  - **map(func)** - Map values of **Series** according to input correspondence.
    - **func** - **function**, **dict** or **Series** to use for mapping
  - **transform(func)** - Transform each column/row of the **DataFrame** using a function.  
Output must have the same shape as the original
    - **axis = 0** - apply function to columns
    - **axis = 1** - apply function to rows
  - **apply(func)** - Apply a function along an axis of the **DataFrame**
    - Similar to **transform()** but without the limitation of preserving the shape

# Lagged values and differences

---

- While analyzing time series, we often refer to values that our time series took **1, 2, 3**, etc time steps in the past
- These are known as lagged values and denoted:

$$x_{t-l}$$

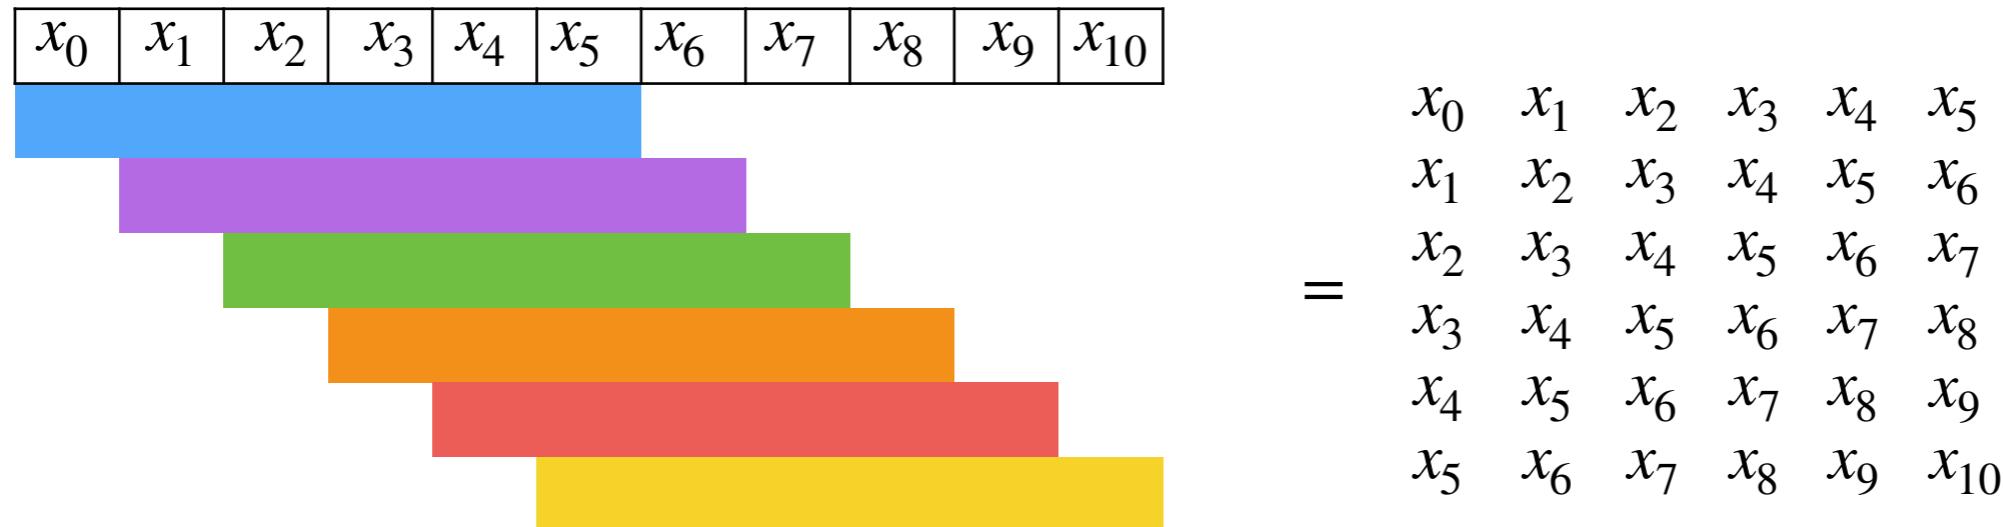
- where  **$l$**  is the value of the lag we are considering.
- pandas supports lagged values with the **shift( $l$ )** DataFrame method
- Perhaps the most common use case for lagged values is for the calculation of **differences** of the form:

$$x_t - x_{t-l}$$

- Where  **$l \geq 1$**  is the value of the lag we are interested in.
- Differences are also a particularly simple way to **detrend** a time series
- pandas supports lagged values with the **diff( $l$ )** DataFrame method

# Windowing

- When analyzing the temporal behavior of a signal, we often need to evaluate if specific quantities are **time varying** or not
- A common approach is to use **sliding windows** of a given length to evaluate the required values
- So a sliding window of width **6** on a series of length **11** would look like:



- and we would calculate the metric of interest **within each window**.
- In Pandas we use the method **rolling(w)** that returns dynamic view of the data allowing us to daisy-chain other common **numpy** operators. **rolling(w).mean()**, **rolling(w).max()**, etc...

# groupby

---

- Sometimes we need to calculate statistics for subsets of our data
- `groupby()` allows us to group data based on the values of a column
- `groupby()` returns a GroupBy object that supports several aggregations functions, including:
  - `max()/min()/mean()/median()`
  - `transform()/apply()`
  - `sum()/cumsum()`
  - `prod()/cumprod()`
  - `quantile()`
- Each of these functions is applied to the contents of each group

# Pivot Tables

---

- Pandas provides an extremely flexible pivot table implementation
- `pd.pivot_table()` - Creates a spreadsheet-style pivot table as a `DataFrame`.
- Three important arguments:
  - **values** - column to aggregate
  - **index** - Keys to group by on the pivot table index. If an array is passed, it is being used as the same manner as column values.
  - **columns** - Keys to group by on the pivot table column.
  - **aggfunc** - Function to use for aggregation. Defaults to `np.mean()`
- **index**, **columns** and **aggfunc** can also be lists of keys or functions to use.

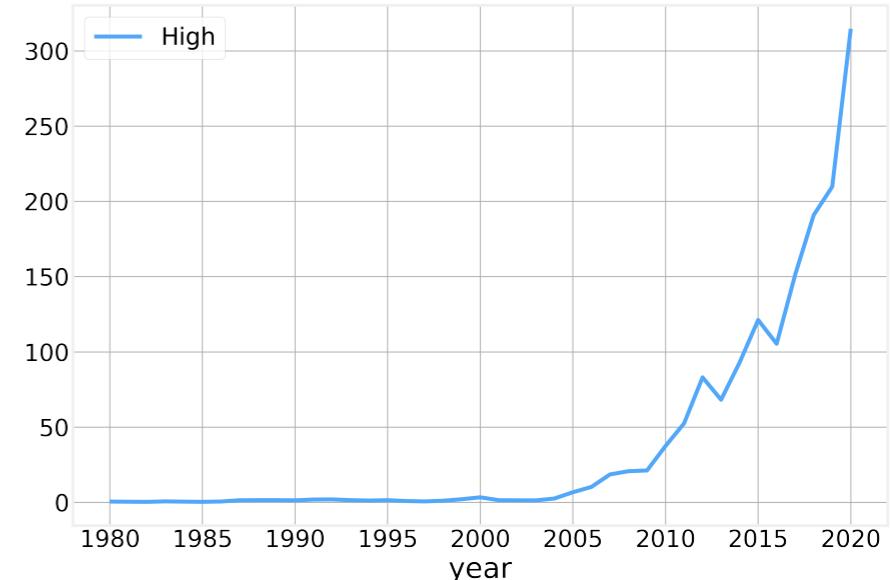
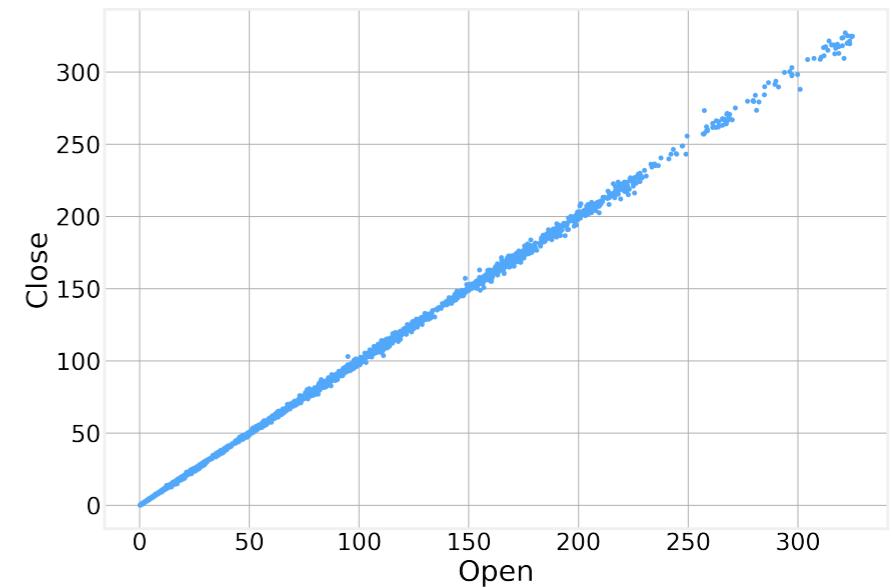
# merge/join

---

- **merge()** and **join()** allows us to perform database-style join operation by columns or indexes (rows)
- Some of the most important arguments are common to both methods:
  - **on** - Column(s) to use for joining, otherwise join on index. Can pass an array as the join key if not already contained in the calling DataFrame. Like an Excel VLOOKUP operation
  - **how** - Type of join to perform: **{'left', 'right', 'outer', 'inner'}**
- **merge()** is more sophisticated and flexible. It also allows us to specify:
  - **left\_on/right\_on** - Field names to join on for each DataFrame
  - **left\_index/right\_index** - Whether or not to use the left/right index as join key(s)

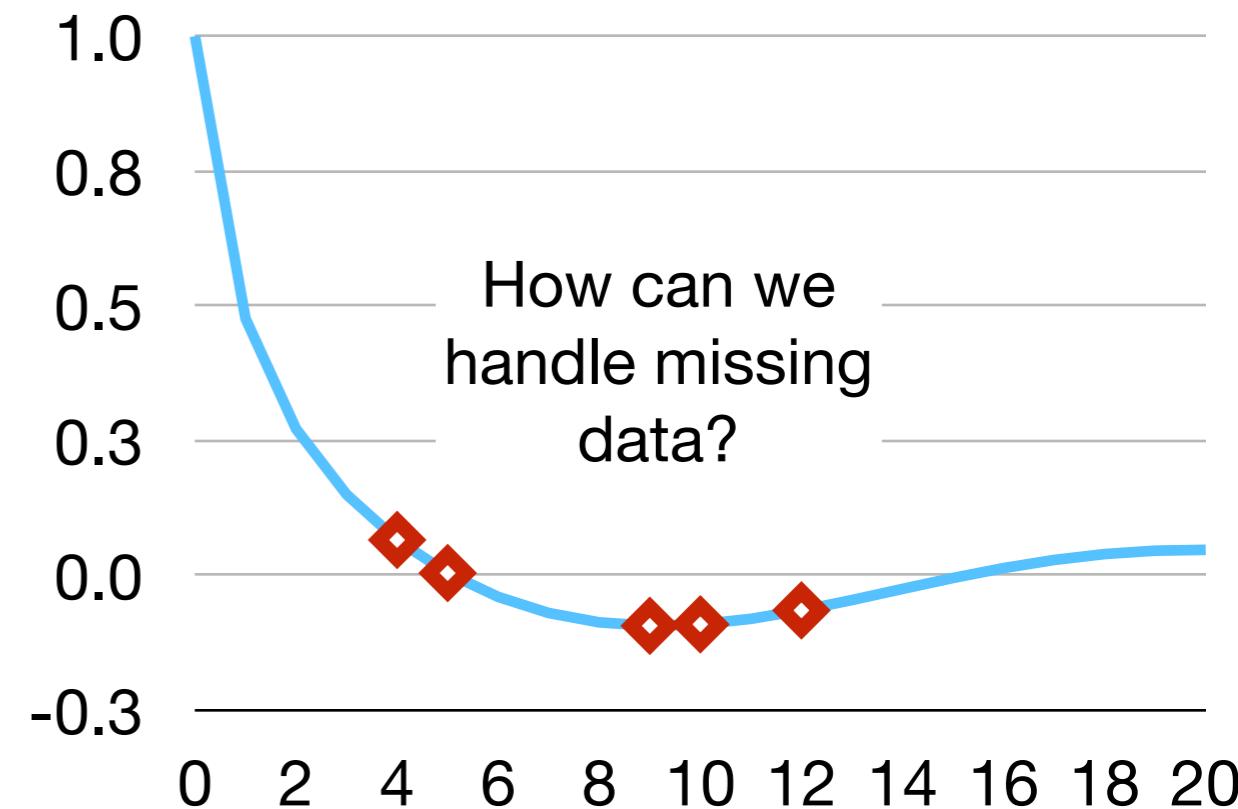
# plotting

- Finally, pandas also provides a simple interface for basic plotting through the `plot()` function
- By specifying some basic parameters the variables plotted and even the kind of plot can be easily modified:
  - **x/y** - column name to use for the **x/y** axis
  - **kind** - type of plot
    - ‘**line**’ - line plot (default)
    - ‘**bar**’/‘**barh**’ - vertical/horizontal bar plot
    - ‘**hist**’ - histogram
    - ‘**box**’ - boxplot
    - ‘**pie**’ - pie plot
    - ‘**scatter**’ - scatter plot
- The `plot()` function returns a `matplotlib Axes` object

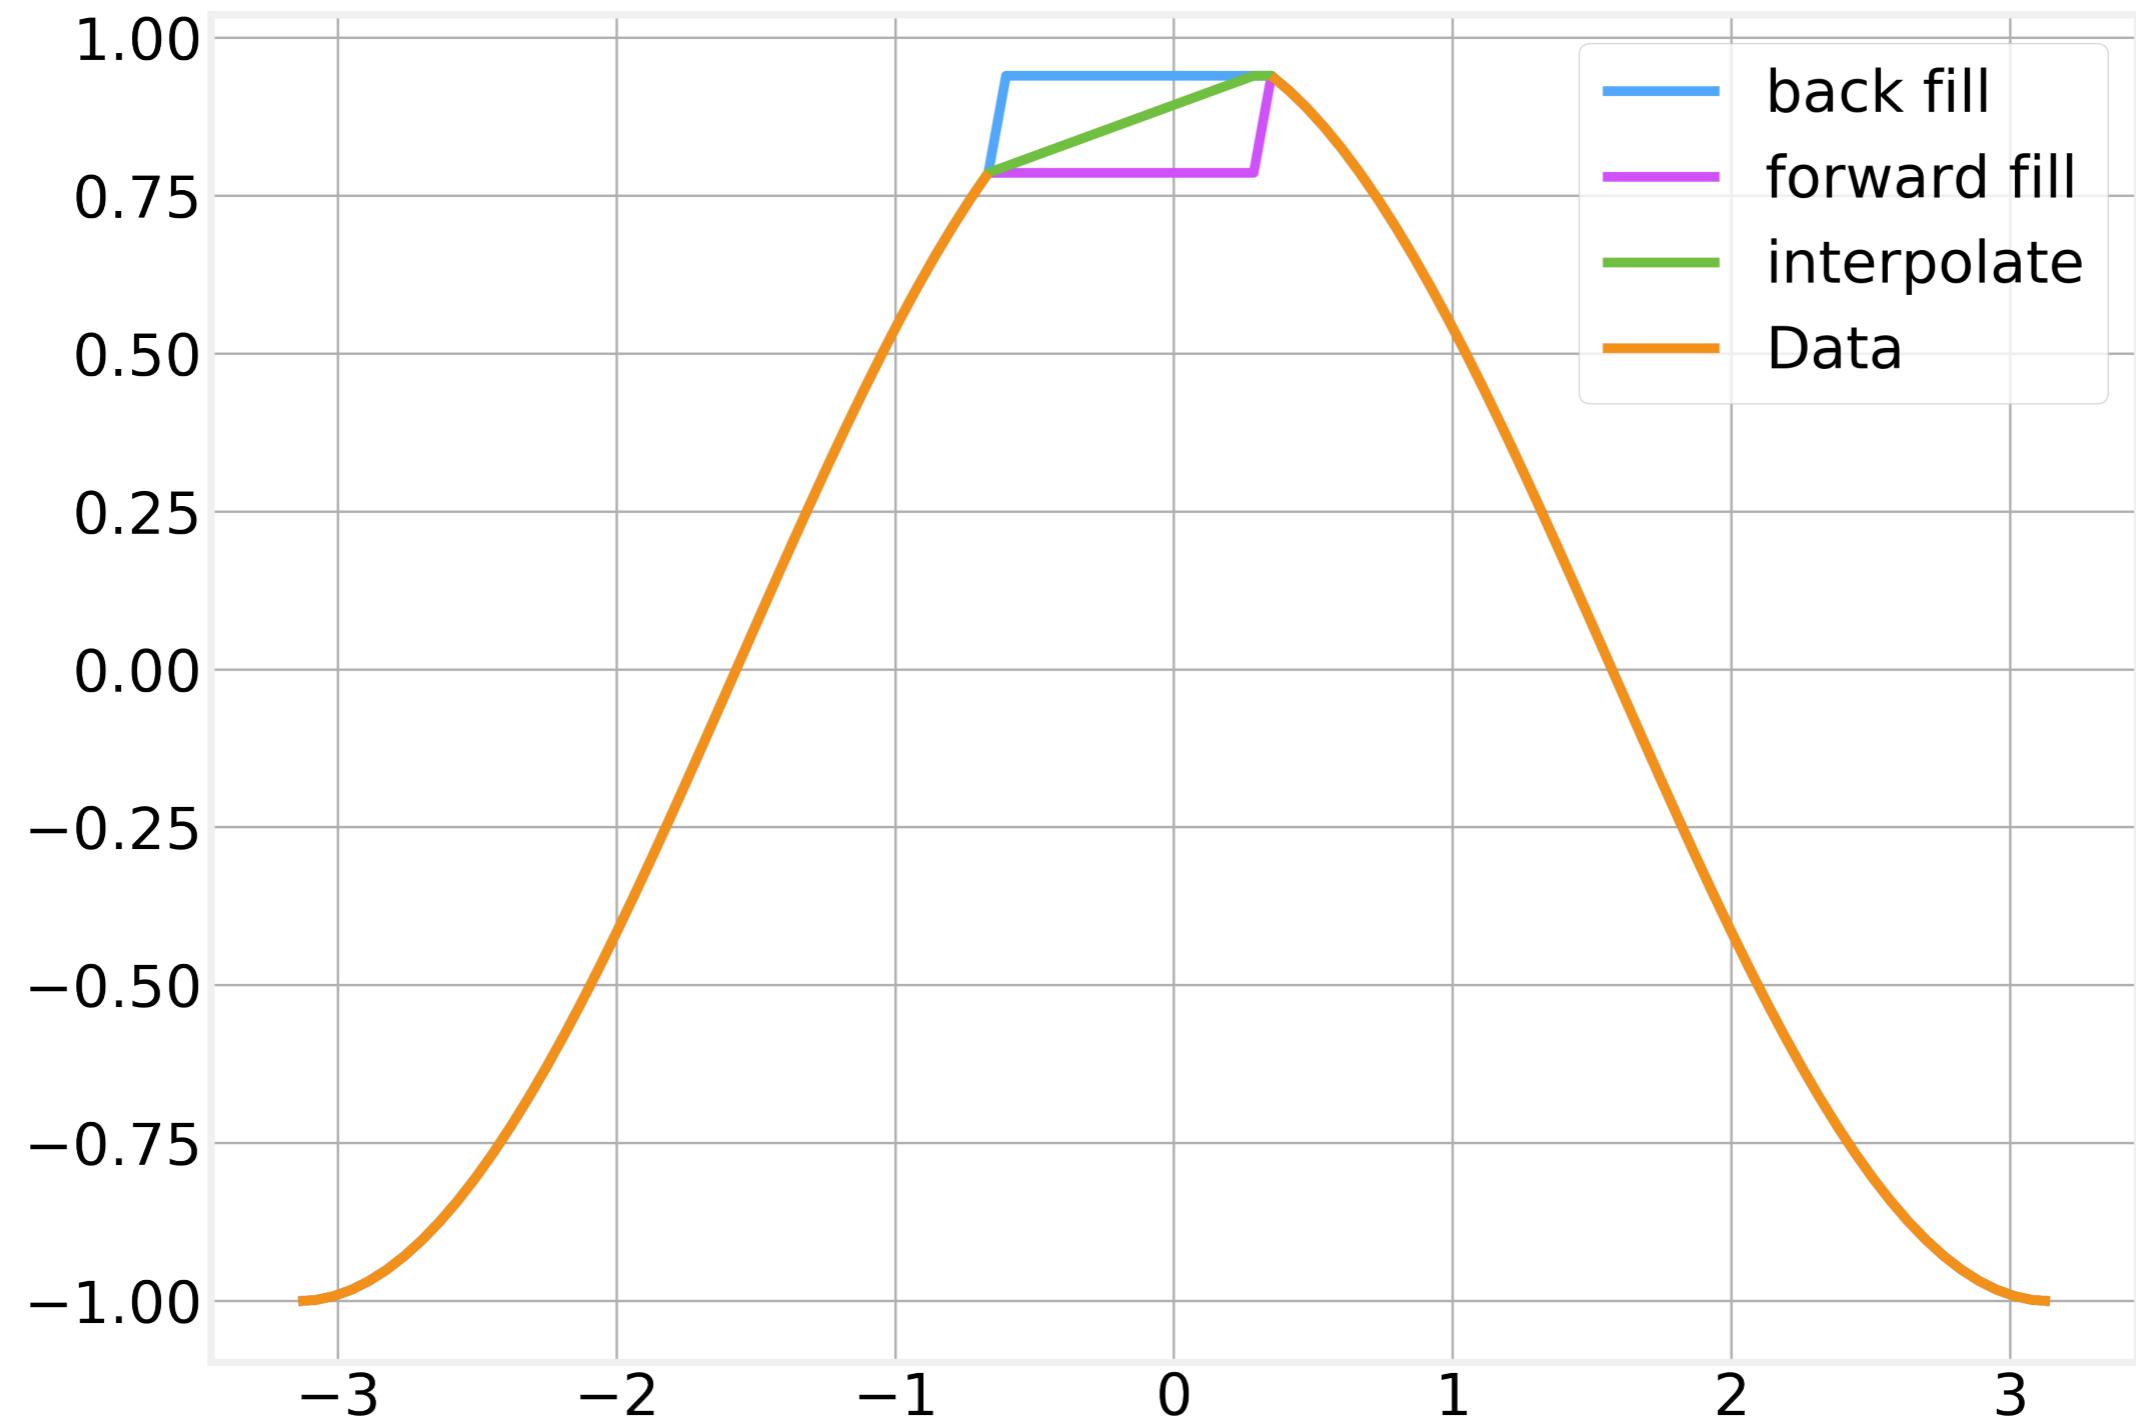


# Fill methods

- Often the time series is [incomplete](#).
- Missing data points can be due to data corruption, data collection issues, etc.
- Missing values are represented as [nan](#)
- Several techniques have been developed to handle this case:
  - `fillna()` - replace all [nan](#) values with a specific value
  - `ffill()` - keep the last valid value
  - `bfill()` - keep the next valid value
  - `interpolate()` - add values by interpolating between the previous and the next value:
    - the default method is "[linear](#)" interpolation, but pandas supports many others, including "[time](#)" to scale the values based on time interval.



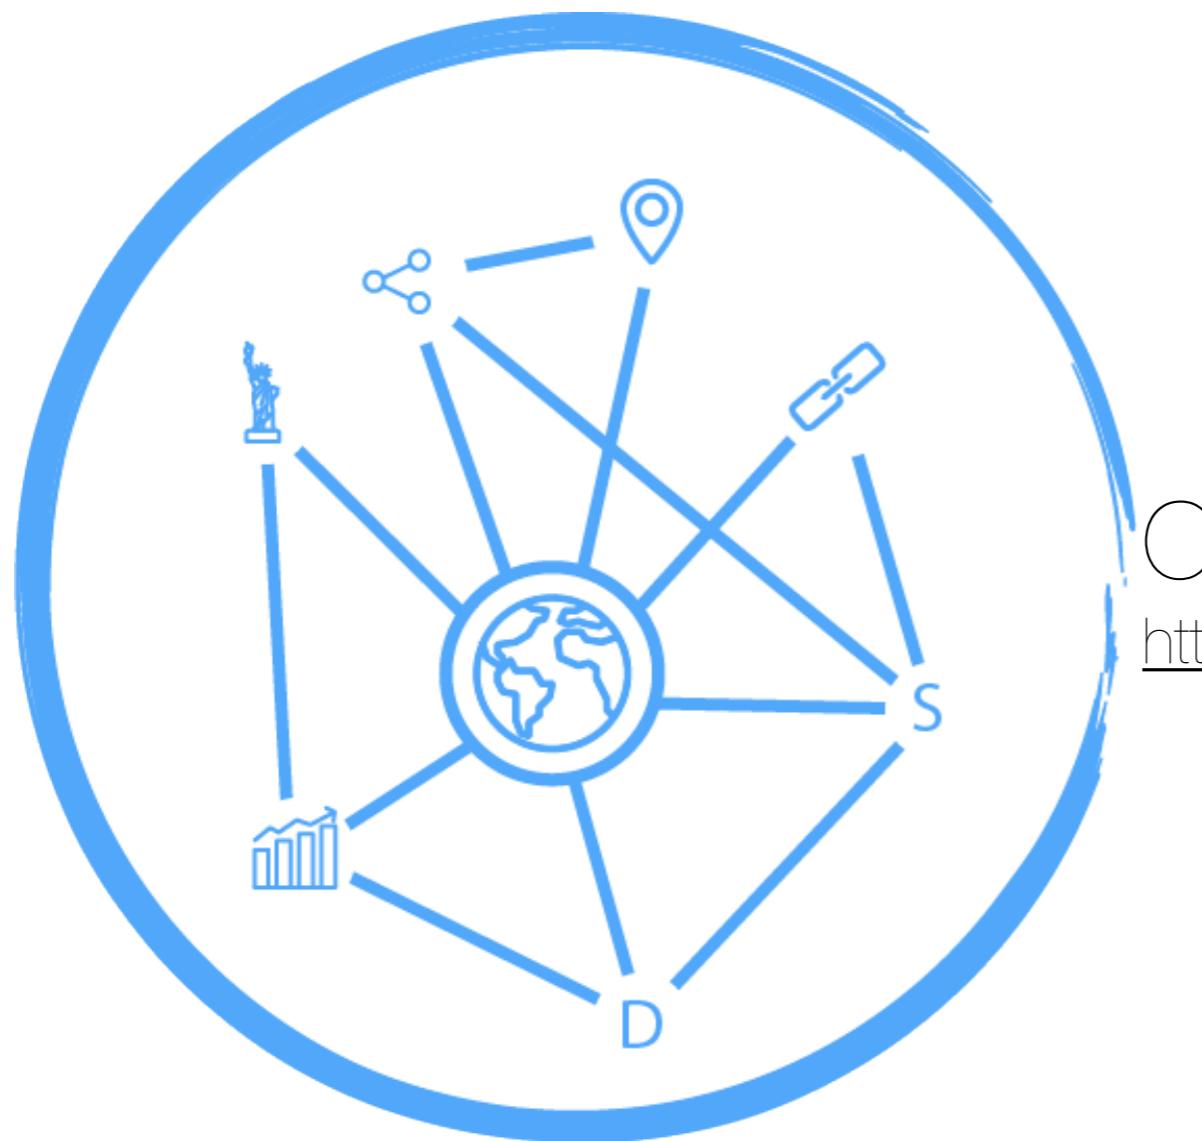
# Fill methods



# Resampling

---

- Time Series typically have an **intrinsic time scale** at which the data was collected: ticks, seconds, days, months, etc
- In many cases, our analysis requires that we **resample** the data to a different time scale
- Resampling to a longer timescale is relatively simple and similar to aggregation:
  - Transforming from daily to weekly frequency requires simply aggregating by week
- Resampling to shorter timescales requires **interpolation or imputation** to make up for the missing values
  - Going from weekly to daily frequency requires **specifying how to allocate** the values for each day of the week



Code - Pandas for Time Series  
<https://github.com/DataForScience/AdvancedTimeSeries>



## Lesson II: statsmodels for Time Series



[statsmodels.org](http://statsmodels.org)

- Python module built on top of **numpy** and **scipy**
- Provides classes and functions for:
  - estimations of various statistical models
  - statistical tests
  - data exploration
- Extensive statistics are available for each estimate produced
- Particularly useful for:
  - Linear Regression
  - Generalized Linear Models
  - Survival Analysis
  - **Time Series!**



statsmodels.org

- Conventionally imported as:

```
import statsmodels.api as sm
```

- Supports **R-style** formula notation for statistical models
- For **R-style** formula notation support, conventionally imported as:

```
import statsmodels.formula.api as smf
```

- And for timeseries:

```
import statsmodels.api as sm
```

-



statsmodels.org

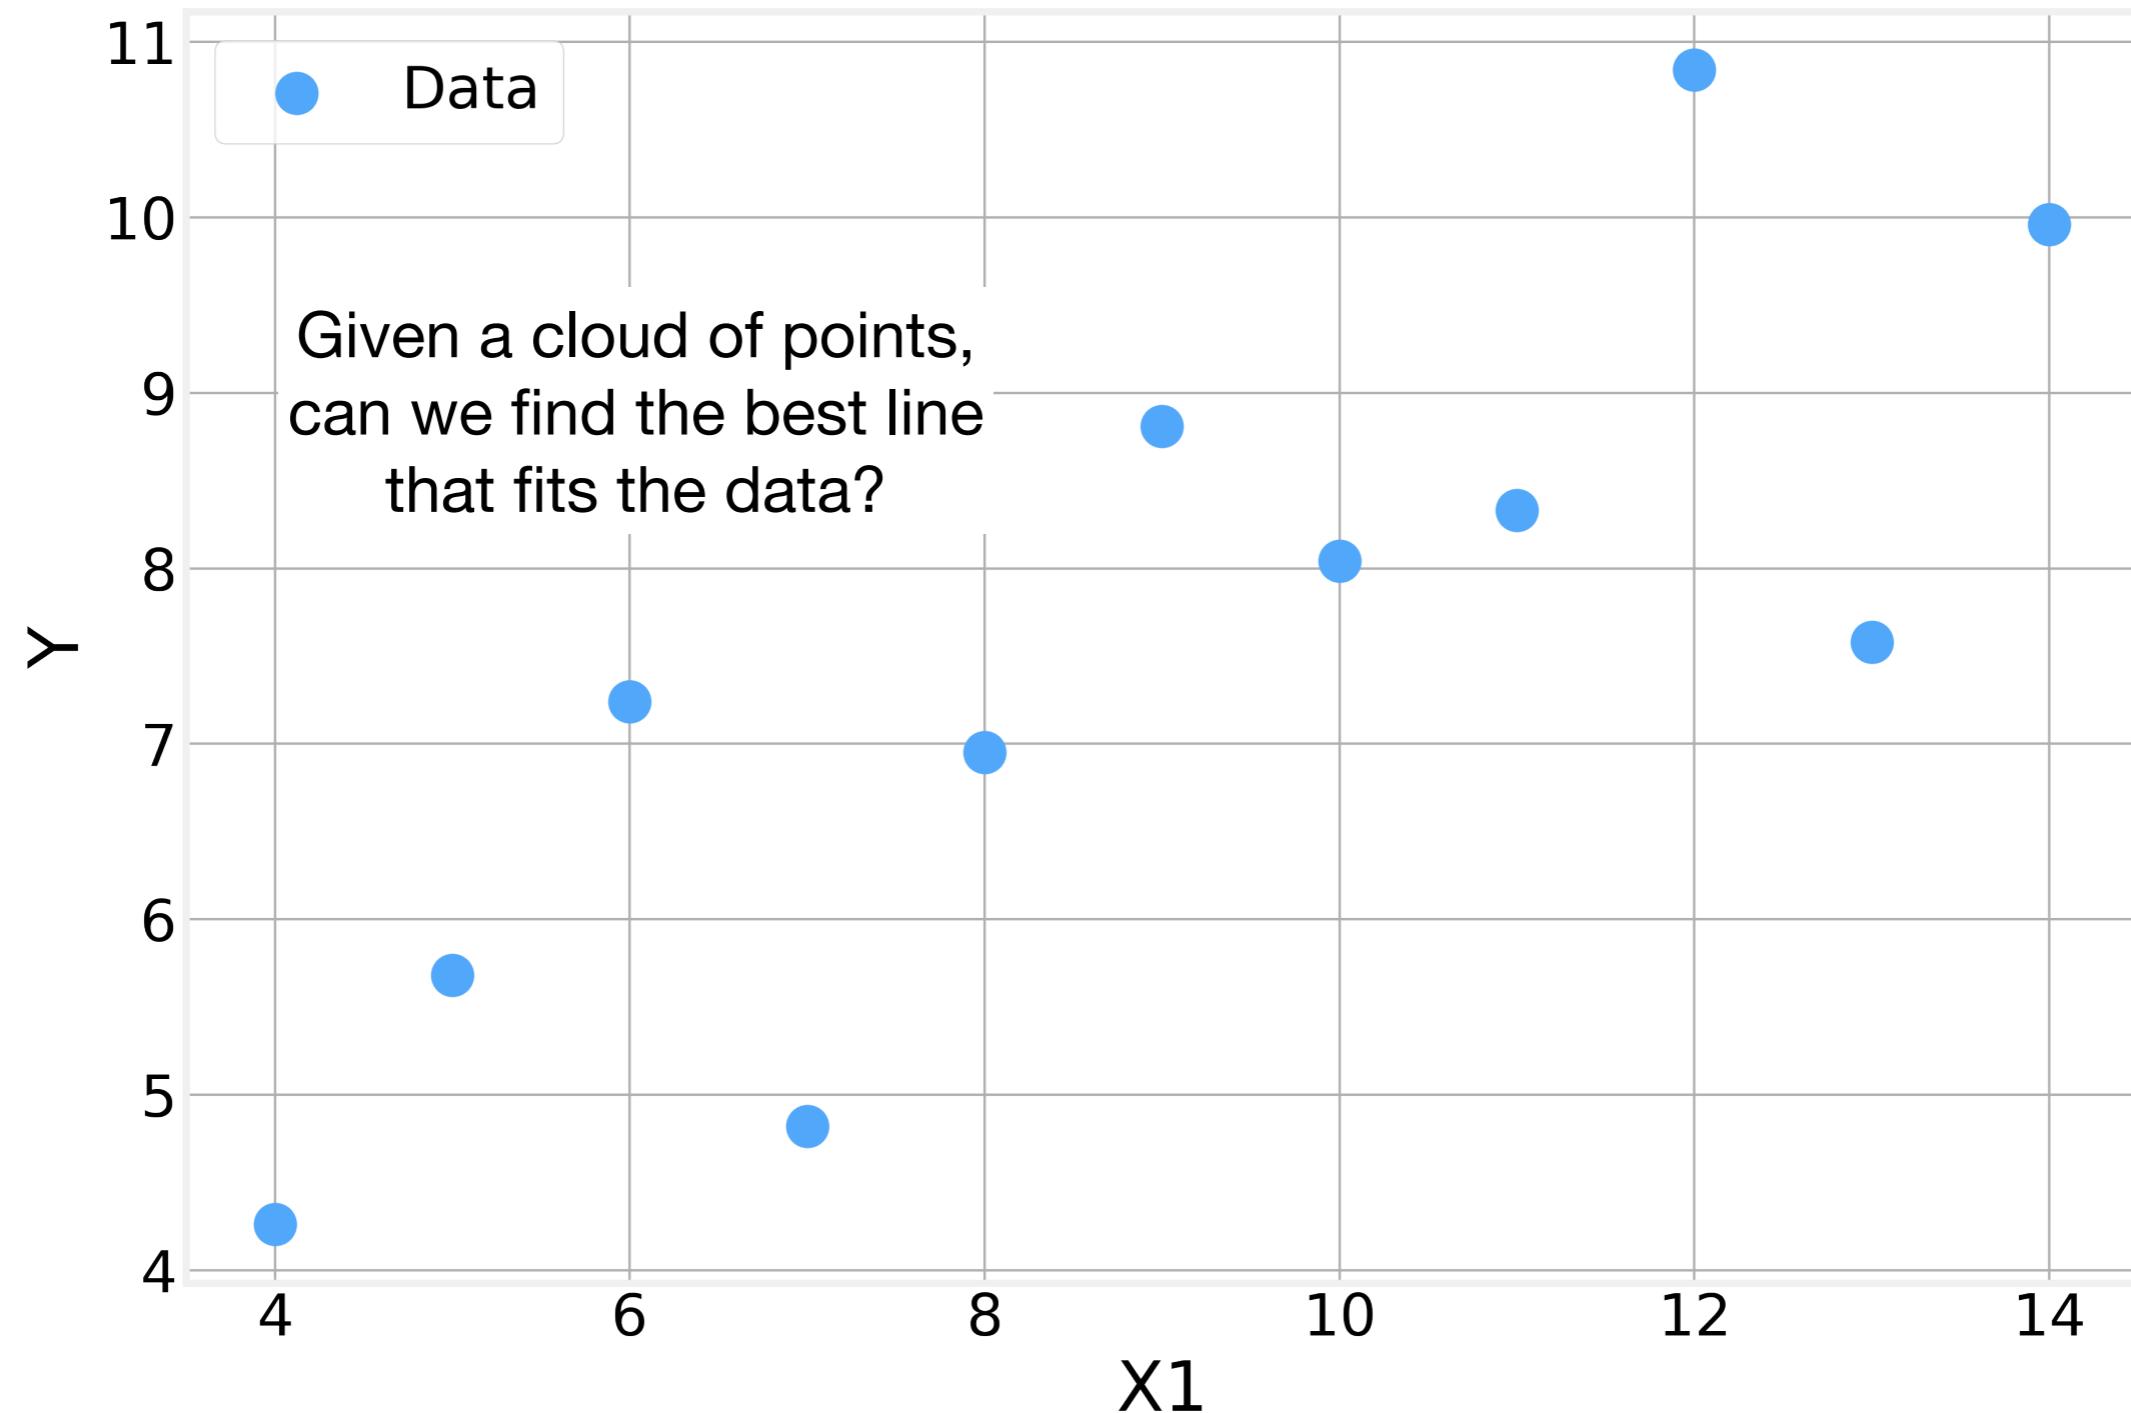
statsmodels.	api.	OLS()	.fit() .predict() .summary()	
		tsa.	seasonal_decompose()	
		stattools.		acf() pacf() adfuller()
		arima.		ARIMA() .fit() .predict() .forecast() .plot_predict() .summary()
		statespace.		SARIMAX() .fit() .predict() .forecast() .summary()
	graphics.	tsa.		plot_acf() plot_pacf()

- statsmodels refers to variables as:
  - “**endog**enous” - caused by factors within the system
  - “**exog**enous” - caused by factors outside the system

endog	exog
y	x
y variable	x variable
left hand side (LHS)	right hand side (RHS)
dependent variable	independent variable
regressand	regressors
outcome	design
response variable	explanatory variable

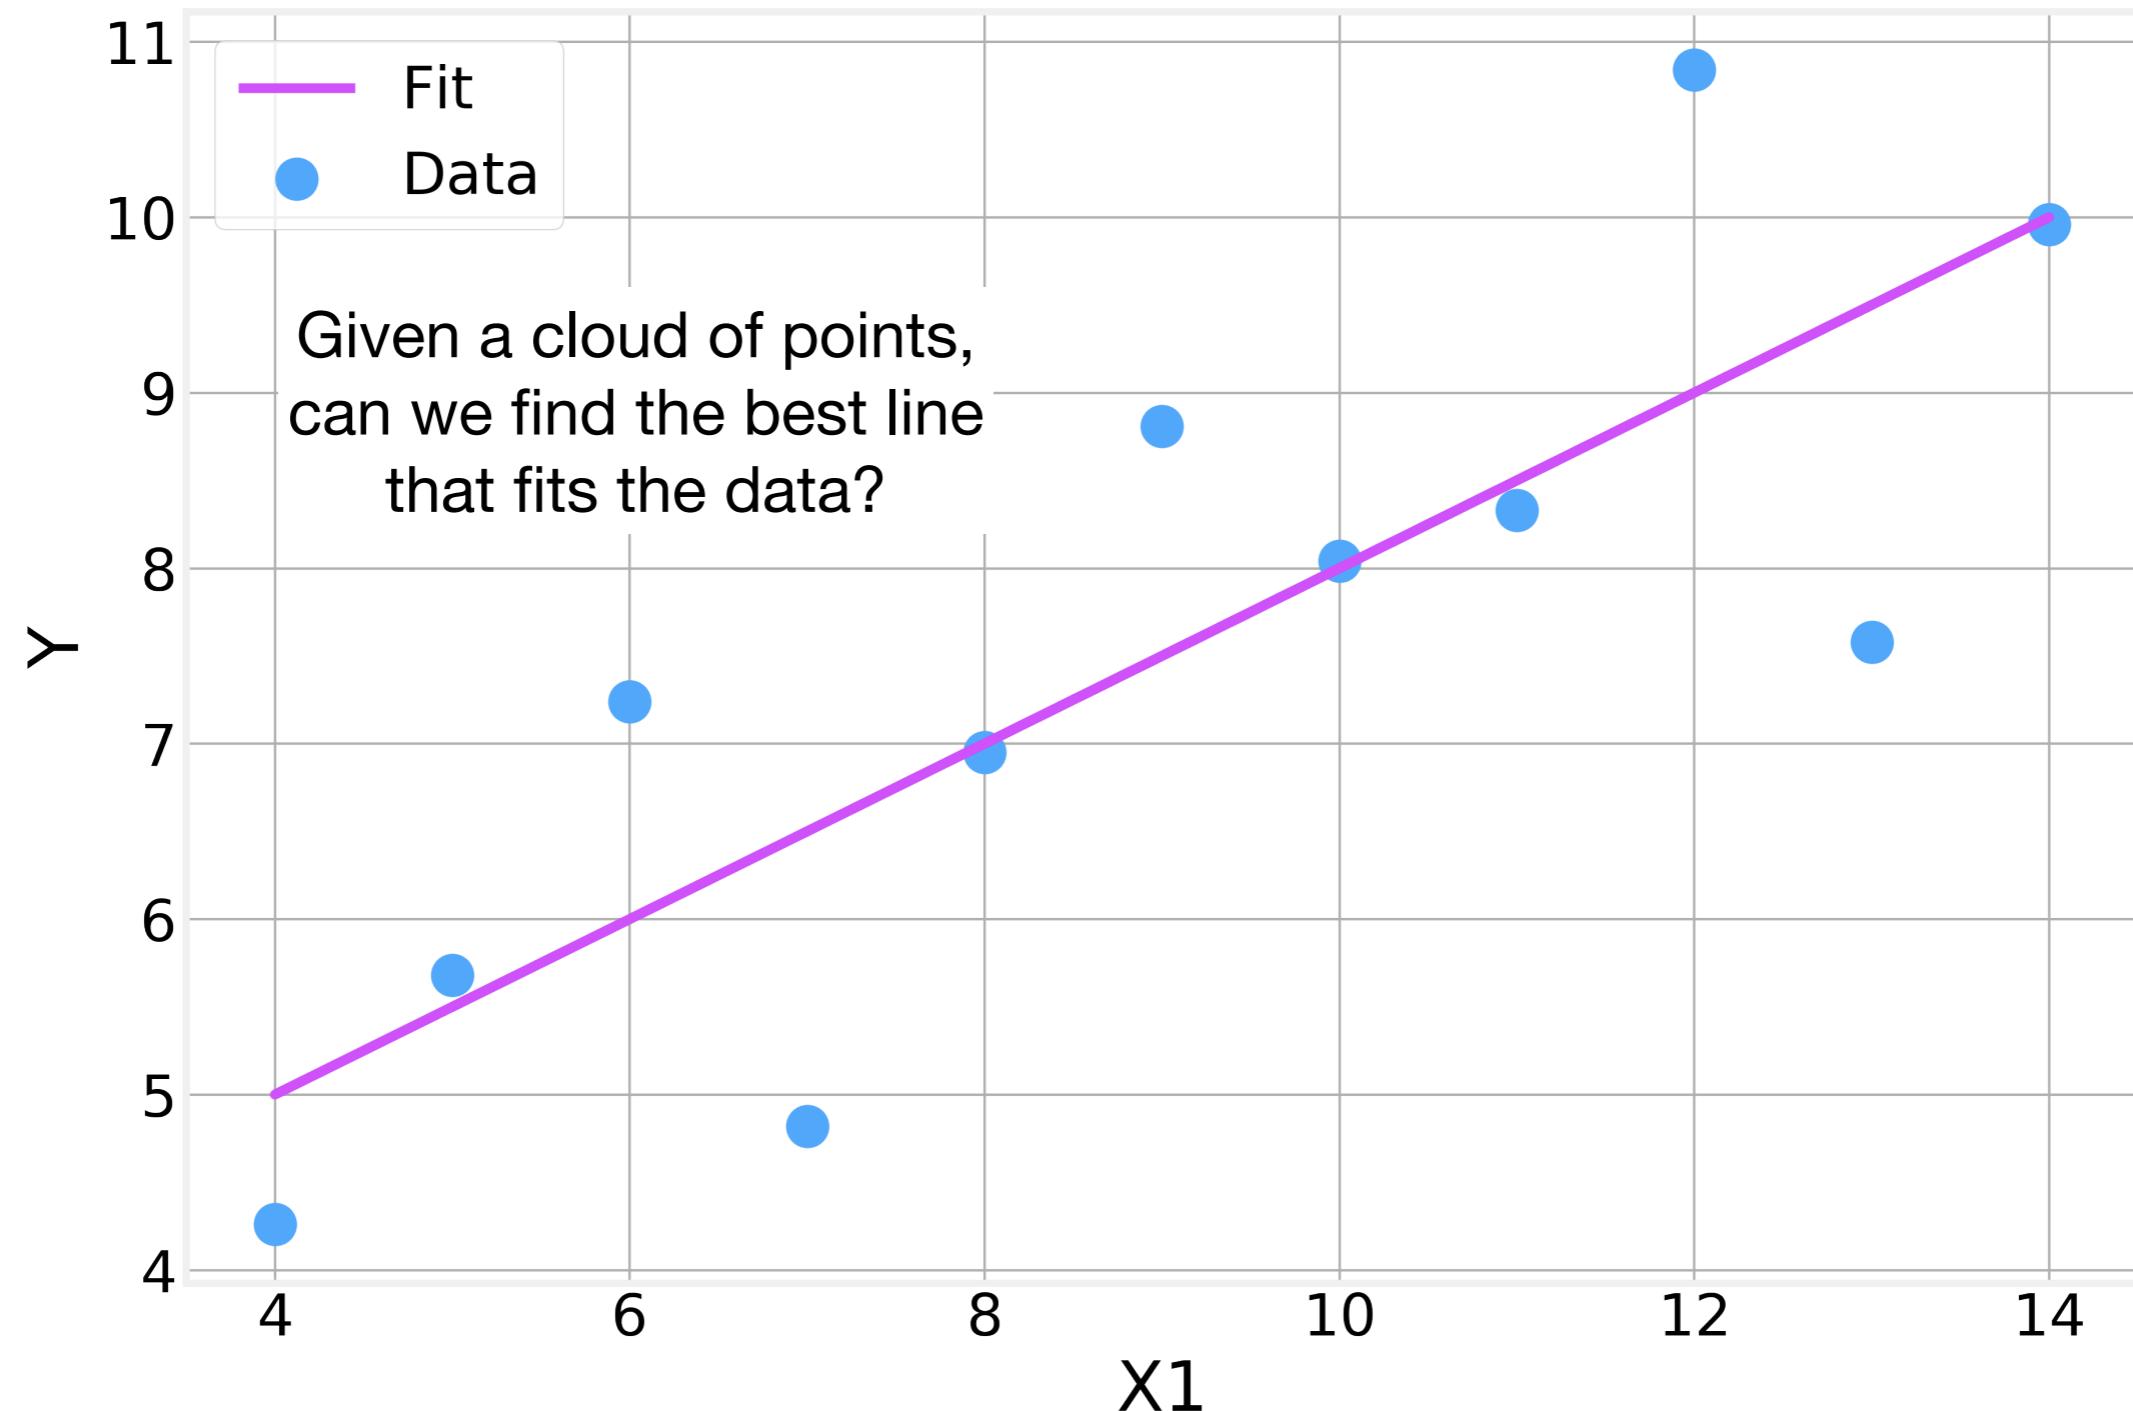
# Linear Regression

[https://en.wikipedia.org/wiki/Linear\\_regression](https://en.wikipedia.org/wiki/Linear_regression)



# Linear Regression

[https://en.wikipedia.org/wiki/Linear\\_regression](https://en.wikipedia.org/wiki/Linear_regression)



# Linear Regression

[https://www.statsmodels.org/devel/generated/statsmodels.regression.linear\\_model.OLS.html](https://www.statsmodels.org/devel/generated/statsmodels.regression.linear_model.OLS.html)

- **statsmodels** returns:

- a description of the model and fitting procedure,
- the **coefficients** of the fitted function,
- estimates of the **error**,
- **statistical significance**,
- and a wealth of statistical test results.

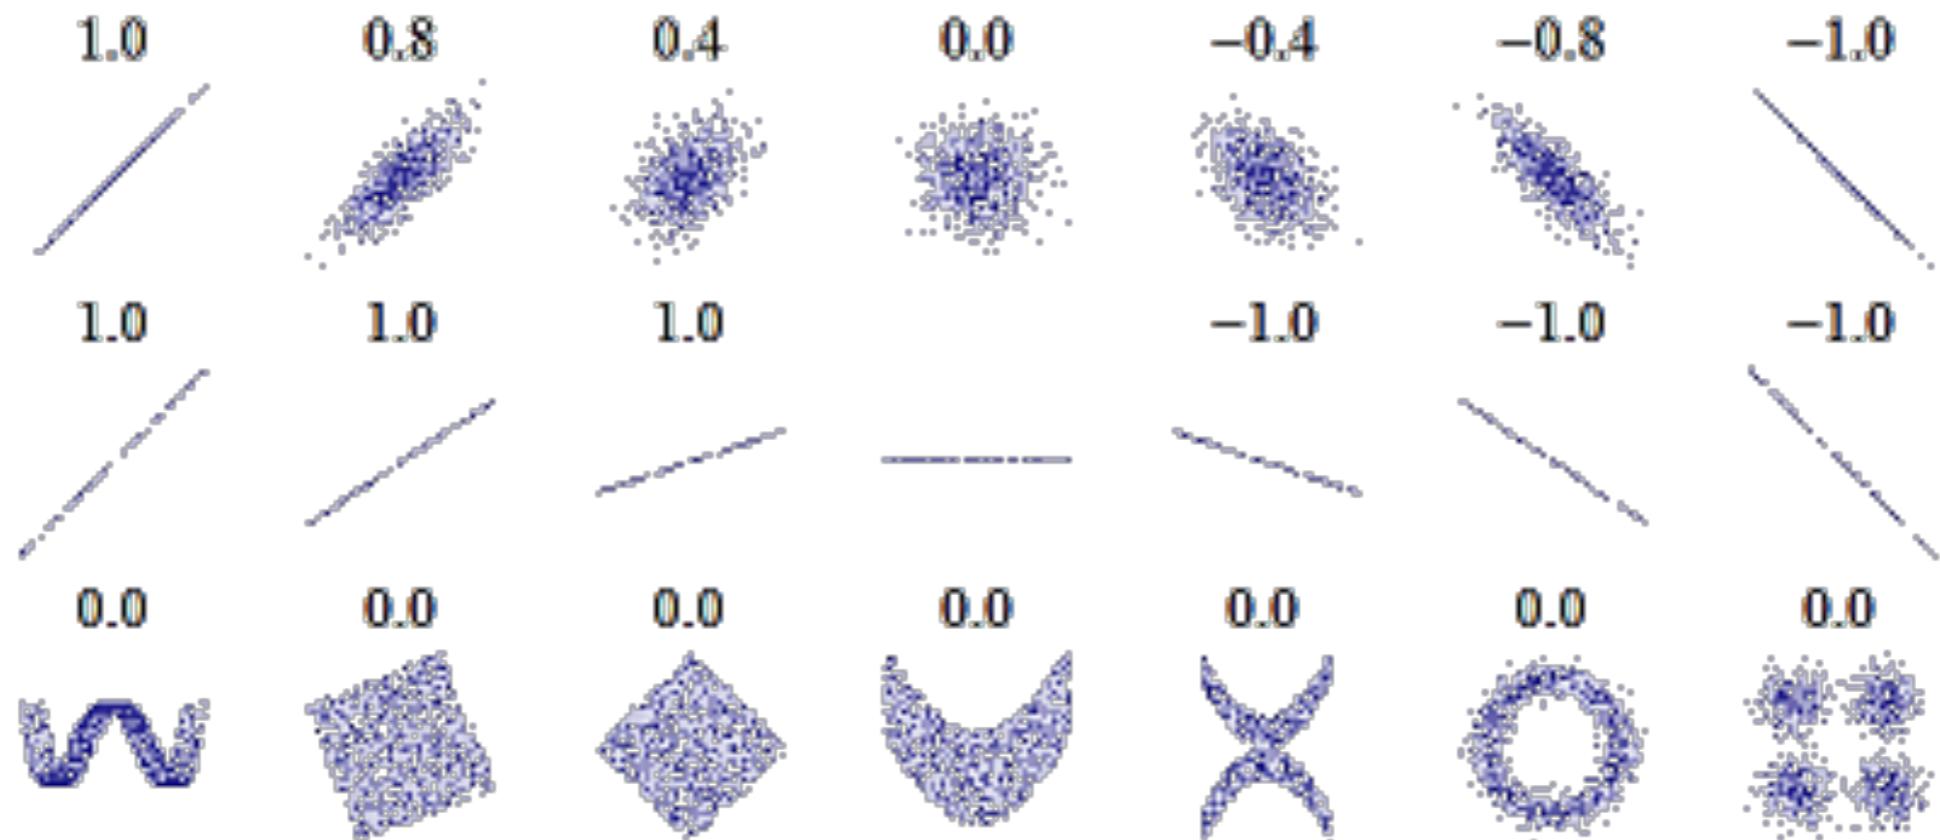
Dep. Variable:		y	R-squared:	0.667
Model:		OLS	Adj. R-squared:	0.629
Method:		Least Squares	F-statistic:	17.99
Date:		Sun, 09 Aug 2020	Prob (F-statistic):	0.00217
Time:		12:37:16	Log-Likelihood:	-16.841
No. Observations:		11	AIC:	37.68
Df Residuals:		9	BIC:	38.48
Df Model:		1		
Covariance Type:		nonrobust		
		coef	std err	t P> t  [0.025 0.975]
<b>const</b>		3.0001	1.125	2.667 0.026 0.456 5.544
<b>x1</b>		0.5001	0.118	4.241 0.002 0.233 0.767
		Omnibus:	0.082	Durbin-Watson: 3.212
		Prob(Omnibus):	0.960	Jarque-Bera (JB): 0.289
		Skew:	-0.122	Prob(JB): 0.865
		Kurtosis:	2.244	Cond. No. 29.1

# Pearson Correlation

- Many correlation measures have been proposed over the years
- The most well known one is the **Pearson Correlation**

$$\rho(x, y) = \sum_{i=1}^N \frac{(x_i - \mu_x)(y_i - \mu_y)}{\sigma_x \sigma_y}$$

- Assumes a **linear relationship** between  $x$  and  $y$ .



# Auto-correlation

- It follows from the previous slide that a series will have a perfect correlation with itself, but what about lagged versions of itself?
- The Auto-correlation function is simply the Pearson correlation between values of the time series at different lags, as a function of the lag:

$$ACF_x(l) = \rho(x_t, x_{t-l})$$

- By definition,  $ACF_x(0) \equiv 1$
- And as the lag  $l$  increases the value of the  $ACF$  tends to decrease.
- We can calculate the confidence interval for the  $ACF$  using:

$$CI = \pm z_{1-\alpha/2} \sqrt{\frac{1}{N} \left( 1 + 2 \sum_{l=1}^k r_l^2 \right)}$$

- where  $z_{1-\alpha/2}$  is the quantile of the normal distribution corresponding to significance level  $\alpha$  and  $r_l$  are the values of the  $ACF$  for a specific lag  $l$

# Auto-correlation

<https://www.statsmodels.org/devel/generated/statsmodels.tsa.stattools.acf.html>

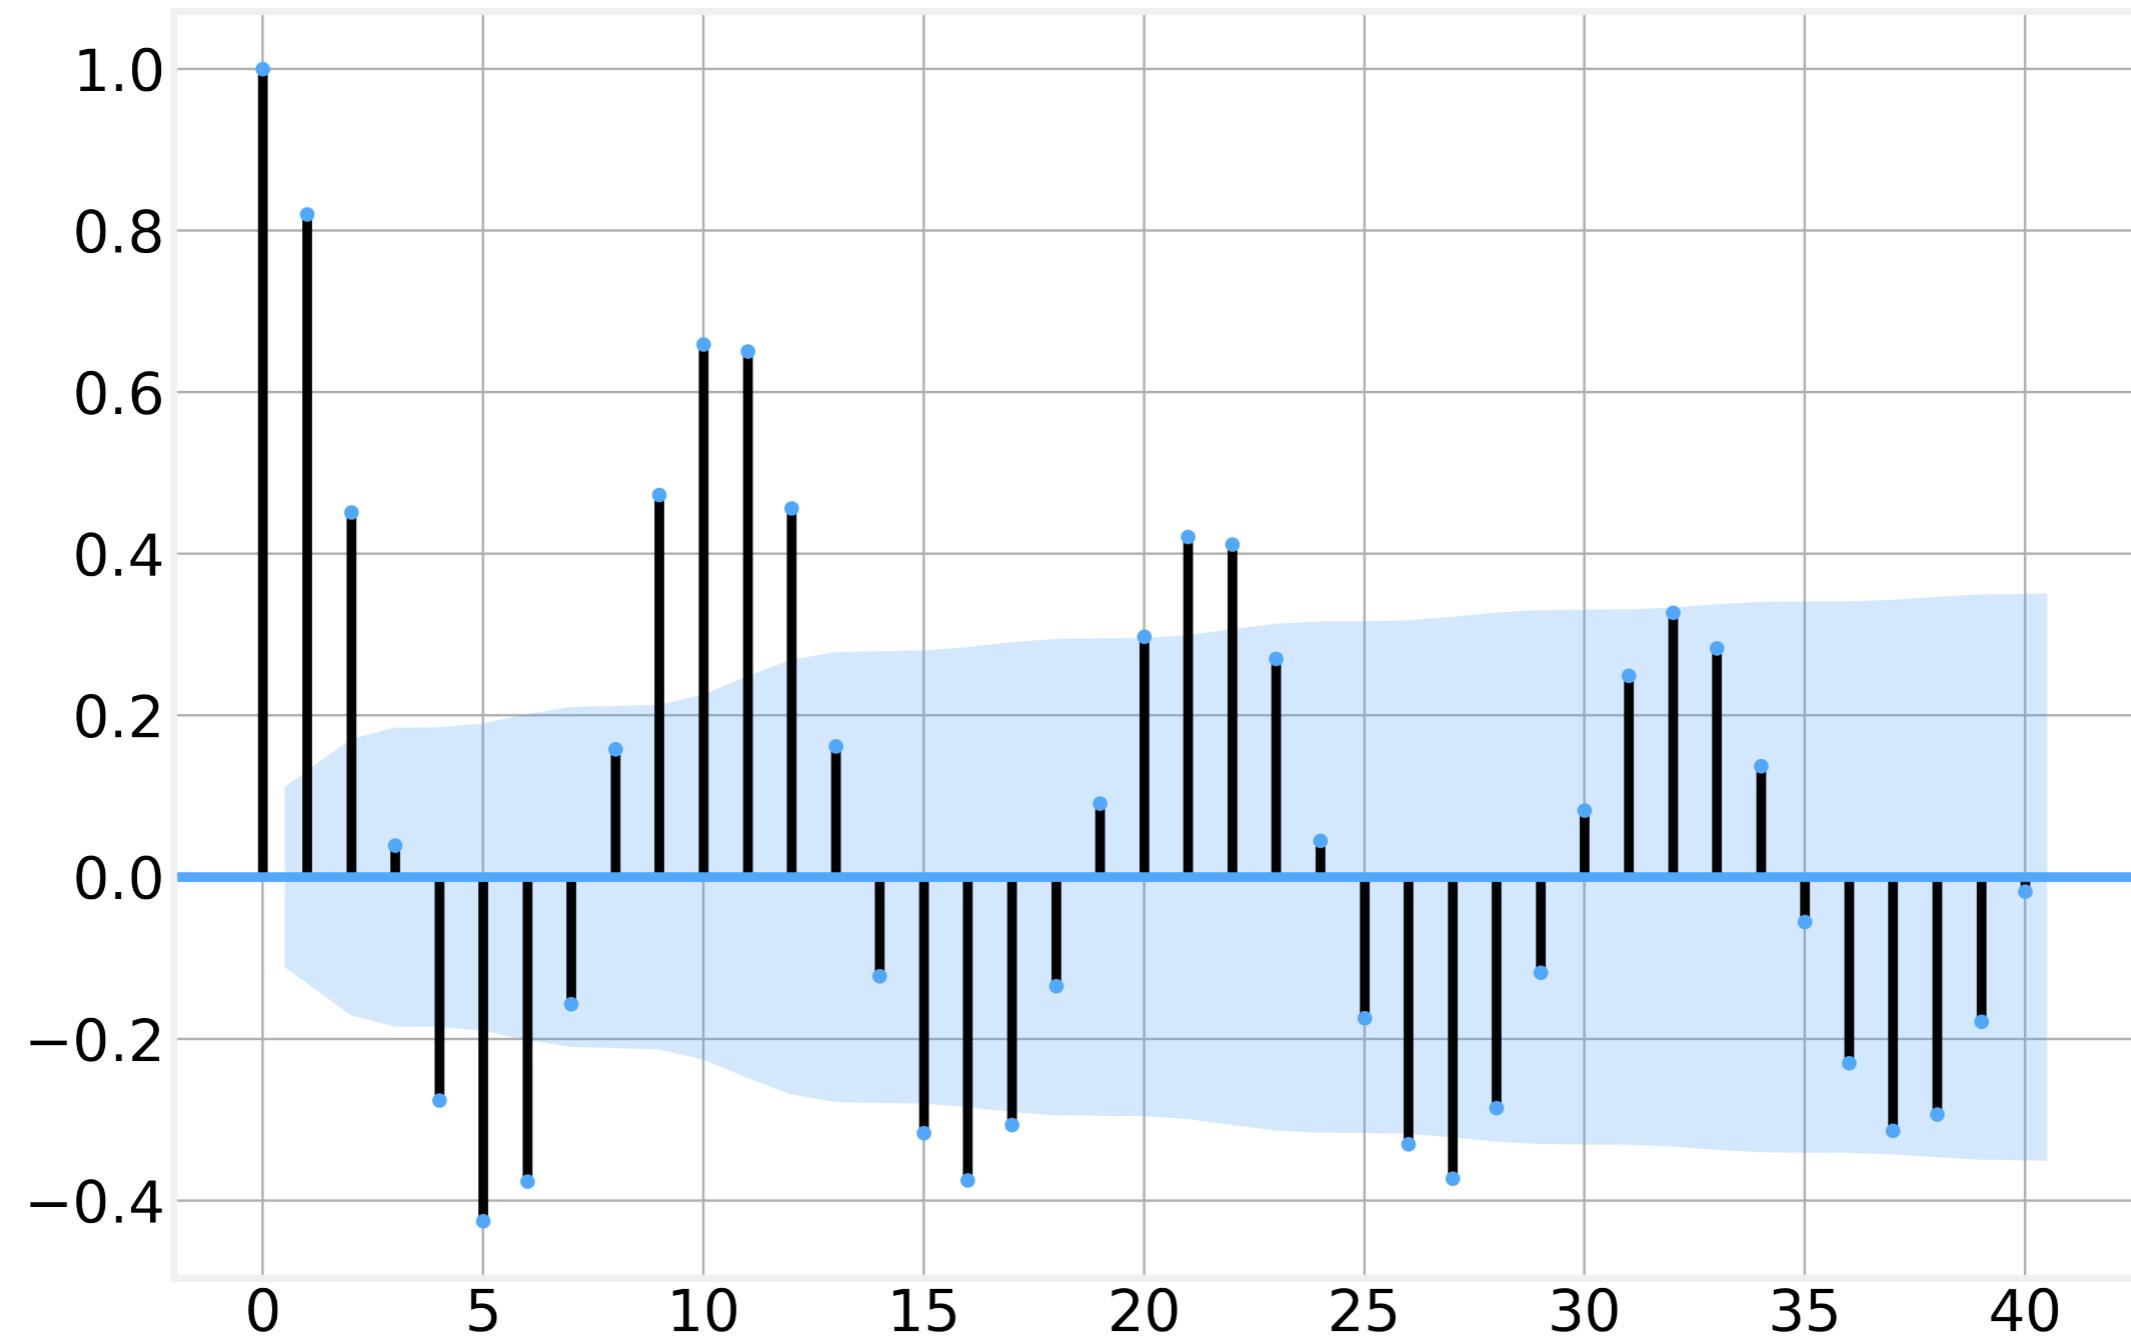
[https://www.statsmodels.org/devel/generated/statsmodels.graphics.tsaplots.plot\\_acf.html](https://www.statsmodels.org/devel/generated/statsmodels.graphics.tsaplots.plot_acf.html)

- **statsmodels** provides us with the **acf()** and **plot\_acf()** methods in the **tsa.stattools** and the **graphics.tsa** modules, respectively.
- **acf()** calculates the values of the auto-correlation function up to **nlags** (default=40)
- **plot\_acf()** calculates the auto-correlation function and the respective confidence interval up to **lags** (default=the number of significant auto-correlation values)
- Both of these methods include various optional arguments to help customize their behavior.

# Auto-correlation

<https://www.statsmodels.org/devel/generated/statsmodels.tsa.stattools.acf.html>  
[https://www.statsmodels.org/devel/generated/statsmodels.graphics.tsaplots.plot\\_acf.html](https://www.statsmodels.org/devel/generated/statsmodels.graphics.tsaplots.plot_acf.html)

## Autocorrelation



# Partial Autocorrelation

[https://en.wikipedia.org/wiki/Partial\\_autocorrelation\\_function](https://en.wikipedia.org/wiki/Partial_autocorrelation_function)

- One of the disadvantages of the Autocorrelation function is that it still considers the intermediate values
- The Partial Autocorrelation function calculate the correlation function between  $x_t$  and  $x_{t-l}$  after explaining away all the intermediate values  $x_{t-1} \cdots x_{t-l+1}$
- Intermediate values are "explained away" by fitting a linear model of the form:

$$\hat{x}_t = f(x_{t-1} \cdots x_{t-l+1})$$

$$\hat{x}_{t-l} = f(x_{t-1} \cdots x_{t-l+1})$$

- And then calculating the Pearson correlation function between the values and their residuals:

$$PACF_x(l) = \rho(x_t - \hat{x}_t, x_{t-l} - \hat{x}_{t-l})$$

- Confidence intervals can be computed using the same formula used for the **ACF**

# Partial Autocorrelation

<https://www.statsmodels.org/devel/generated/statsmodels.tsa.stattools.pacf.html>

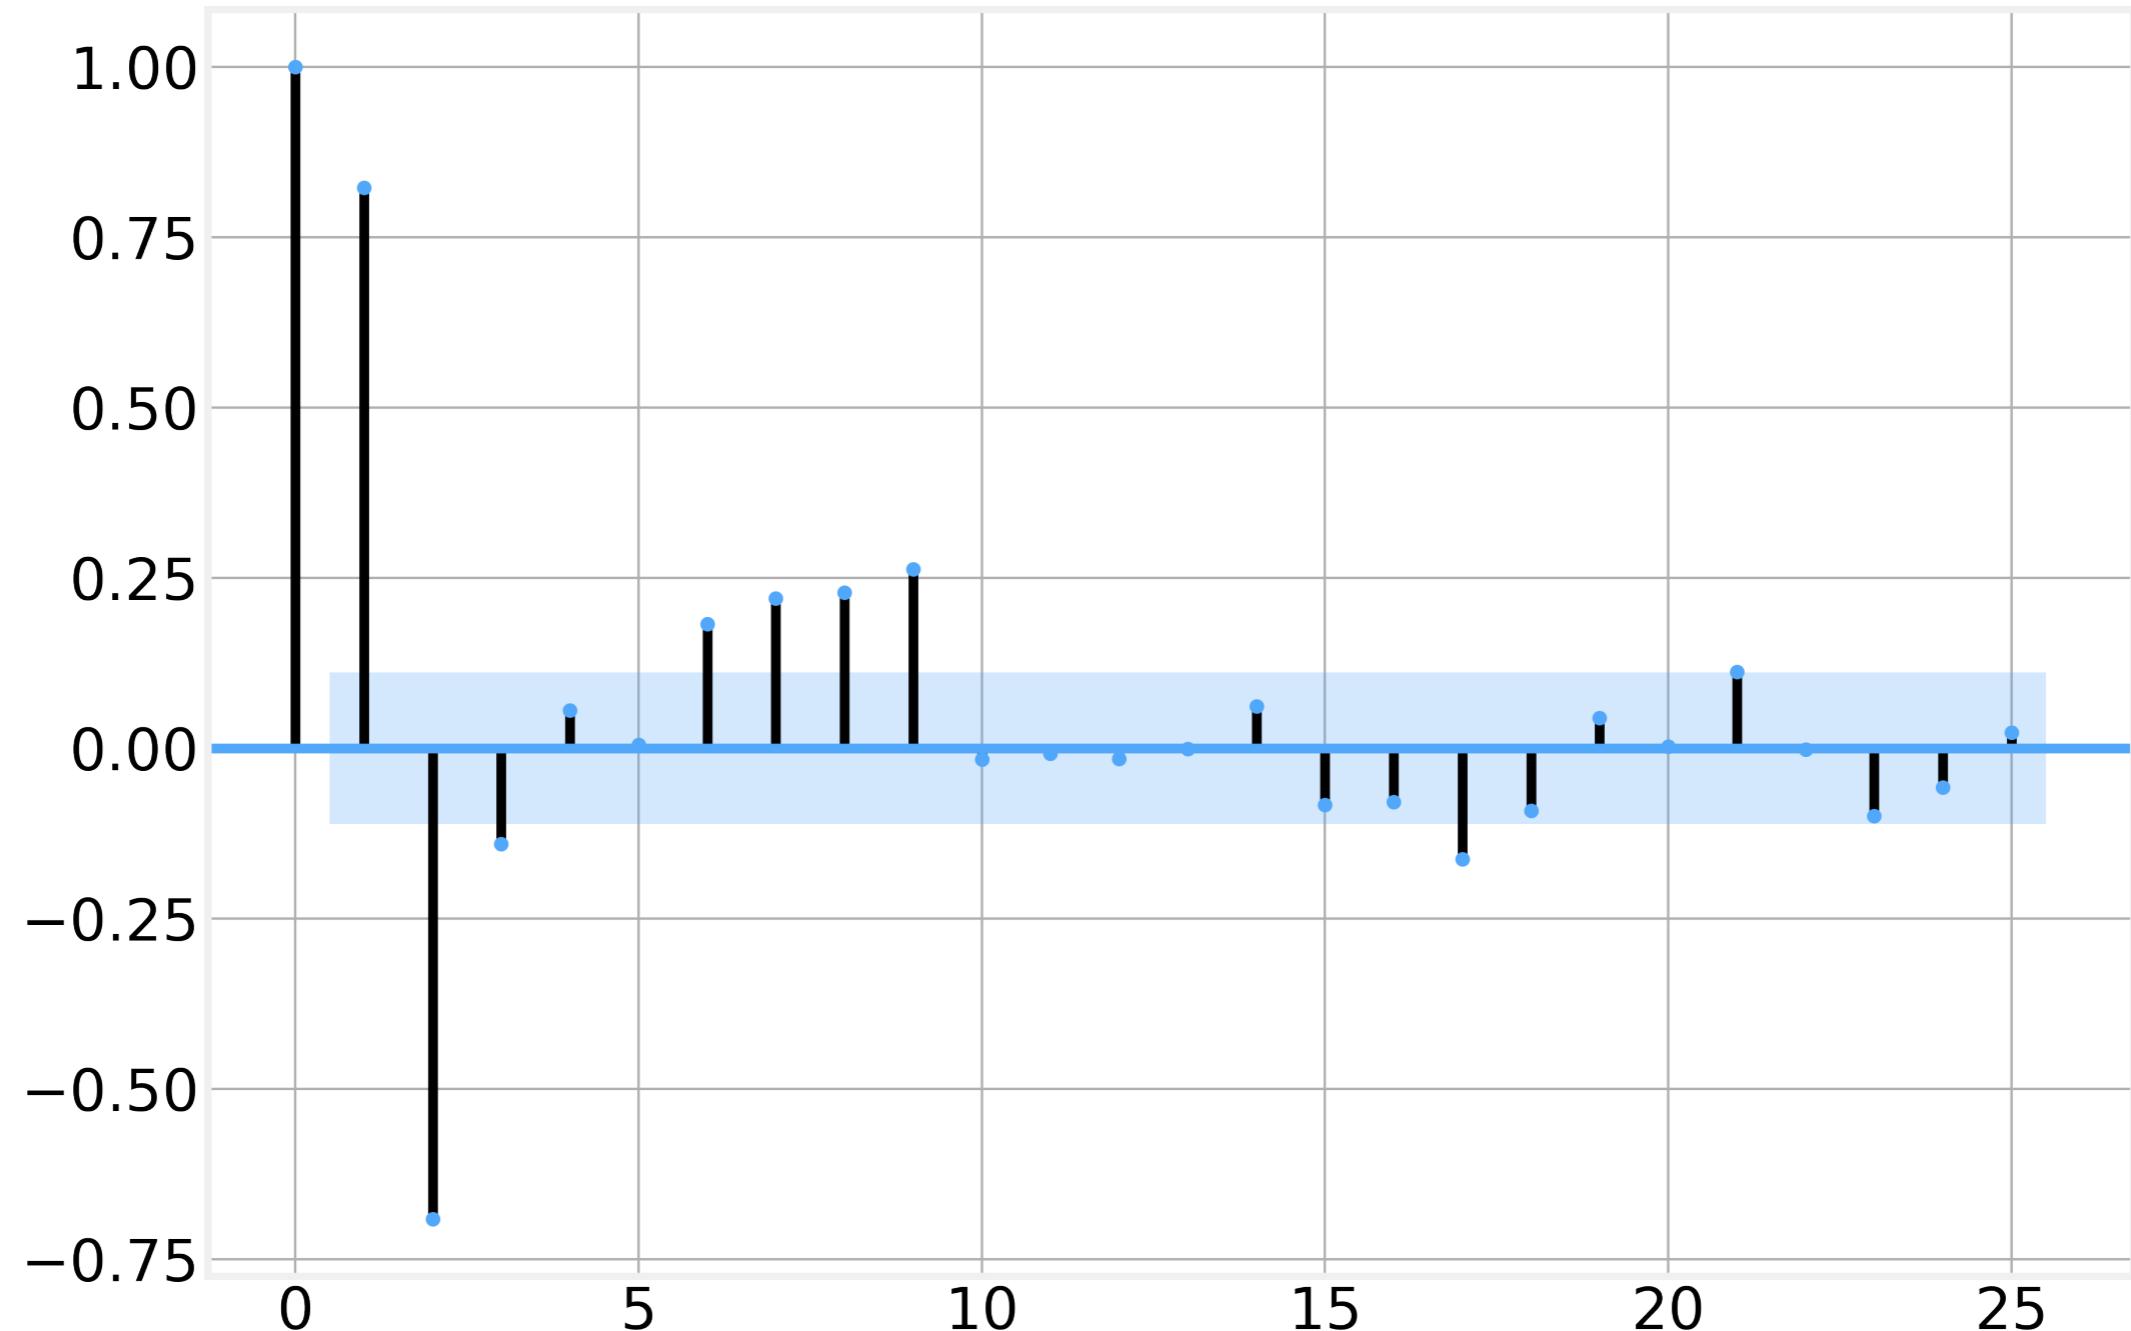
[https://www.statsmodels.org/devel/generated/statsmodels.graphics.tsaplots.plot\\_pacf.html](https://www.statsmodels.org/devel/generated/statsmodels.graphics.tsaplots.plot_pacf.html)

- **statsmodels** uses a similar structure for the partial auto-correlation function with a **pacf()** and **plot\_pacf()** functions provided by the **tsa.statstools** and the **graphics.tsa** modules, respectively.
- **pacf()** calculates the values of the partial auto-correlation function up to **nlags** (default=40)
- **plot\_acf()** calculates the auto-correlation function and the respective confidence interval up to **lags** (default=the number of significant auto-correlation values)
- And as with the **acf/plot\_acf** duo, both of these methods include various optional arguments to help customize their behavior.

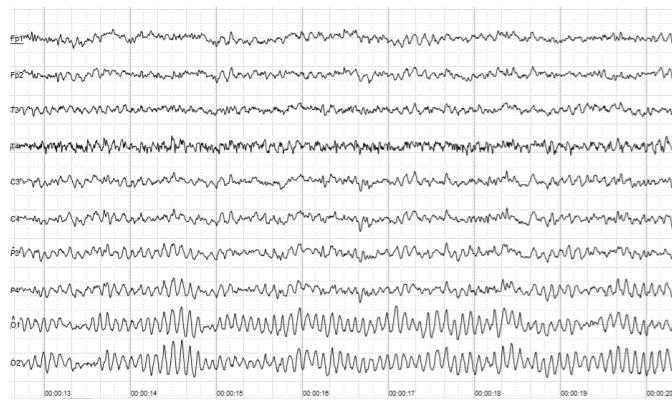
# Partial Autocorrelation

<https://www.statsmodels.org/devel/generated/statsmodels.tsa.stattools.pacf.html>  
[https://www.statsmodels.org/devel/generated/statsmodels.graphics.tsaplots.plot\\_pacf.html](https://www.statsmodels.org/devel/generated/statsmodels.graphics.tsaplots.plot_pacf.html)

## Partial Autocorrelation

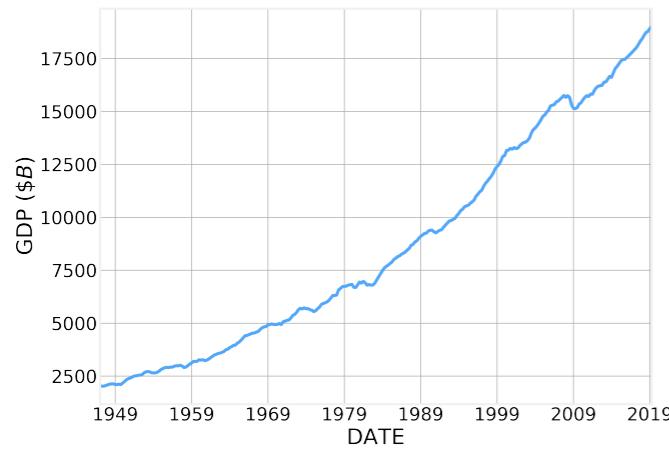


# Three fundamental behaviors



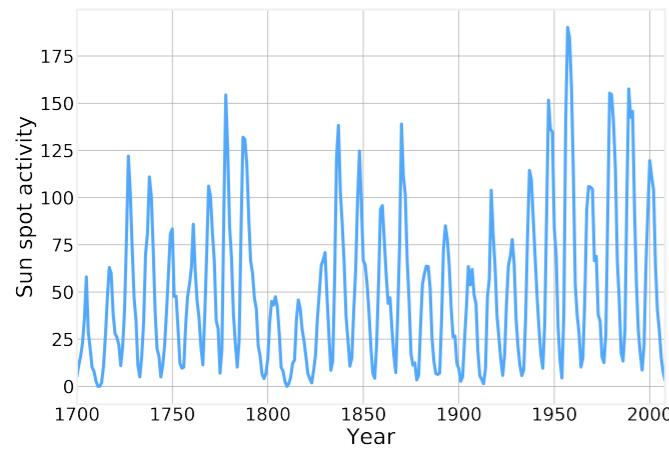
Stationarit

$$\langle x_t \rangle \approx \text{constant}$$



Trend

$$\langle x_t \rangle \approx ct$$



Seasonalit

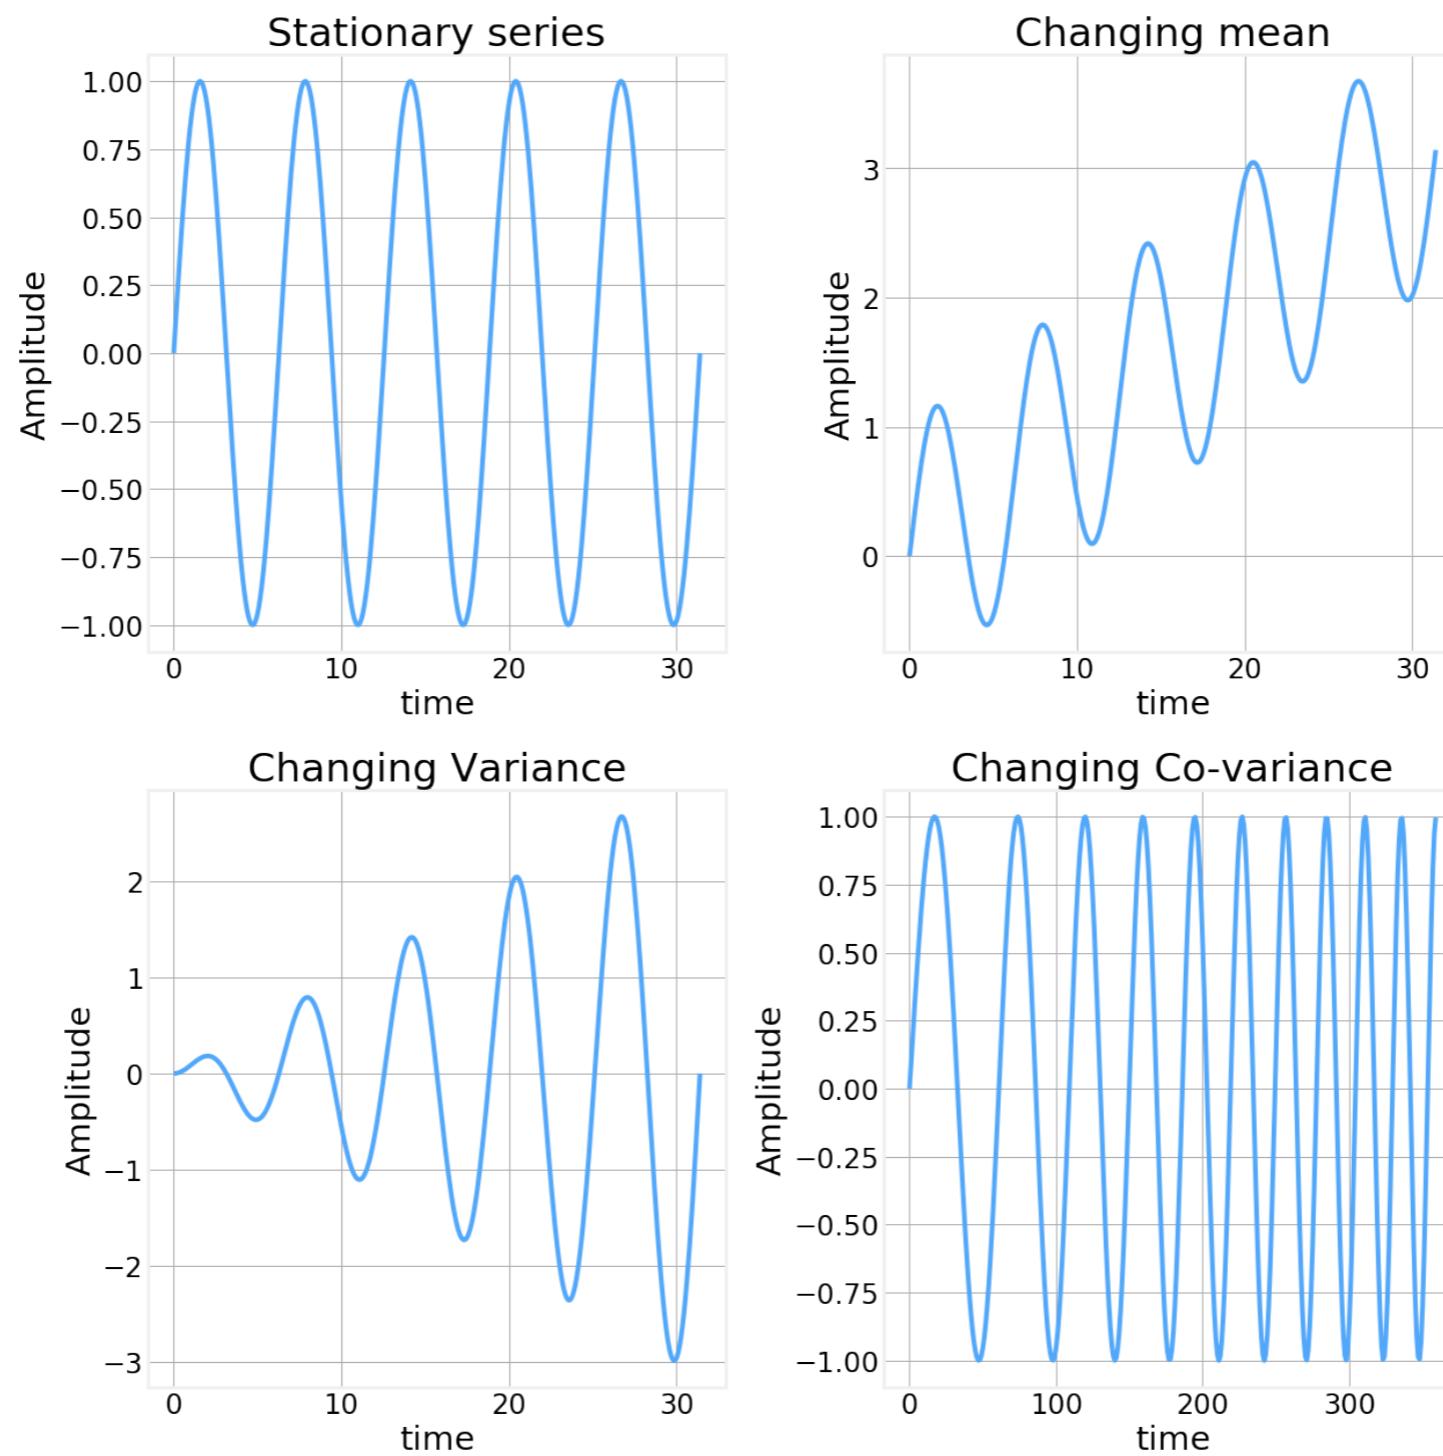
$$x_{t+T} \approx x_t$$

# Stationarity

---

- A time series is said to be stationary if its basic statistical properties are **independent of time**
- In particular:
  - **Mean** - Average value stays constant
  - **Variance** - The width of the curve is bounded
  - **Covariance** - Correlation between points is independent of time
- Stationary processes are **easier** to analyze. If the parameters of the model vary in time, the number of parameters we must estimate **increases significantly**.
- Many time series analysis algorithms **assume stationarity**.
- Several rigorous tests for stationarity have been developed such as the **(Augmented) Dickey-Fuller** and **Hurst Exponent**

# Stationarity



# Dickey-Fuller Test

- The **Dickey-Fuller Test** is a test of stationarity inspired by a simple: Can we show that

$$\rho \neq 1$$

- with some degree of certainty?
- Numerically, we can express this as a **linear regression** fit

$$x_t - x_{t-1} = \gamma x_{t-1} + \epsilon_t$$

- where  $\gamma = \rho - 1$
- The slope of the regression is then our expected value for  $\rho - 1$ .
- If the process is **non-stationary** then we expect  $\gamma \neq 0$ .

# Dickey-Fuller Test

---

- **statsmodels** provides an implementation of the **Augmented Dickey-Fuller** test with the function **adftest()** in the **tsa.stattools** module
- **adftest()** is able to automatically handle multiple lags, adjust for trends, etc.
- **adftest()**, when using the parameter **regresults=True**, returns a tuple with 4 elements:
  - **adfstat** - the value of the augmented dickey fuller statistic
  - **p-value** - the p-value (statistical significance) of the result
  - **criticalvalues** - the critical values for the test statistic at the 1%, 5%, and 10% levels
  - **results** - an object containing more information, including the results for the linear regression
- We encourage you to explore the various arguments to this function to fully understand its functionality.

# Time series decomposition

---

- A time series can be decomposed into **three components**:
  - Trend,  $T_t$
  - Seasonality,  $S_t$
  - Residuals,  $R_t$
- Decompositions can be
  - **additive** -  $x_t = T_t + S_t + R_t$
  - **multiplicative** -  $x_t = T_t \cdot S_t \cdot R_t$
- The residuals are simply what is left of the original signal after we **remove the trend and the seasonality**
- Residuals are typically **stationary**

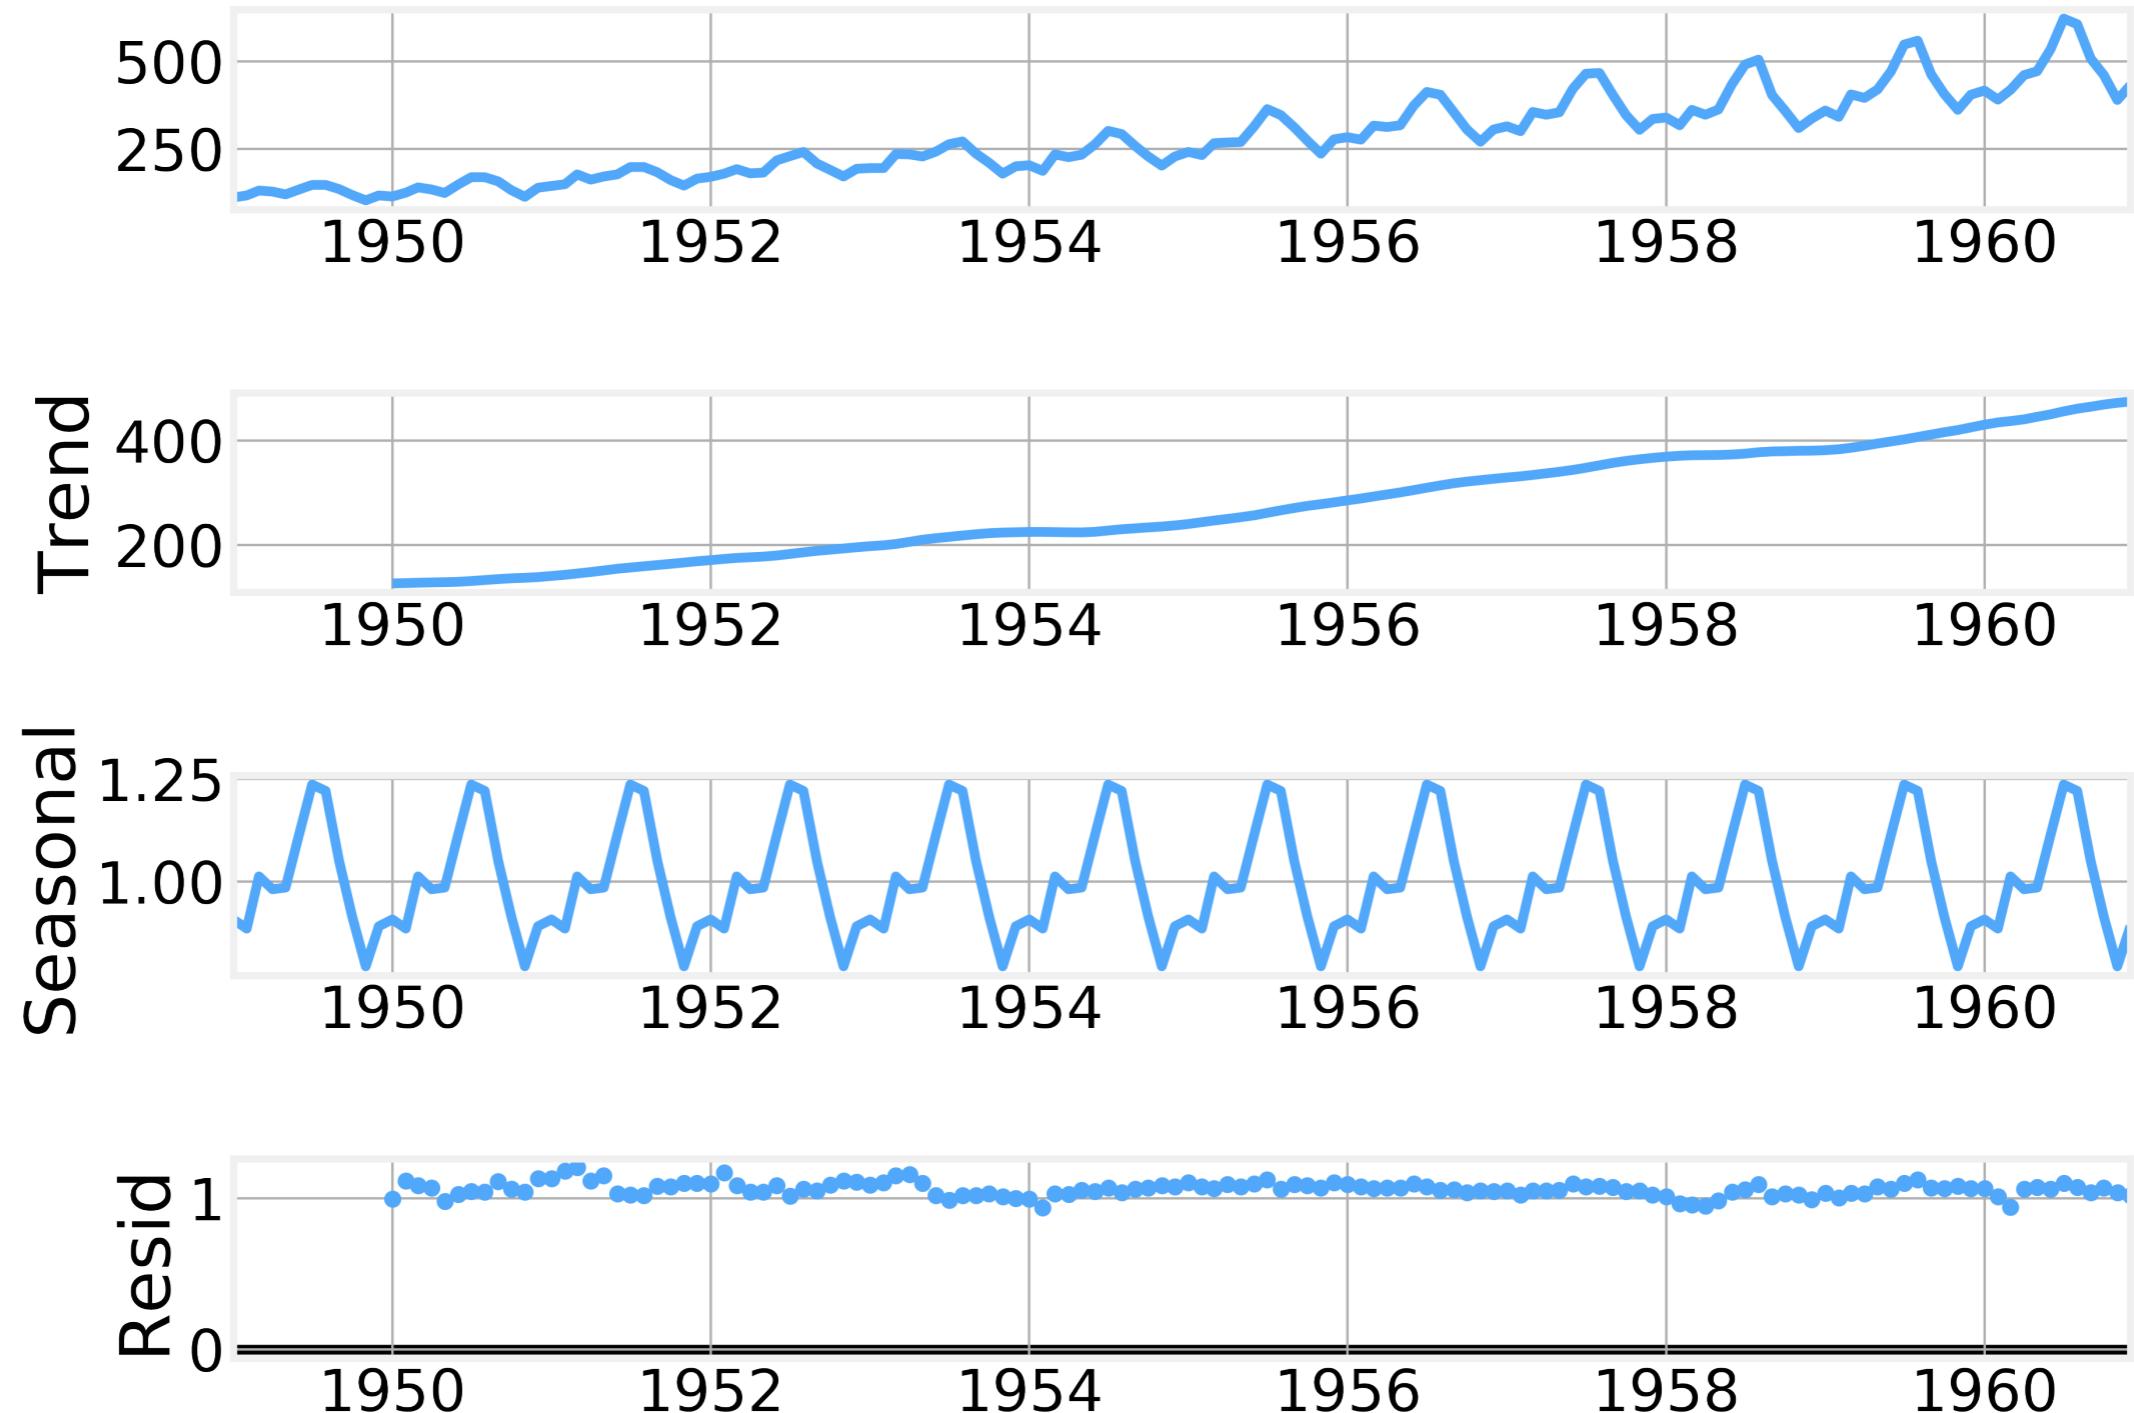
# Time series decomposition

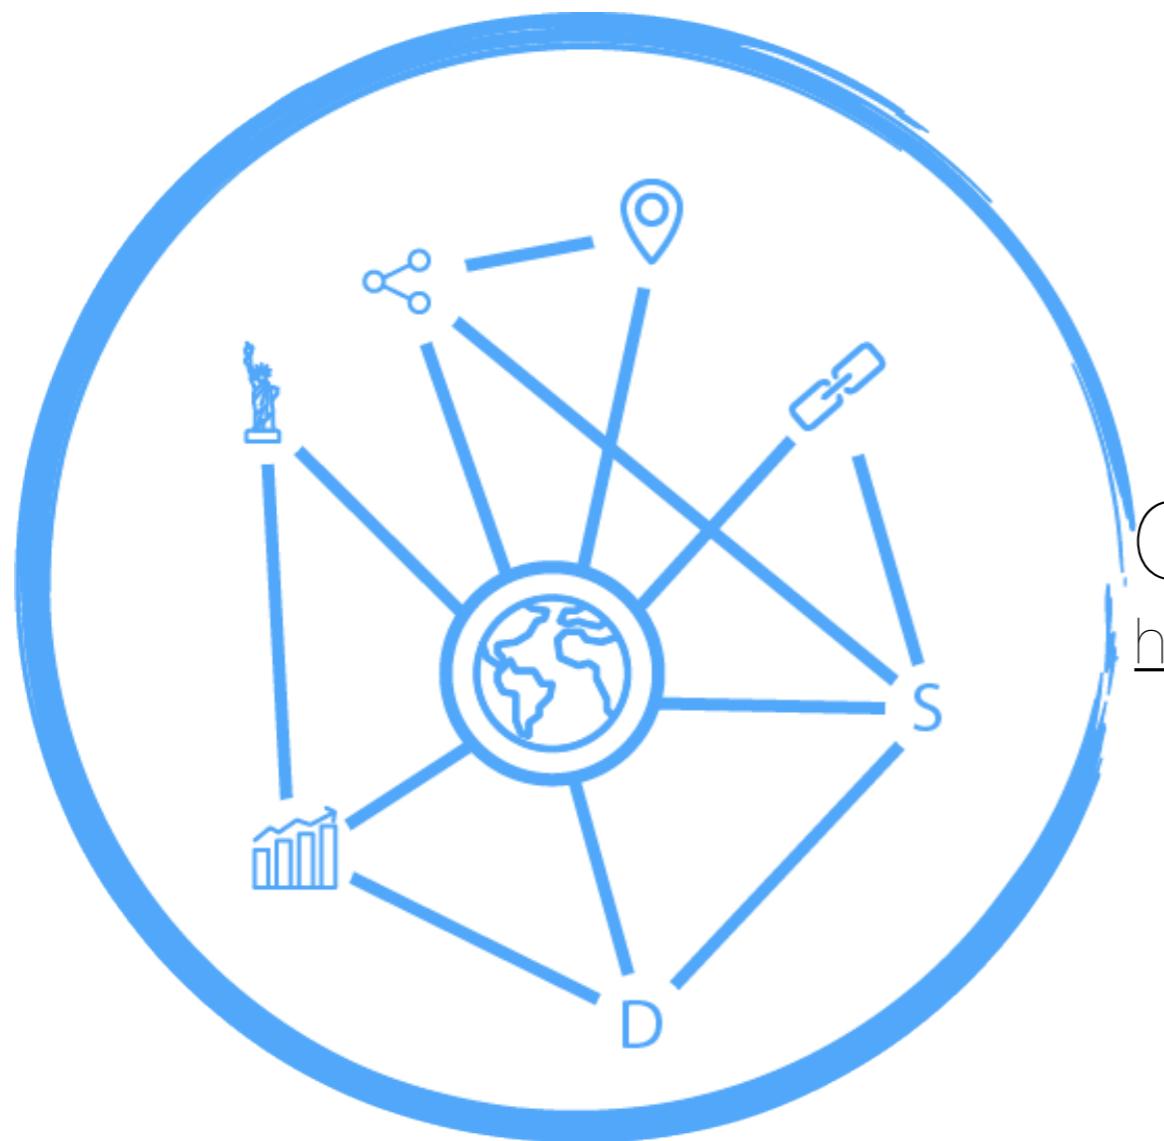
[https://www.statsmodels.org/dev/generated/statsmodels.tsa.seasonal.seasonal\\_decompose.html](https://www.statsmodels.org/dev/generated/statsmodels.tsa.seasonal.seasonal_decompose.html)

- **statsmodels** provide the **seasonal\_decompose()** function in the **tsa.seasonal** module to perform the seasonal decomposition described above.
- The **model** argument allows us to select between “**additive**” or “**multiplicative**” decomposition.
- The **two\_sided** argument lets us decide how to calculate the moving average. We should set it to **False** to make sure the moving average values are computed only past values.
- It returns an object with **observed**, **seasonal**, **trend**, and **resid** attributes as **pandas** DataFrames. The **plot()** method generates the usual decomposition plot.

# Time series decomposition

[https://www.statsmodels.org/dev/generated/statsmodels.tsa.seasonal.seasonal\\_decompose.html](https://www.statsmodels.org/dev/generated/statsmodels.tsa.seasonal.seasonal_decompose.html)





Code - statsmodels

<https://github.com/DataForScience/AdvancedTimeSeries>



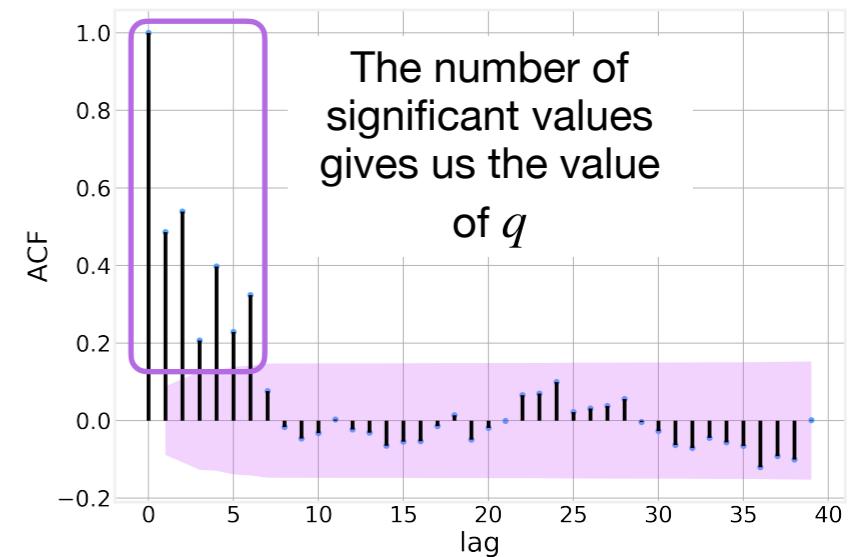
## Lesson III: ARIMA Models

# Moving Average (MA) Model

- A more general case can be written as:

$$x_t = \beta + \sum_{l=0}^q \omega_l \epsilon_{t-l}$$

- where  $\beta$  is a constant offset,  $\omega_l$  are the weights for the values at lag  $l$  and  $q$  is the moving window size.  $\omega_0 \equiv 1$ .
- The  $\epsilon_t$  values are **stochastic variables** (often referred to as “errors”) rather than the actual **observed values**  $x_t$
- The generated  $x_t$  is **uncorrelated** with itself for any lag  $l > q$

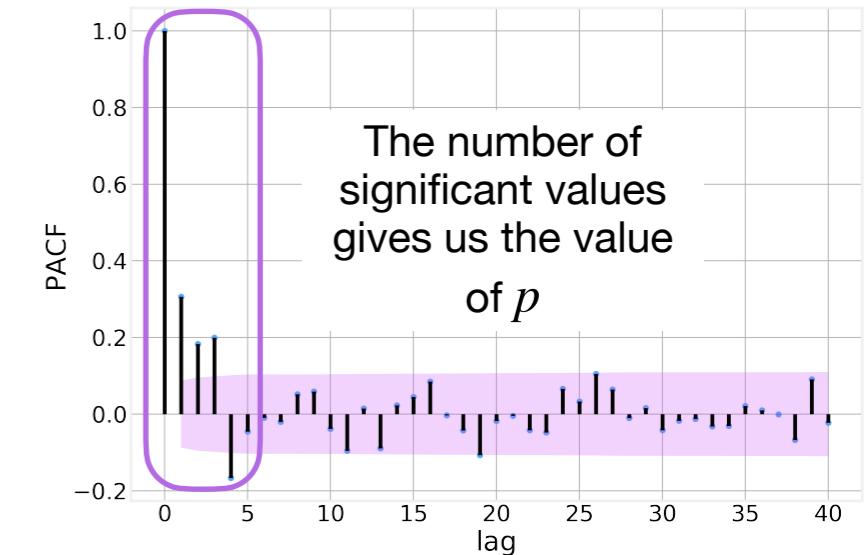


# Auto-Regressive (AR) Models

- **Auto-Regressive** models can be defined as:

$$x_t = \alpha + \epsilon_t \sum_{l=1}^p \phi_l x_{t-l}$$

- Where the constant  $\alpha$  represents the process' average value and the  $x_{t-l}$  are the observed values at a given lag  $l$  and  $\phi_l$  are the corresponding weights.
- The generated  $x_t$  has negligible partial correlation (**uncorrelated**) for any lag  $l > p$
- The series “reverts to the mean”.



## Integrative (I) “model”

- We already saw that we can take differences to “stationarize” the time series.
- To recover the original values, we must then integrate
- While not a model by itself, it is often an important first step in modeling time series

# ARIMA model

- The **ARIMA** class of models is a combination of these three distinct models, each of which accounts for a specific feature of classical time series:

A<sub>uto</sub>  
R<sub>egressive</sub>  
I<sub>ntegrated</sub>  
M<sub>oving</sub>  
A<sub>verage</sub>

$$x_t = c + \sum_i^p \phi_i x_{t-i} + \epsilon_t$$

$$x_t = \mu + \epsilon_t + \sum_j^q \theta_j \epsilon_{t-j}$$

**Regression to the mean** - relying on positions up to lag ***p***

**Integrated** - considering only differences up to order ***d***

**Random Walk** - relying on random variables up to lag ***q***

- The complete model can be written as:

$$\hat{x}_t = c + \mu + \sum_i^p \phi_i x_{t-i} + \sum_j^q \theta_j \epsilon_{t-j} + \epsilon_t$$

- where  $\hat{x}_t$  is the properly differentiated time series
- Naturally, by selecting the proper parameters, we're able to recover any combination of the three models considered

# ARIMA model

---

- **ARIMA (0,1,0)** - Random Walk (with or without drift)

- $$\bullet x_t - x_{t-1} = c + \epsilon_t$$

- **ARIMA (0,0,0)** - White noise (the sequence of stochastic variables)

- $$\bullet x_t = \epsilon_t$$

- **ARIMA (0,1,1)** - Exponential Smoothing

- $$\bullet x_t - x_{t-1} = \epsilon_t + \theta_1 \epsilon_{t-1}$$

- **ARIMA (0,2,2)** - Double exponential Smoothing

- $$\bullet x_t - 2x_{t-1} + x_{t-2} = \epsilon_t + \theta_1 \epsilon_{t-1} + \theta_2 \epsilon_{t-2}$$

# ARIMA model

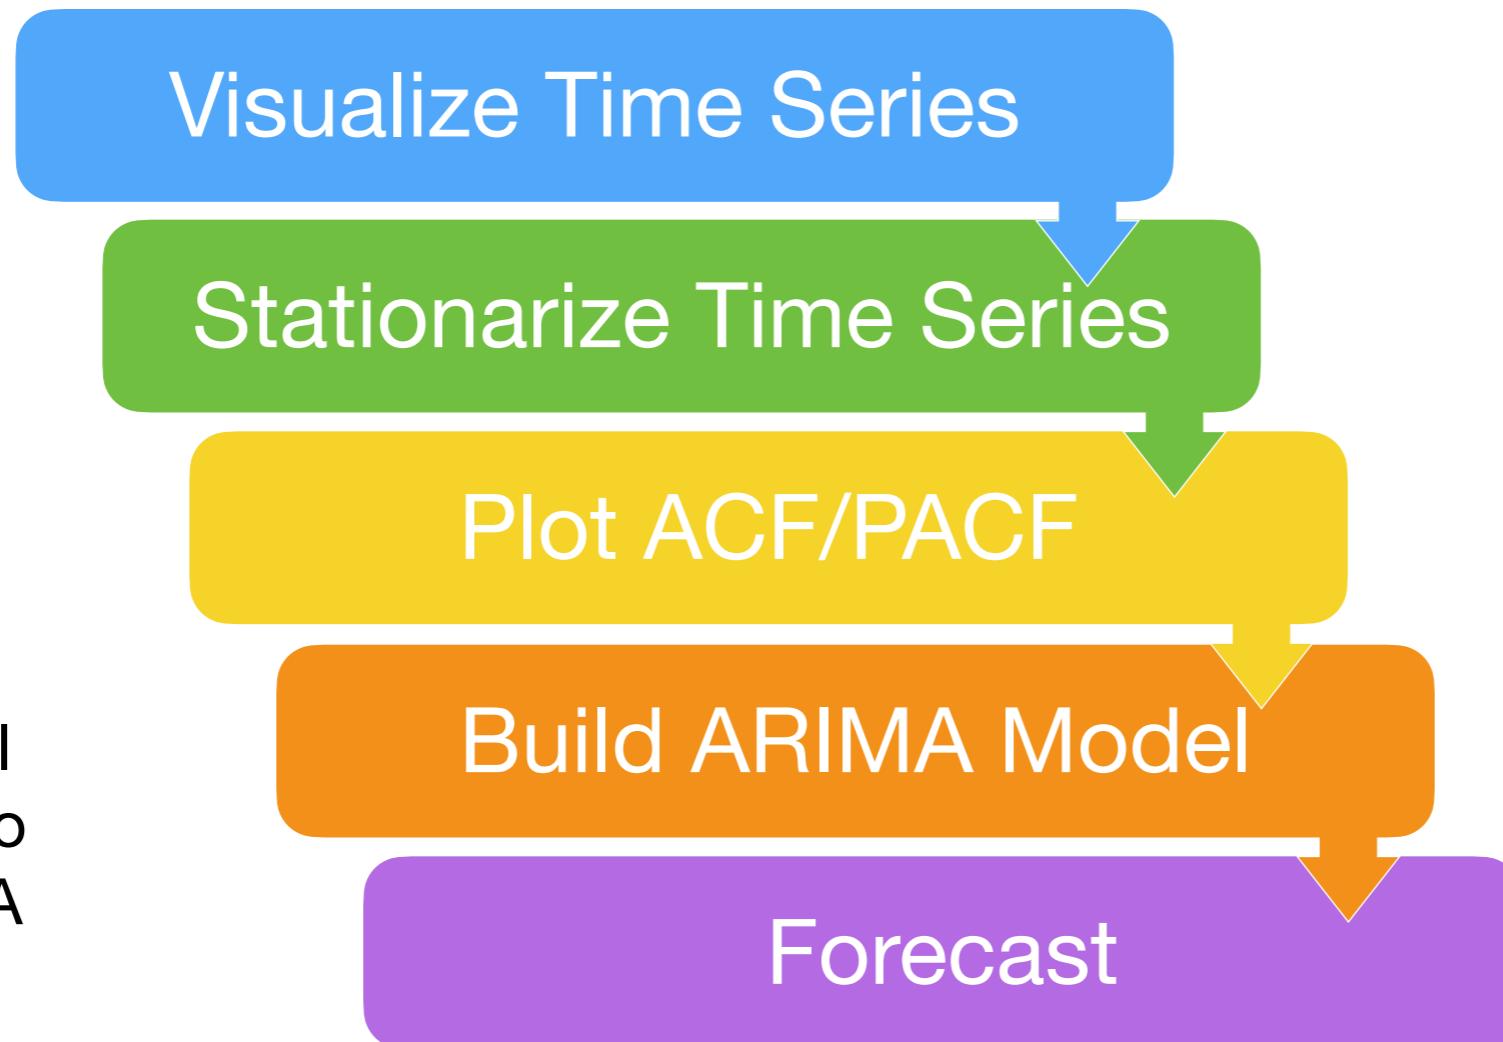
<https://www.statsmodels.org/devel/generated/statsmodels.tsa.arima.model.ARIMA.html>

- **statsmodels** provides the **ARIMA** object in the **tsa.arima** module
- This object allows us to fit a specific model (specified by the argument **order=(p, q, d)**) to a given dataset, i.e. find the specific values of  $\theta_j$  and  $\phi_i$
- The procedure to fit the model is similar to that of **sklearn**, relying two specific methods:
  - **fit()** - Determine the parameters of the model
  - **predict()** - "Run" the model to determine the predicted values of the "position"  $x_t$
- Plus our friendly **summary()** to summarize the results and **plot\_predict()** to compare the original data with the fitted model.

How can we find the values of  $p$ ,  $q$ , and

# Fitting ARIMA models

The general procedure to fit an ARIMA model was originally proposed by Box-Jenkins



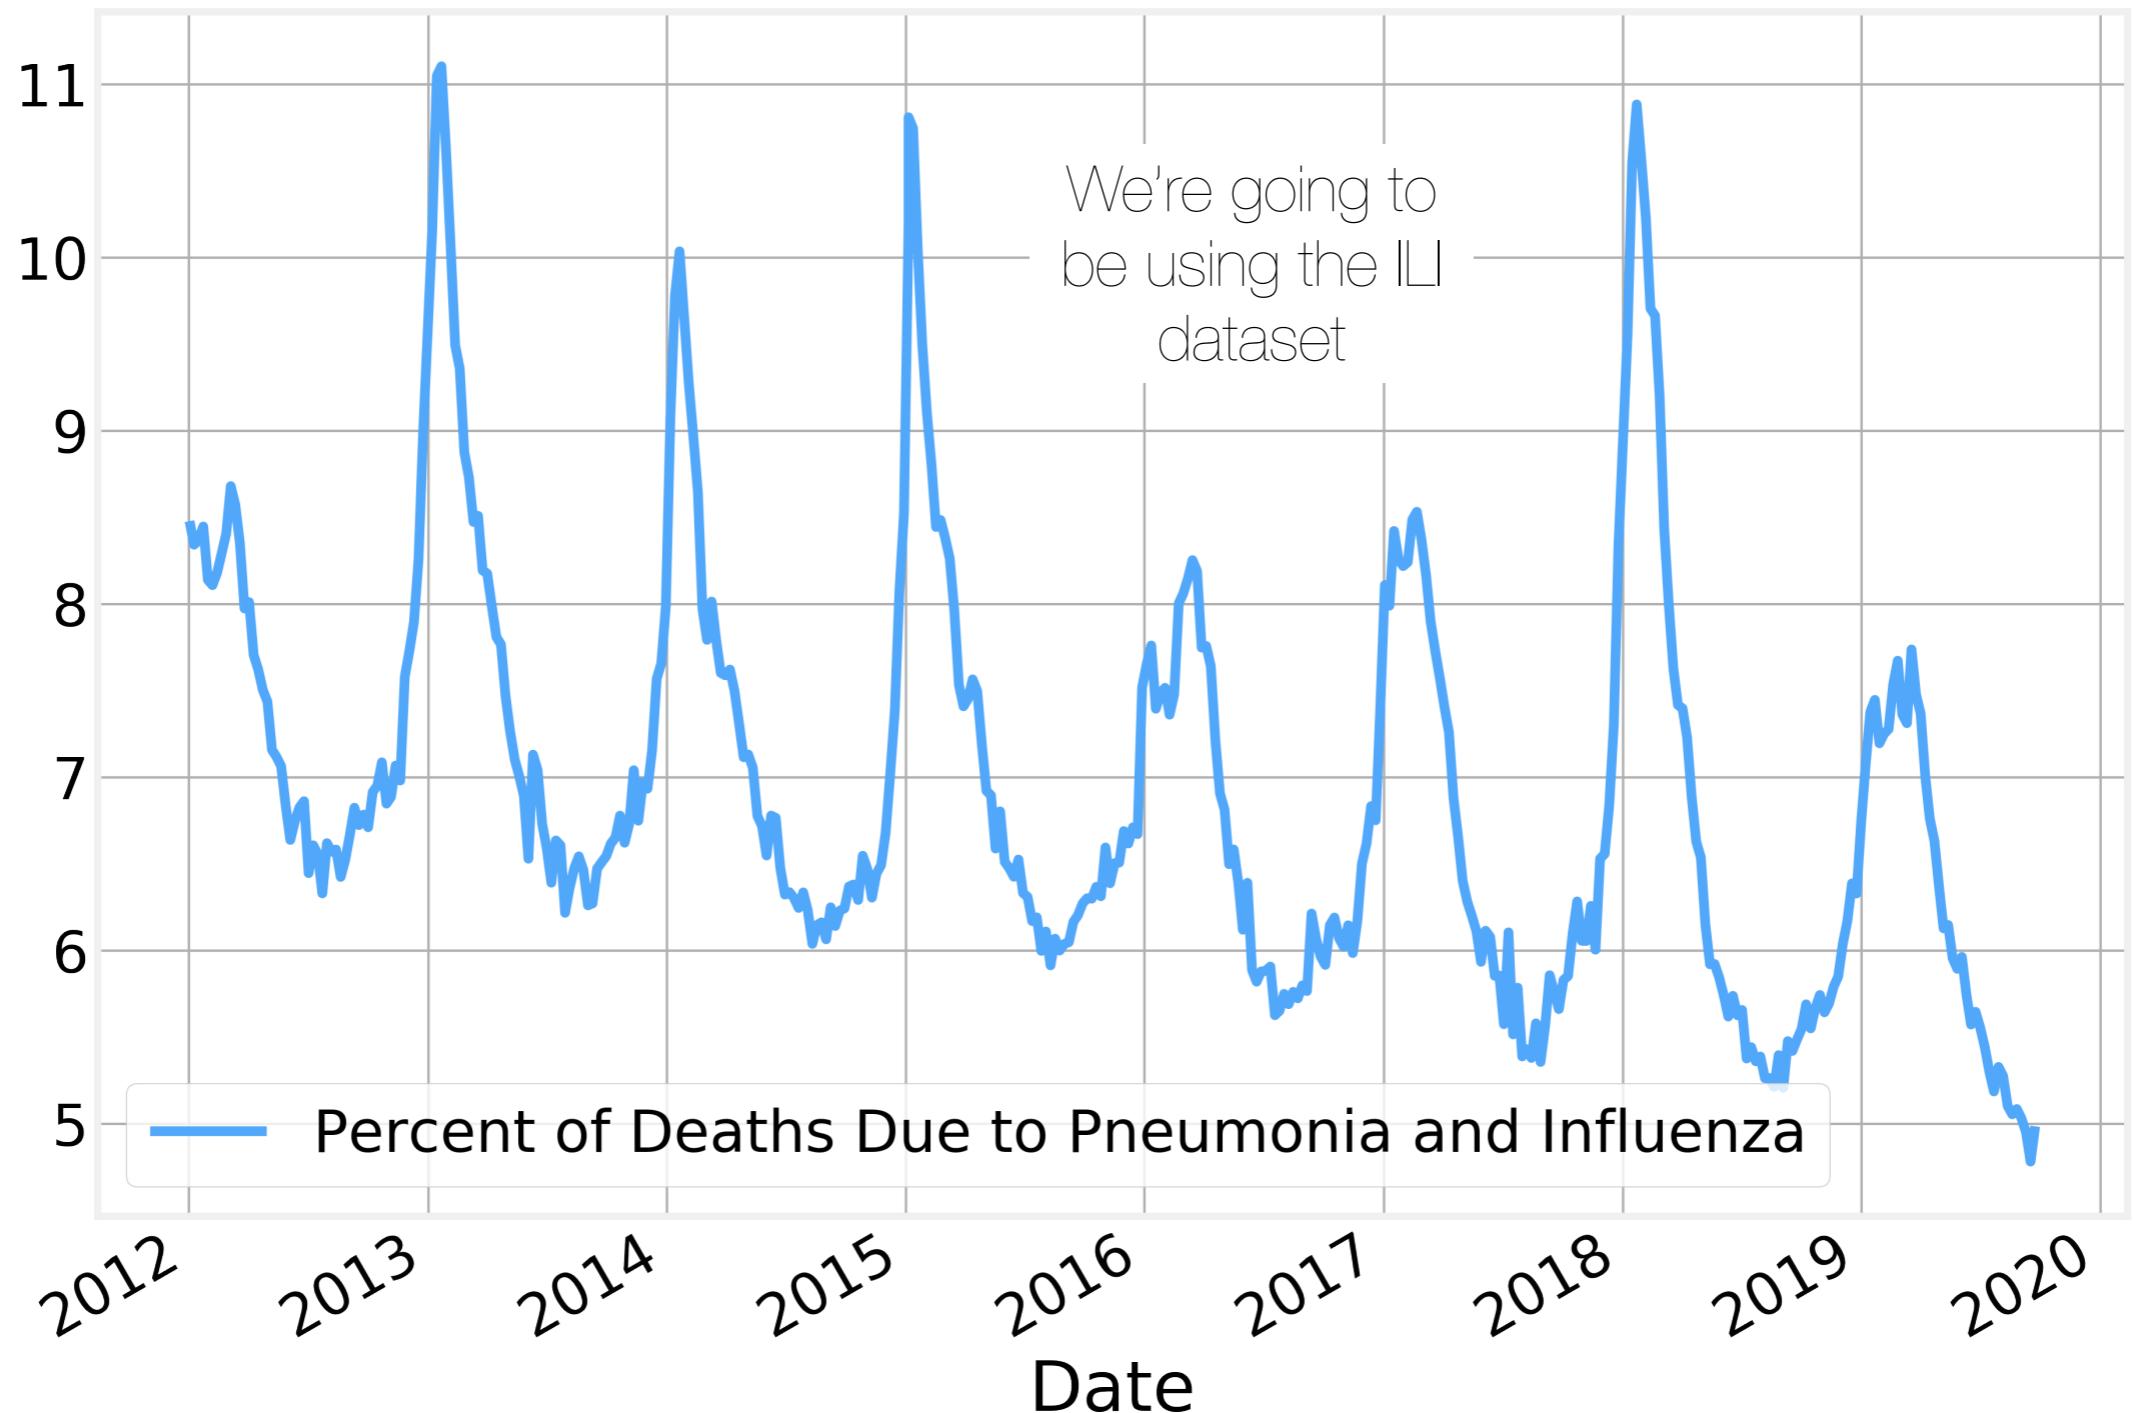
- $d$  - degree of differencing
- $p$  - number of lag observations included in the model ( $PACF$ )
- $q$  - size of the moving average window ( $ACF$ )

# Fitting ARIMA models

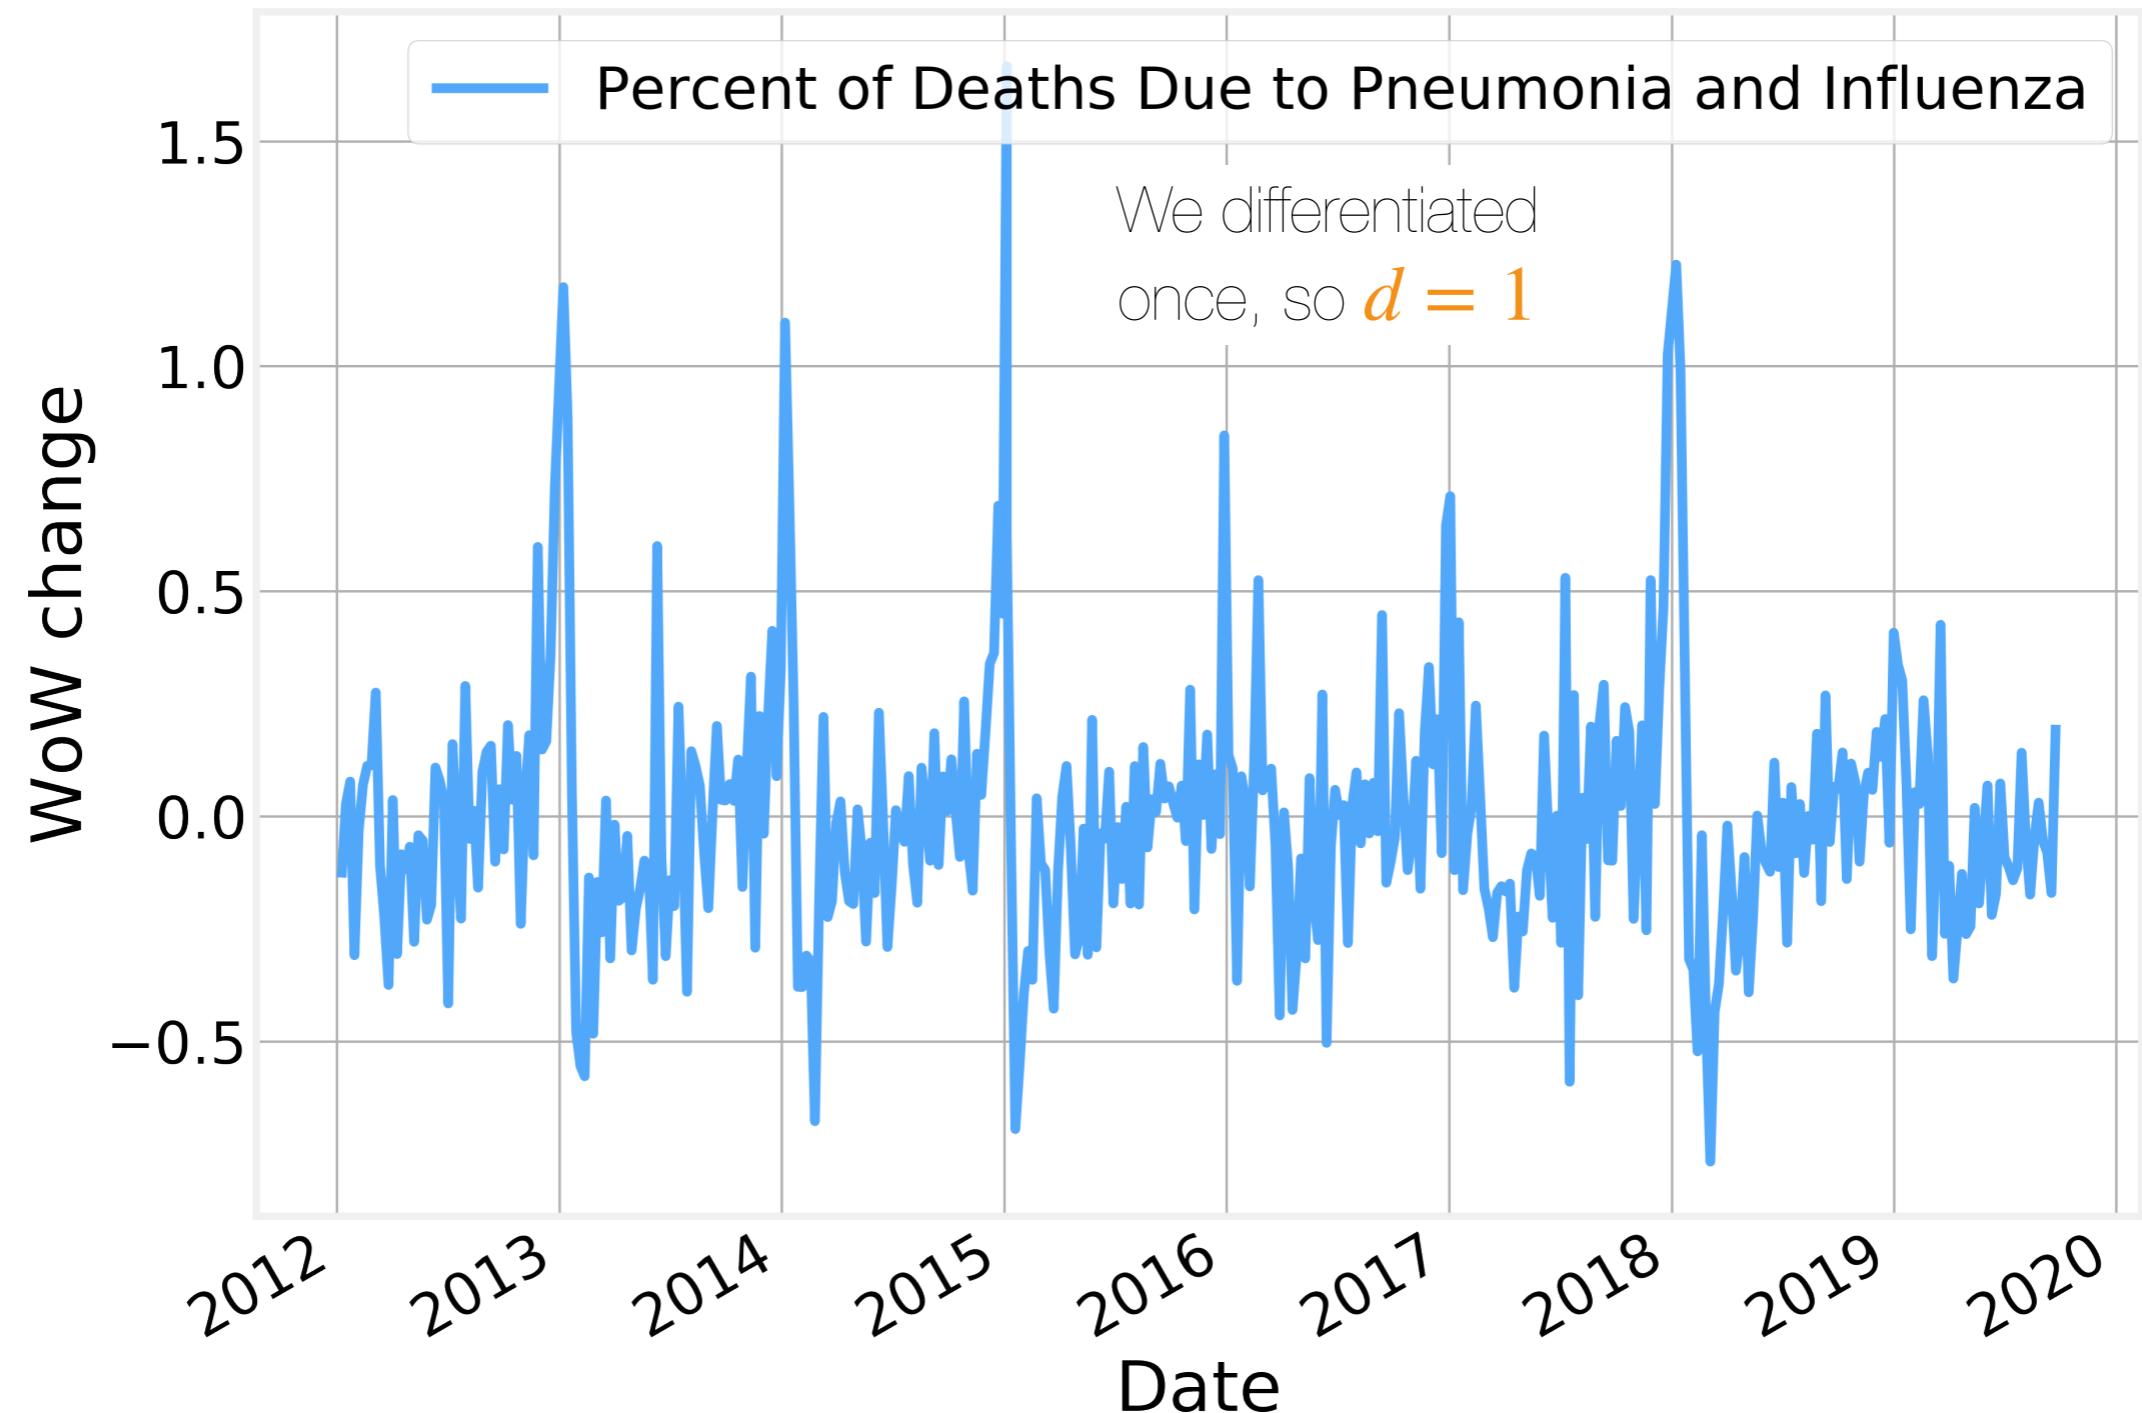


- $d$  - degree of differencing
- $p$  - number of lag observations included in the model (**PACF**)
- $q$  - size of the moving average window (**ACF**)

# Visualize Time Series

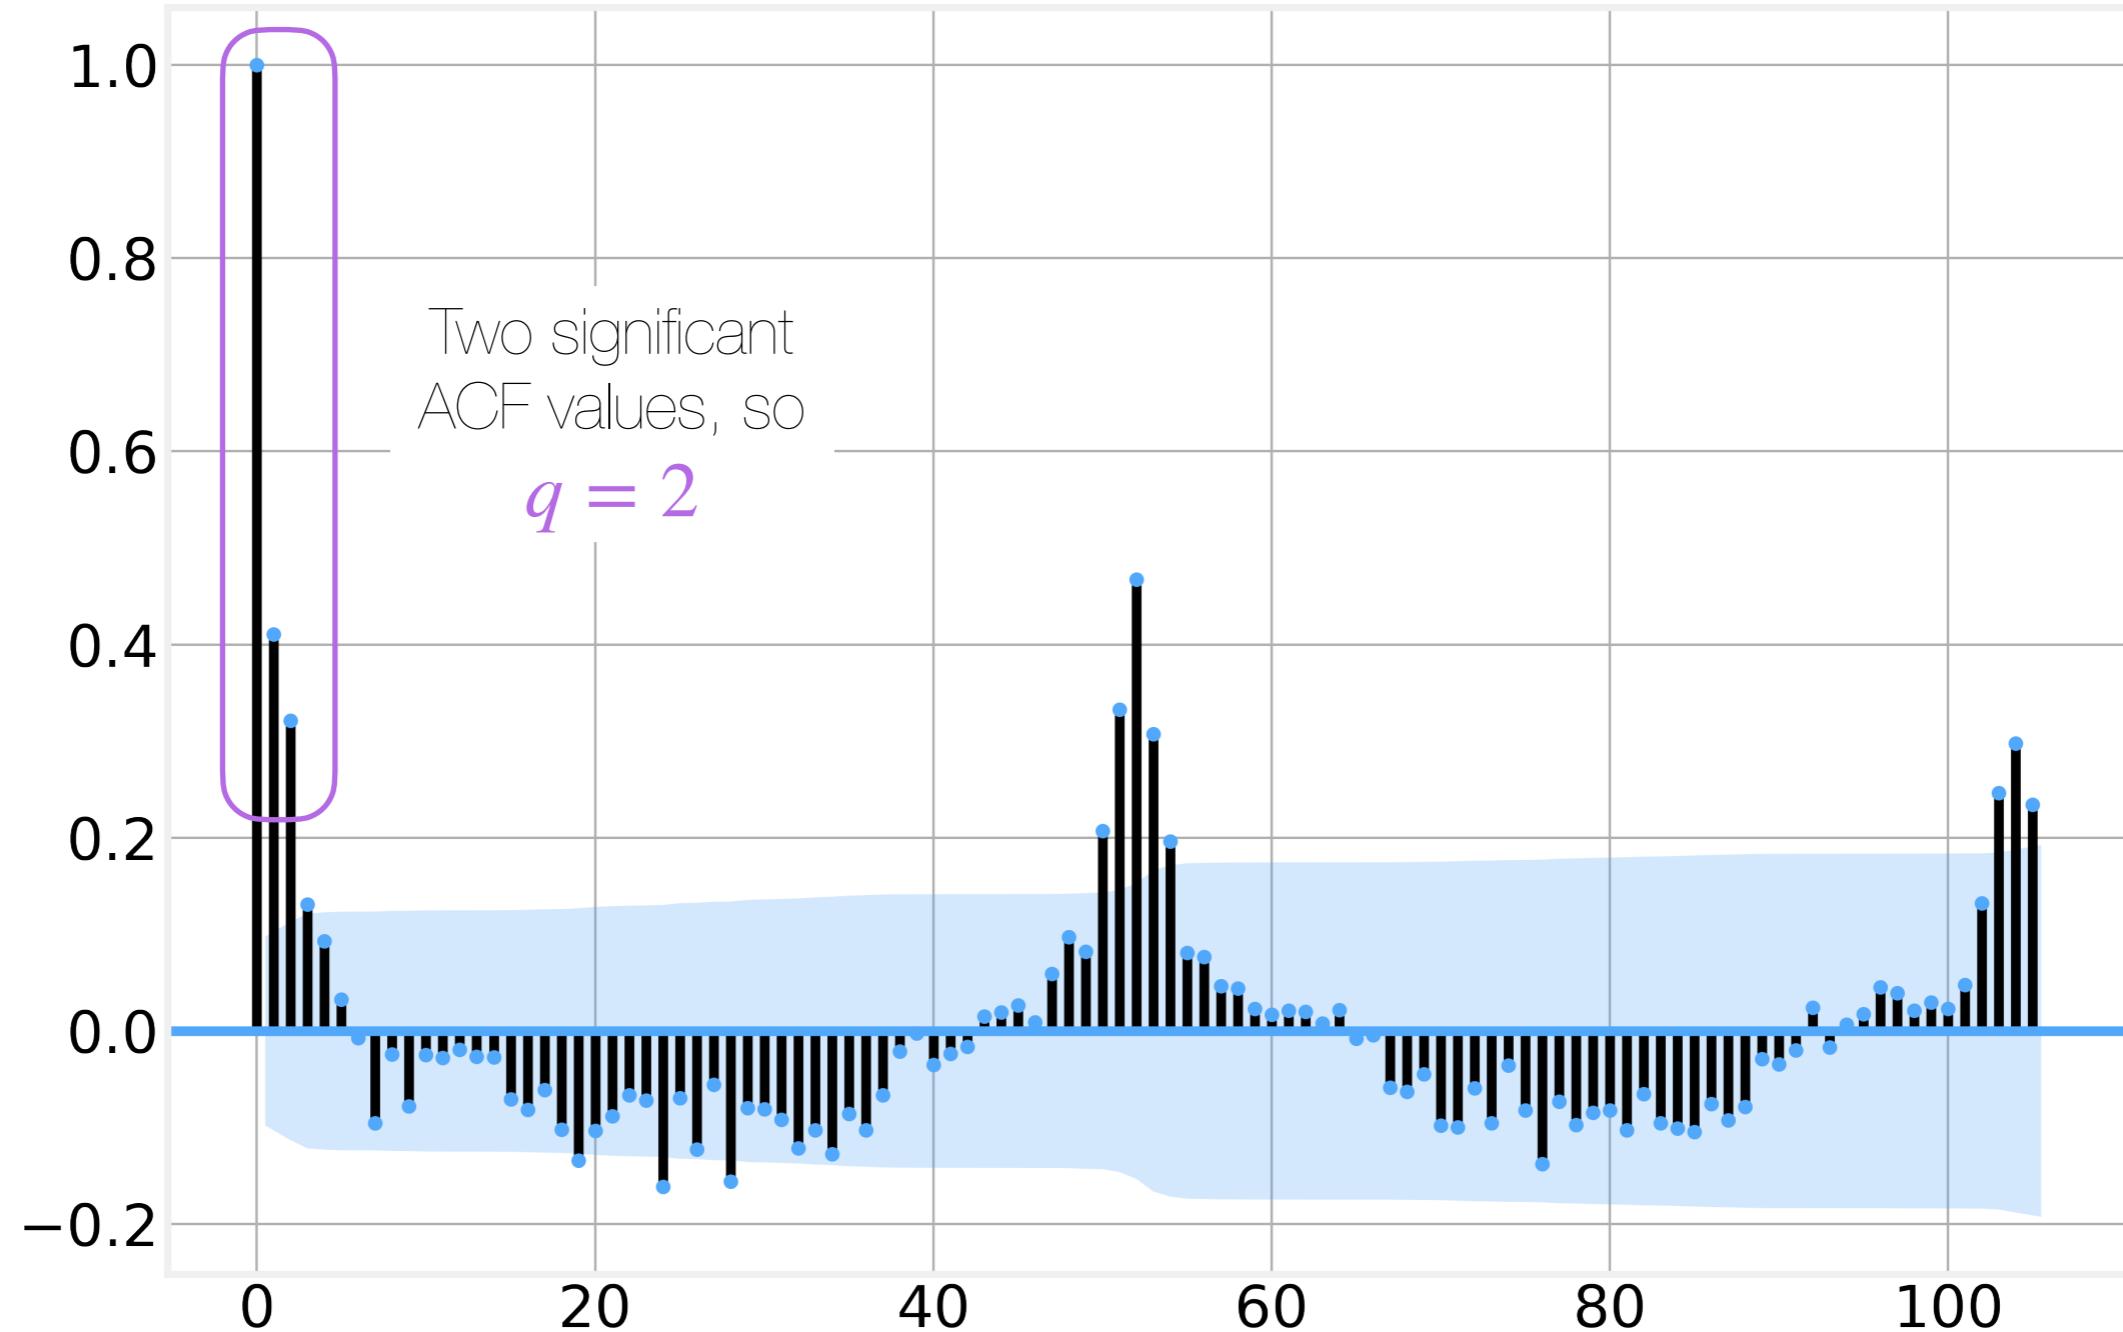


## Stationarize Time Series



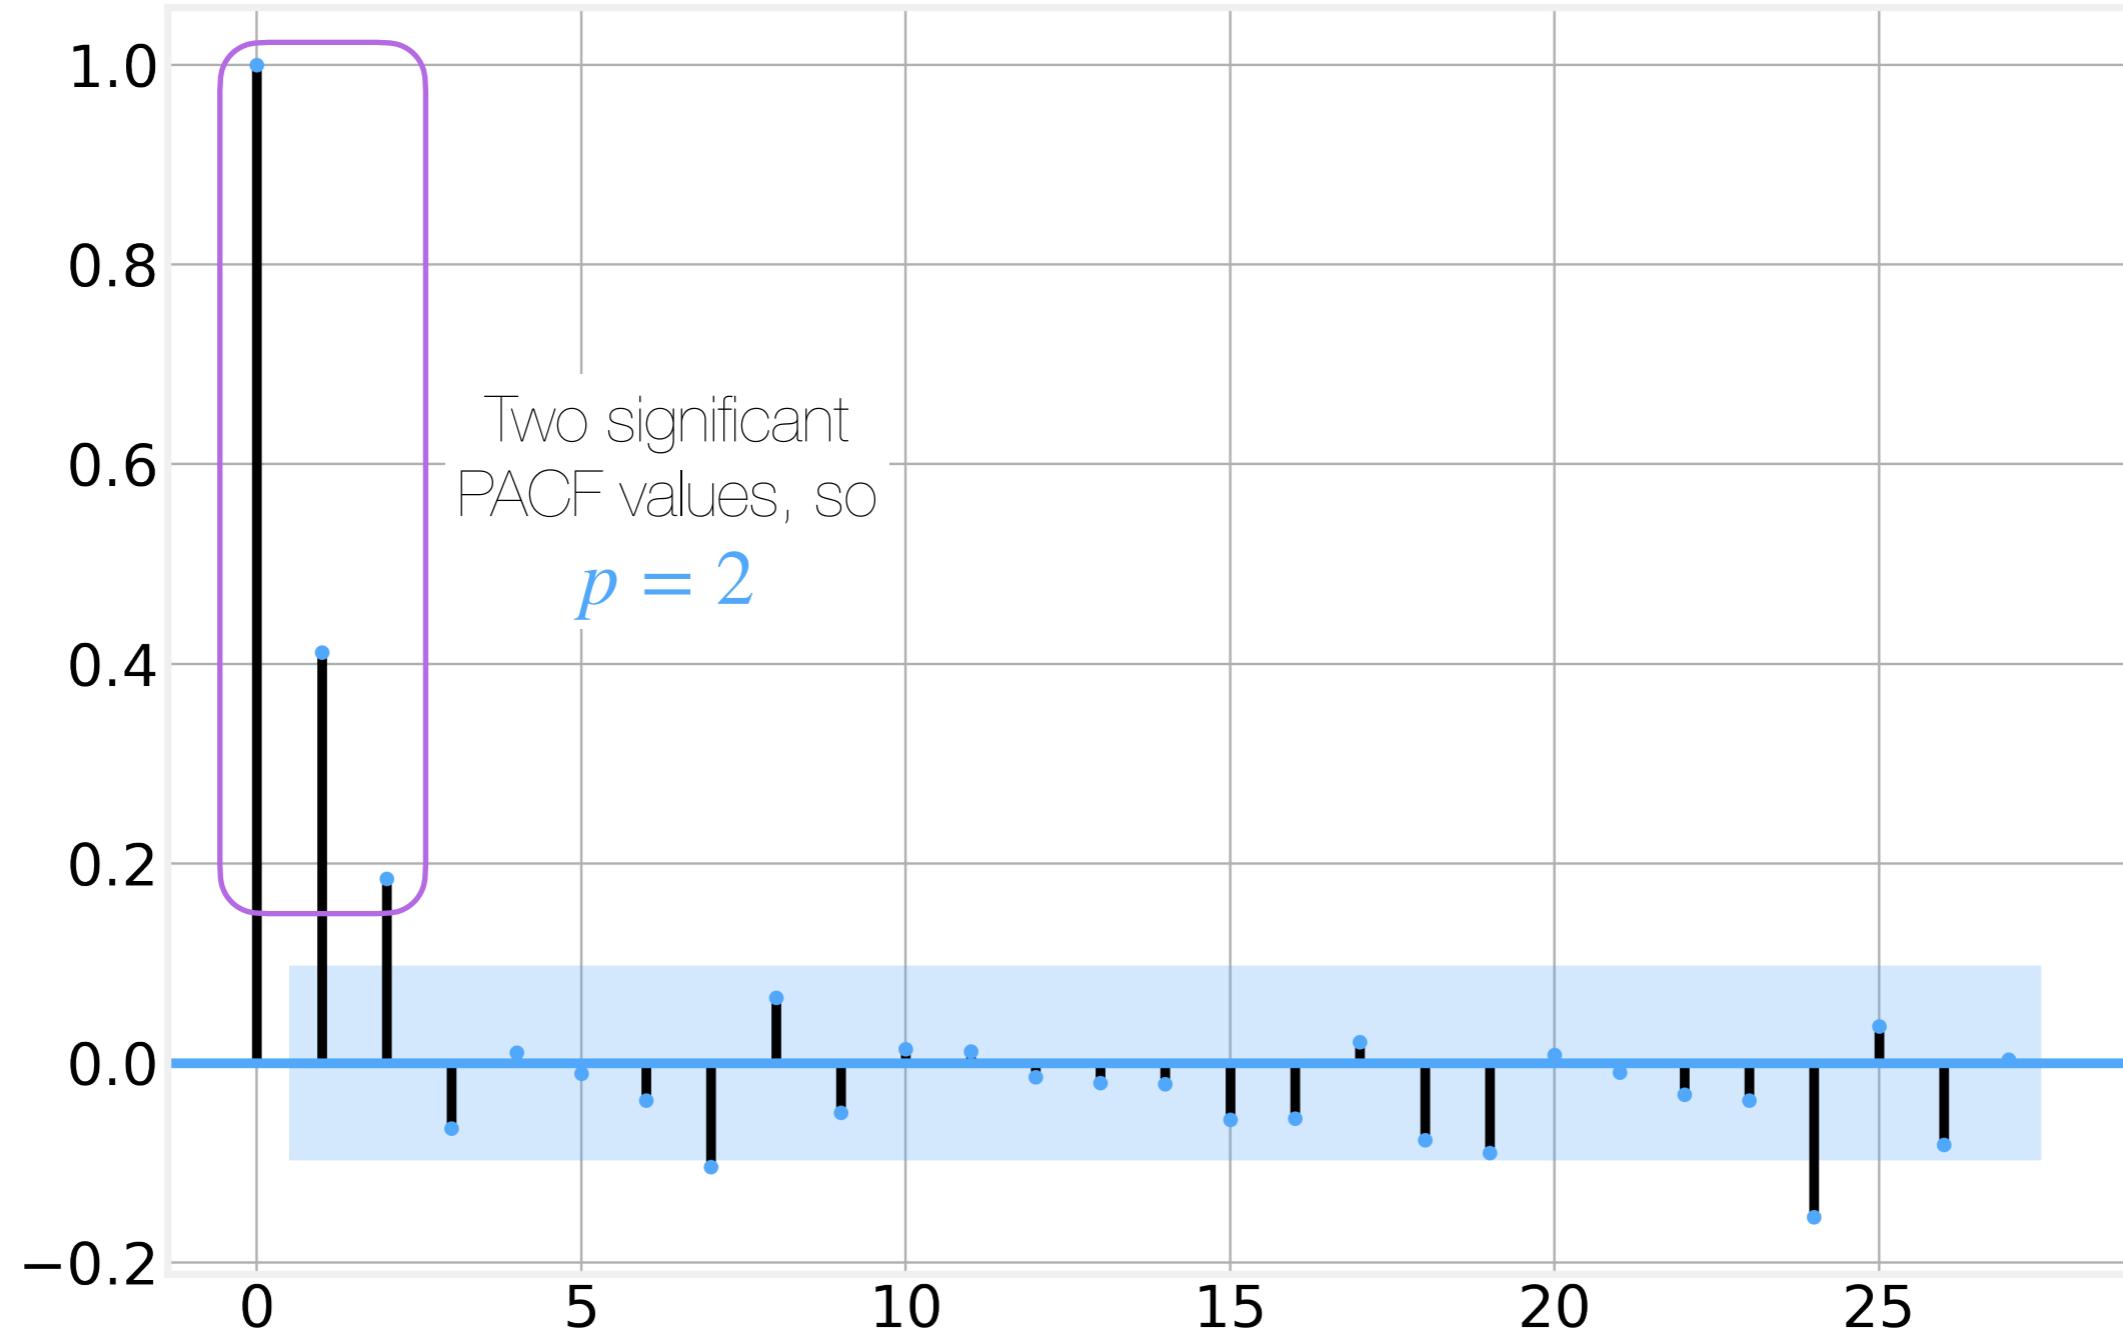
## Plot ACF

### Autocorrelation



## Plot PACF

### Partial Autocorrelation



# Build ARIMA Model

- Using **statsmodels** we just have to do:

```
model = sm.tsa.ARIMA(ILI, (2, 1, 2))
results = model.fit()
```

# Build ARIMA Model

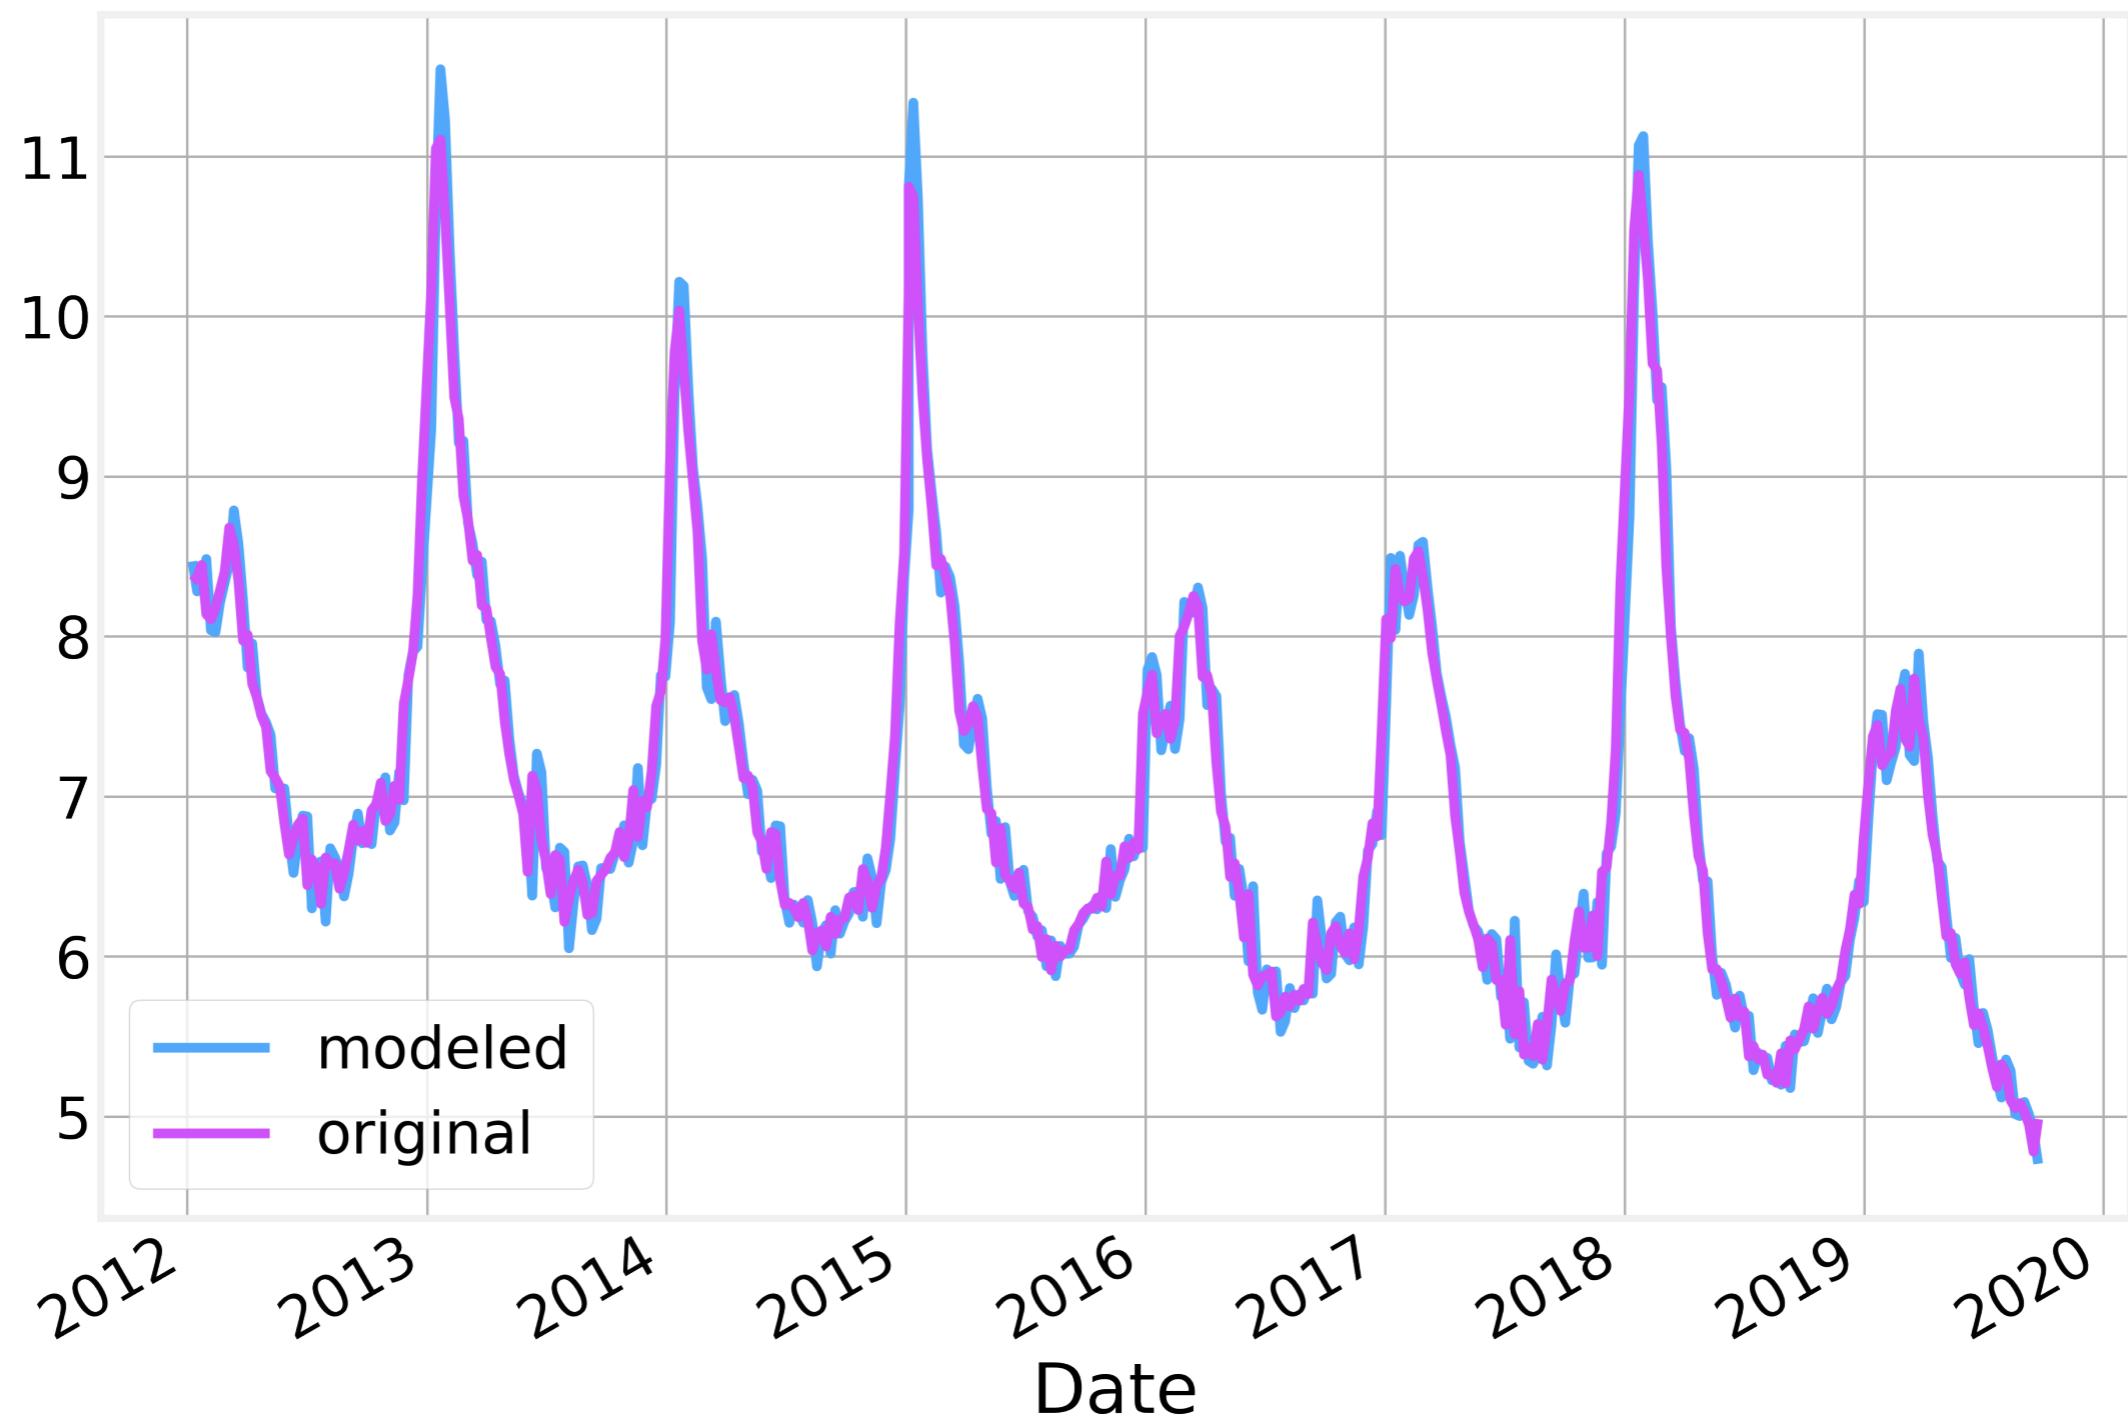
- Using **statsmodels** we just have to do:

```
model = sm.tsa.ARIMA(ILI, (2, 1, 2))
results = model.fit()
```

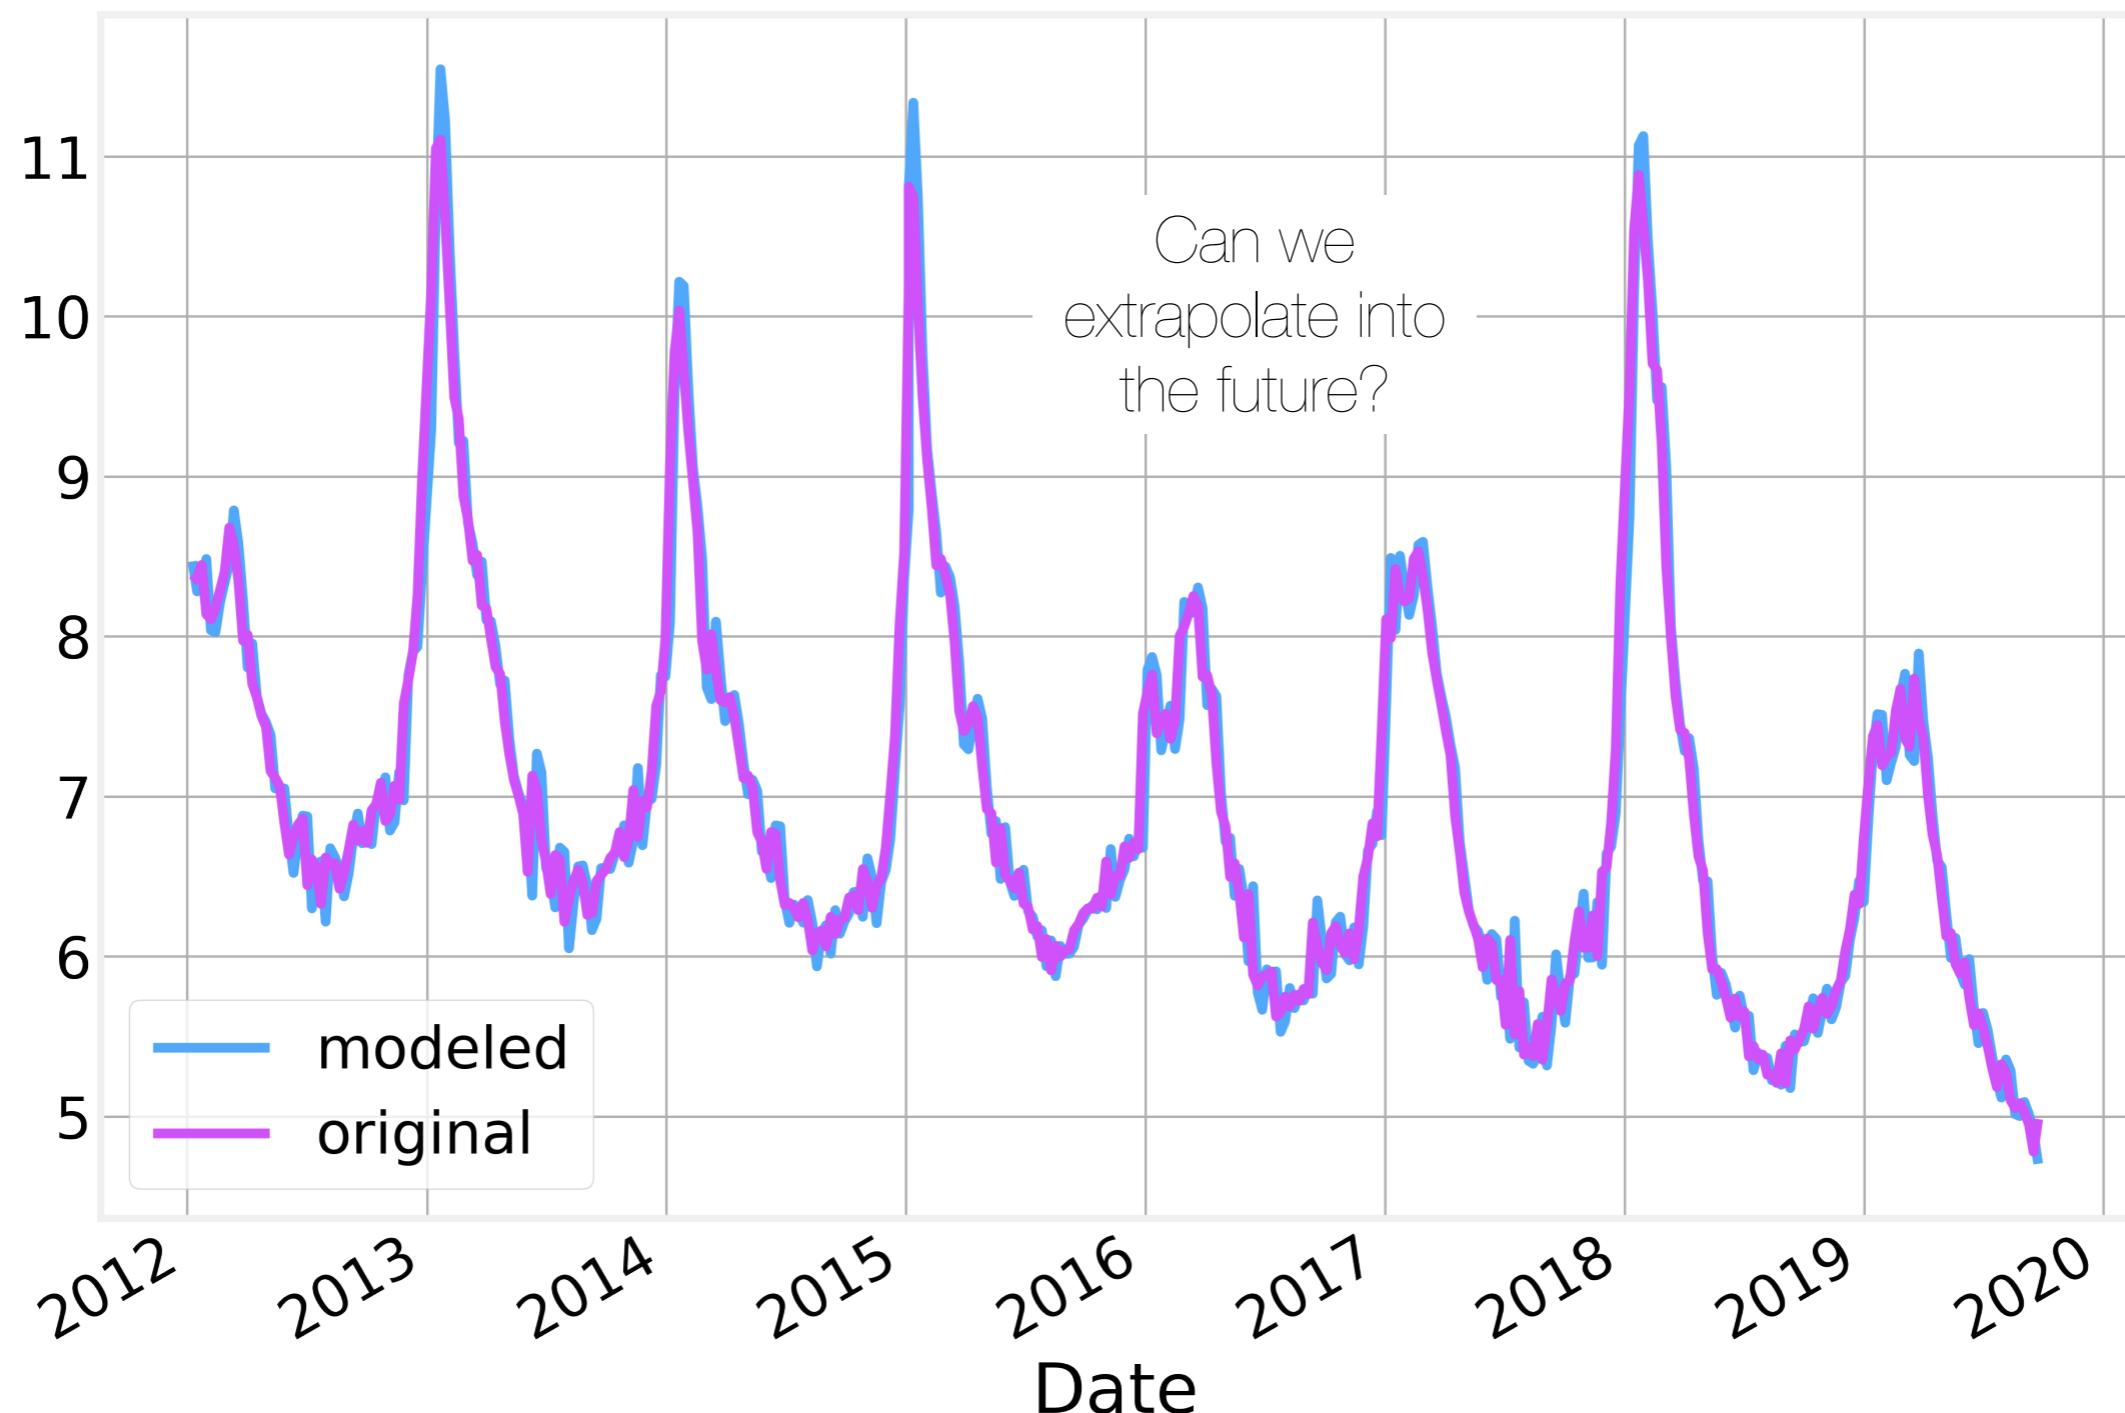
- From where we can easily extract the parameter values using **summary()**:
- and, finally, using **plot\_predict()** we obtain

ARIMA Model Results							
Dep. Variable:	D.Percent of Deaths Due to Pneumonia and Influenza	No. Observations:	402				
Model:	ARIMA(2, 1, 2)	Log Likelihood	-18.777				
Method:	css-mle	S.D. of innovations	0.253				
Date:	Sun, 09 Aug 2020	AIC	49.553				
Time:	20:15:20	BIC	73.532				
Sample:	1	HQIC	59.047				
Coefficients							
	const	coef	std err	z	P> z	[0.025	0.975]
ar.L1.D.Percent of Deaths Due to Pneumonia and Influenza	0.0506	0.561	0.090	0.928	-1.050	1.151	
ar.L2.D.Percent of Deaths Due to Pneumonia and Influenza	0.2785	0.476	0.585	0.559	-0.655	1.212	
ma.L1.D.Percent of Deaths Due to Pneumonia and Influenza	0.2971	0.564	0.527	0.598	-0.808	1.403	
ma.L2.D.Percent of Deaths Due to Pneumonia and Influenza	0.0261	0.318	0.082	0.935	-0.597	0.649	
Roots							
	Real	Imaginary	Modulus	Frequency			
AR.1	1.8061	+0.0000j	1.8061	0.0000			
AR.2	-1.9878	+0.0000j	1.9878	0.5000			
MA.1	-5.7000	-2.4260j	6.1948	-0.4360			
MA.2	-5.7000	+2.4260j	6.1948	0.4360			

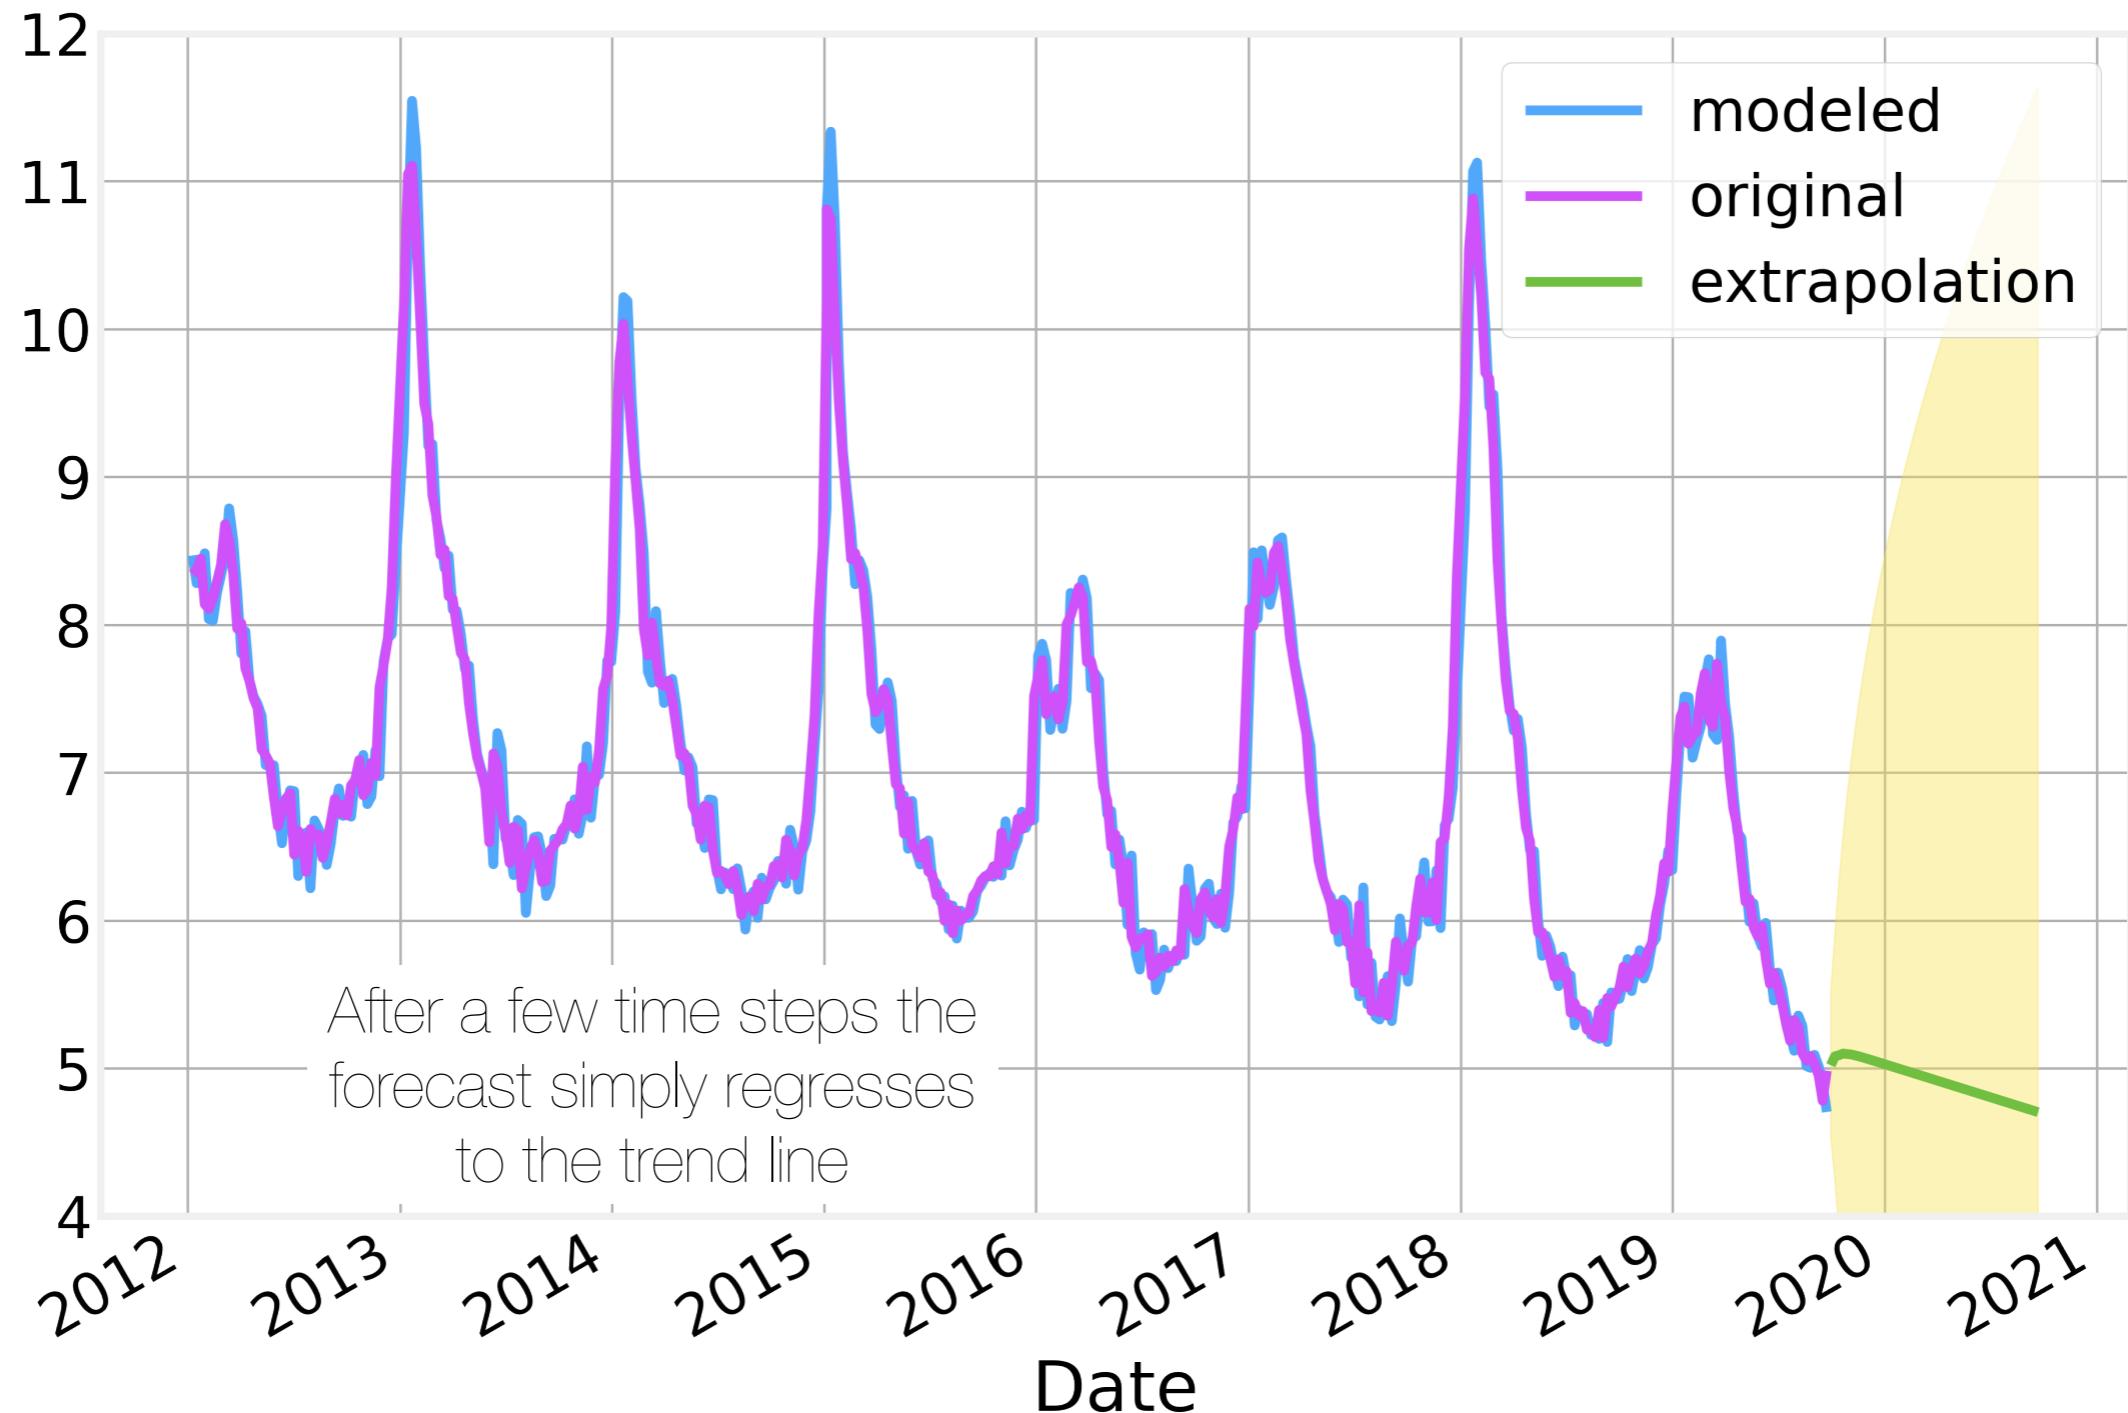
# Forecast



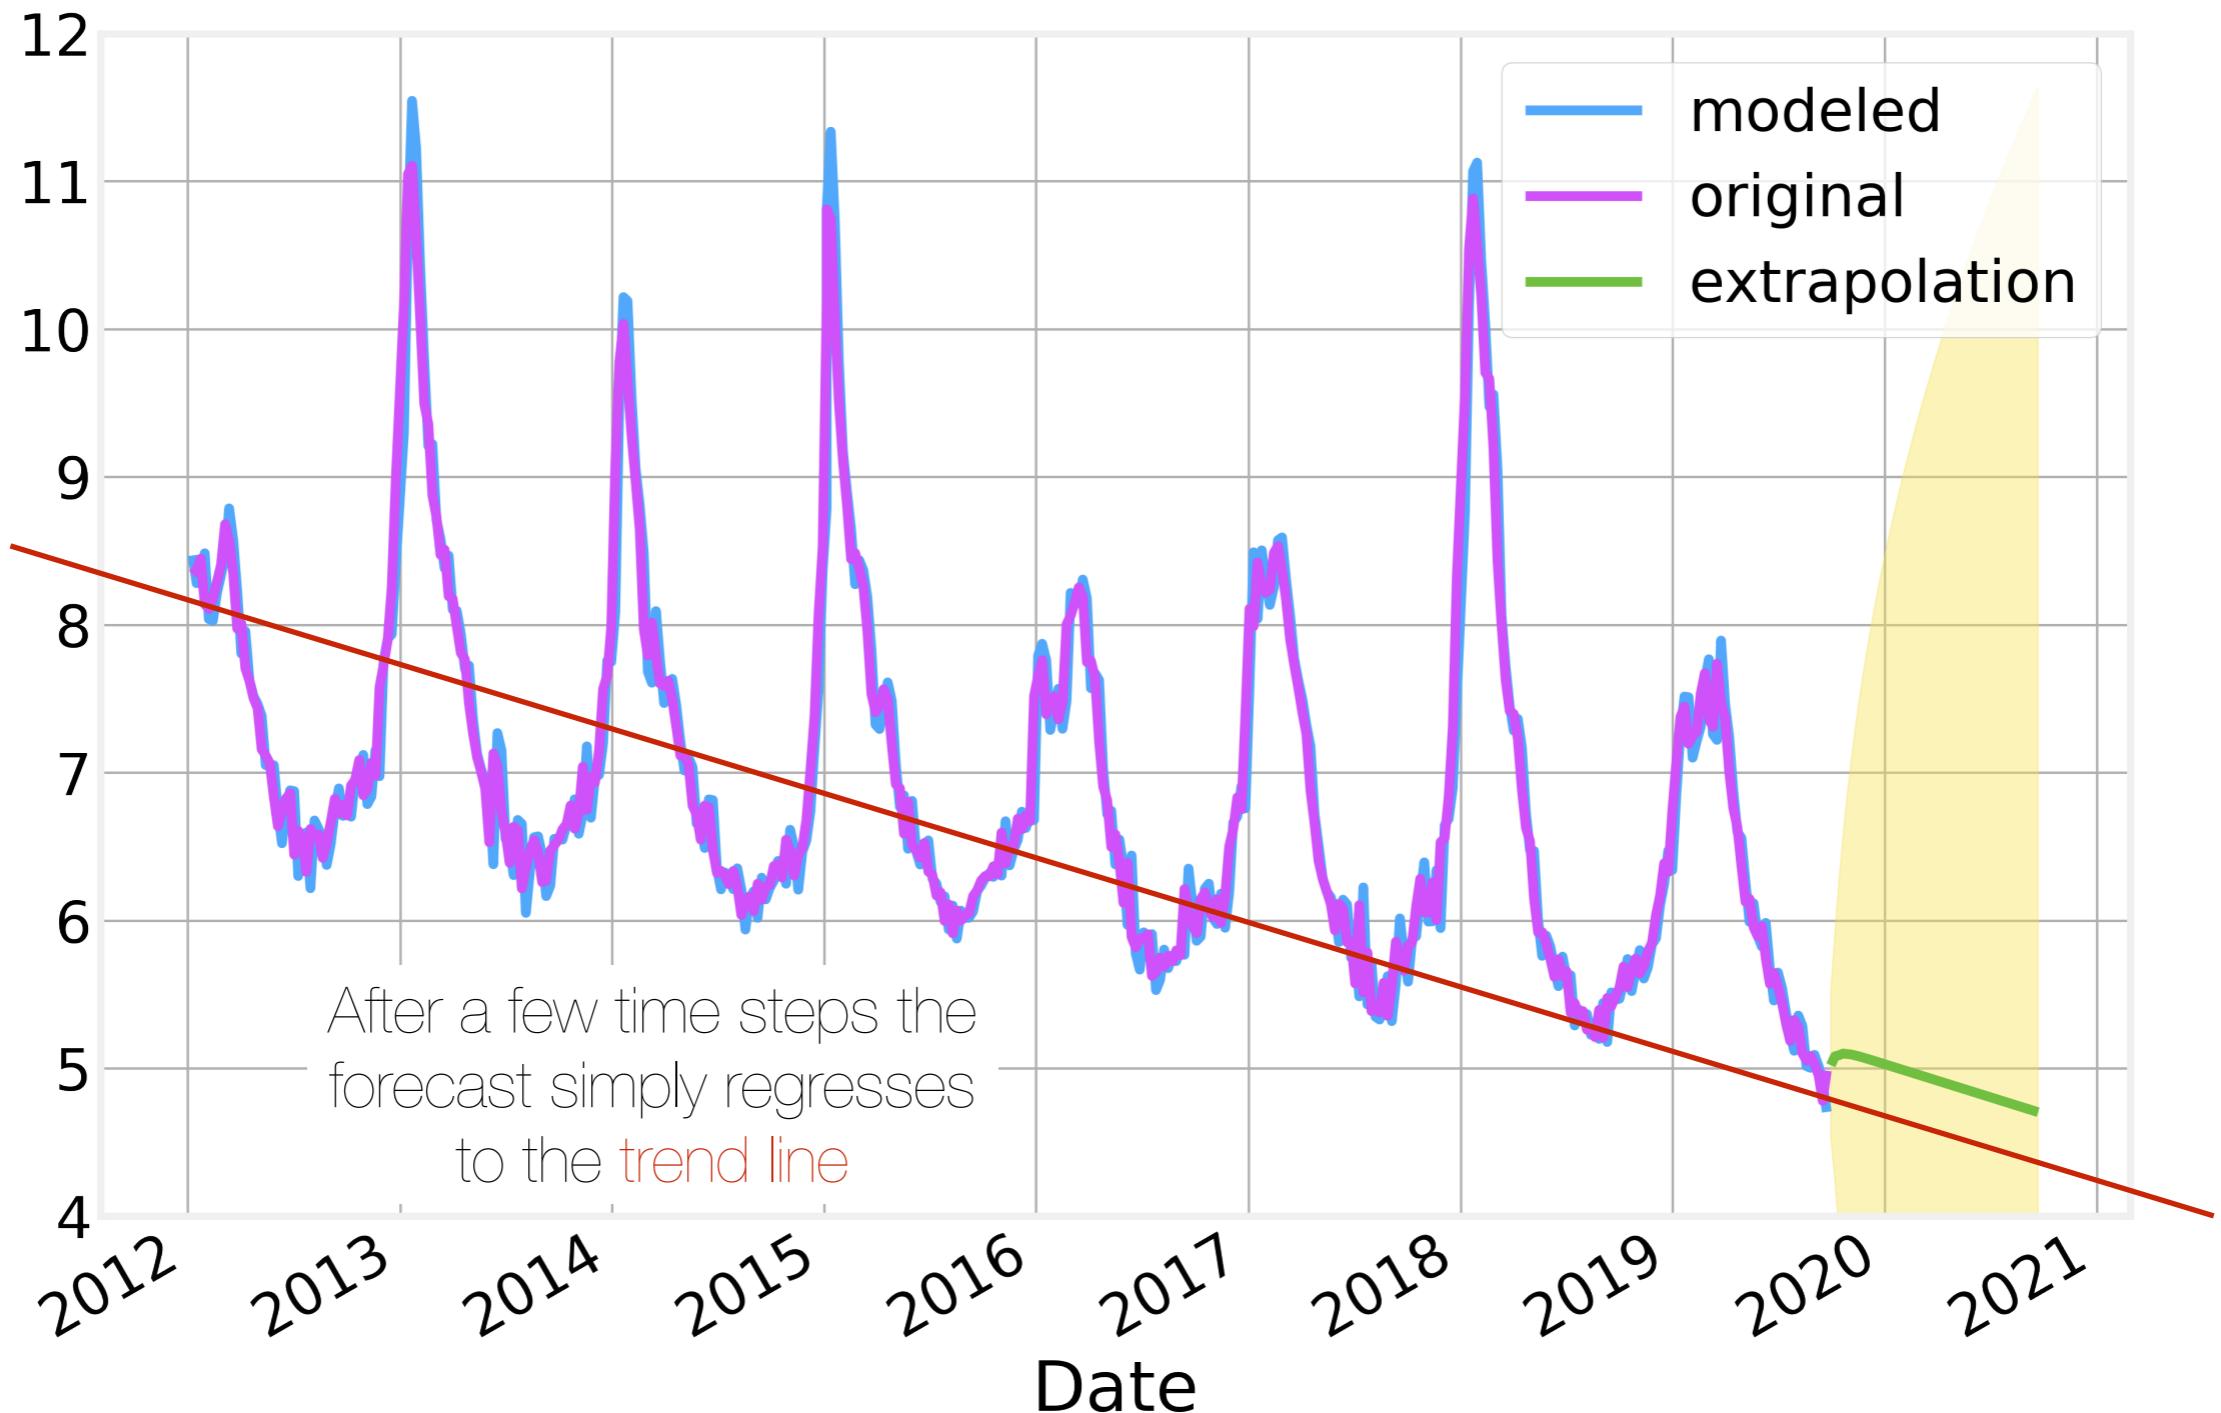
# Forecast



# Forecast



# Forecast



# Seasonal Models

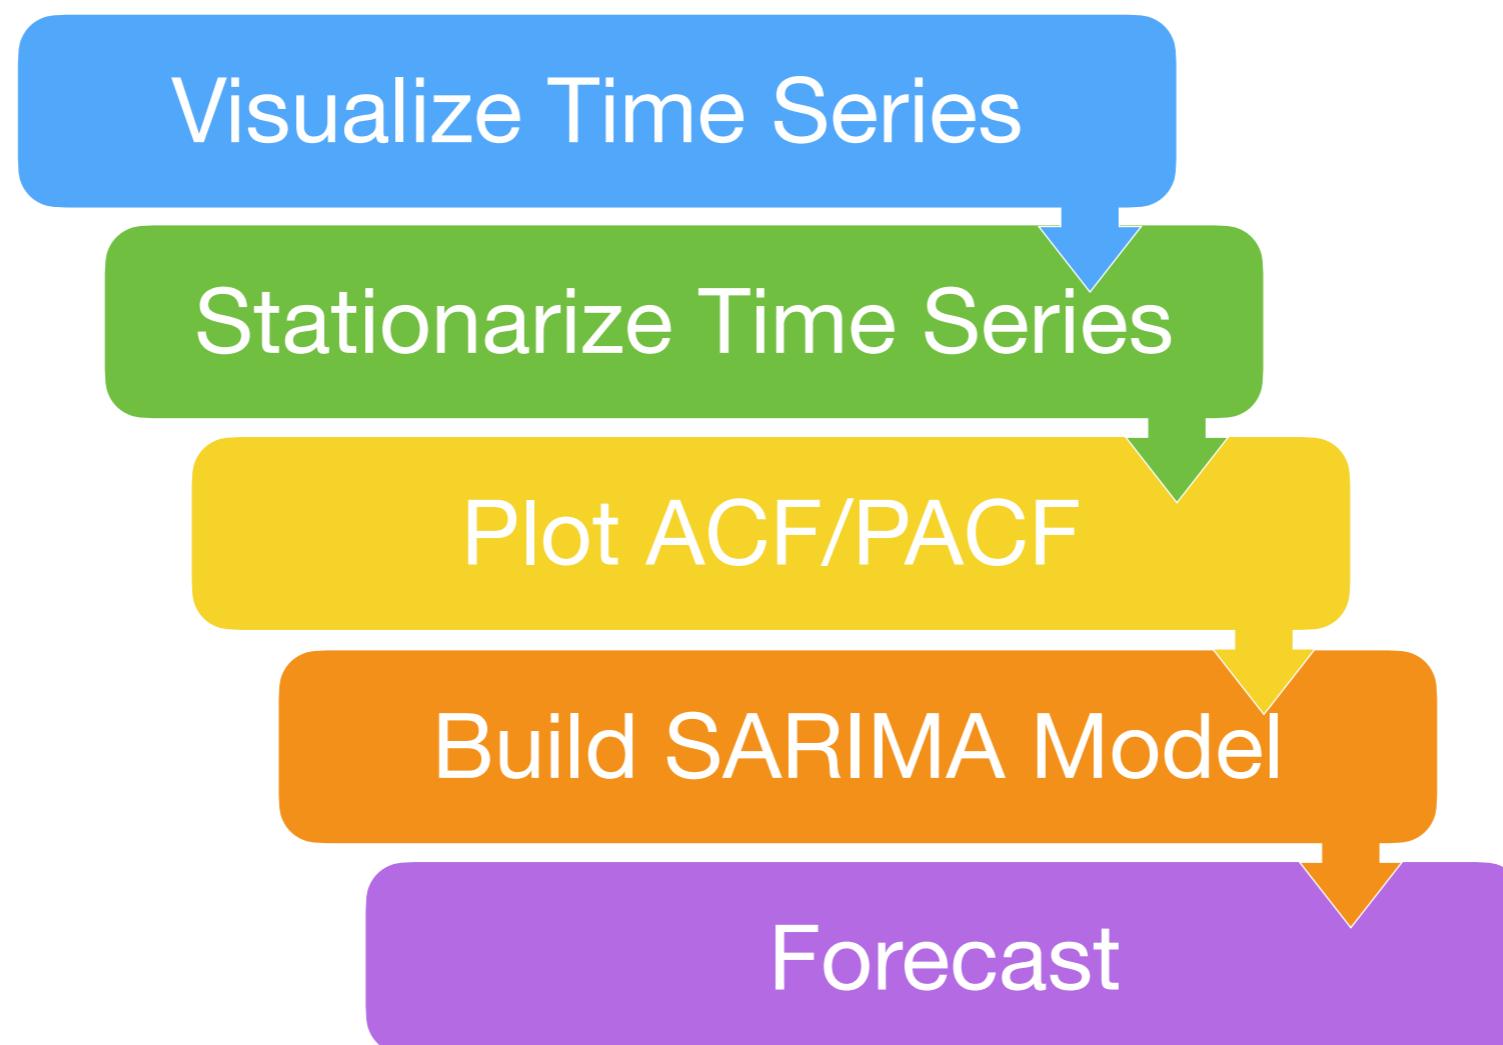
---

- The problem with our previous approach is that our model isn't taking into account the full seasonal structure of the data.
- A simple approach to taking seasonal variations into account is to train the **ARIMA** model just on the residuals of the time series, after properly decomposing it.
  - We can recover the original signal by "manually" adding the **trend** and **seasonality** we obtained in our decomposition
- A more sophisticated approach is to use the **SARIMA** (Seasonal **ARIMA**) class of models.
- **SARIMA** explicitly models the seasonal variations as a full fledged **ARIMA** model at the cost of 4 extra parameters:  $\text{SARIMA } (p, d, q) \times (P, D, Q, s)$ :
  - $p$ ,  $d$  and  $q$ , are defined exactly as above
  - $s$ , is the seasonal periodicity
  - $P$ ,  $D$  and  $Q$ , are the parameters of an **ARIMA** model trained on the  $s$ -lagged time series

# SARIMA

<https://www.statsmodels.org/devel/generated/statsmodels.tsa.statespace.sarimax.SARIMAX.html>

- To fit a SARIMA model, we follow a similar procedure to the one described above:

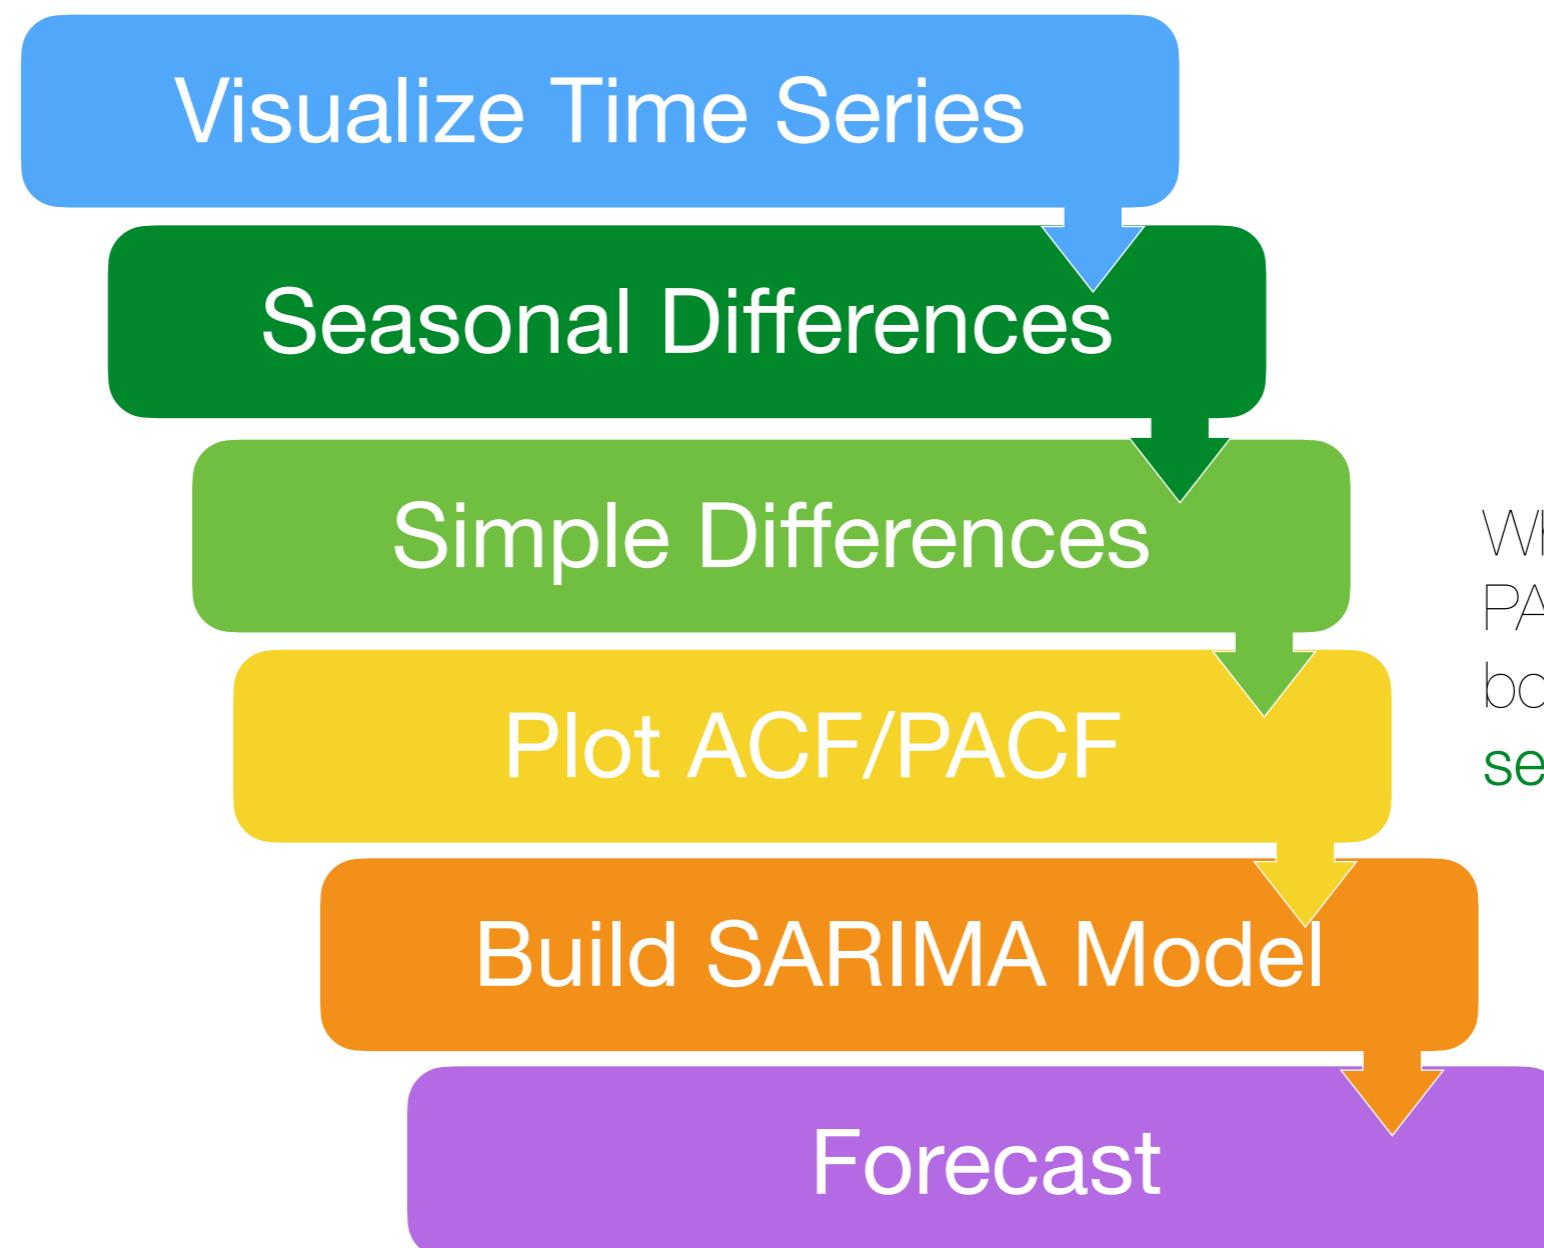


- With one important difference, in the Stationarize step, we differentiate on simple and seasonal lags

# SARIMA

<https://www.statsmodels.org/devel/generated/statsmodels.tsa.statespace.sarimax.SARIMAX.html>

- Our procedure is then:

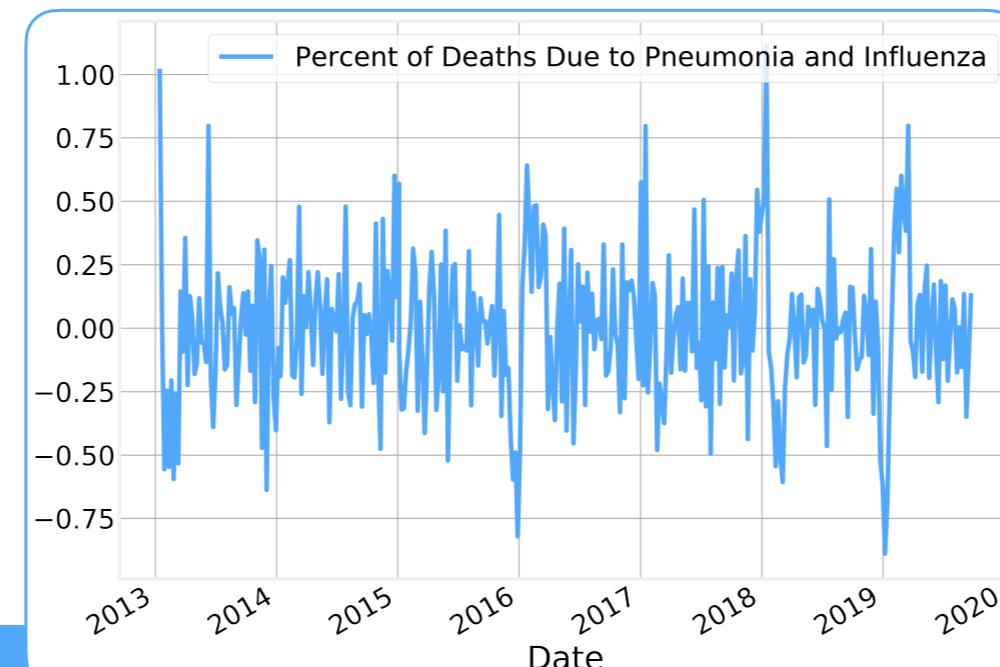


When plotting the ACD/  
PACF we must look for  
both **simple** and  
**seasonal** lags

# SARIMA

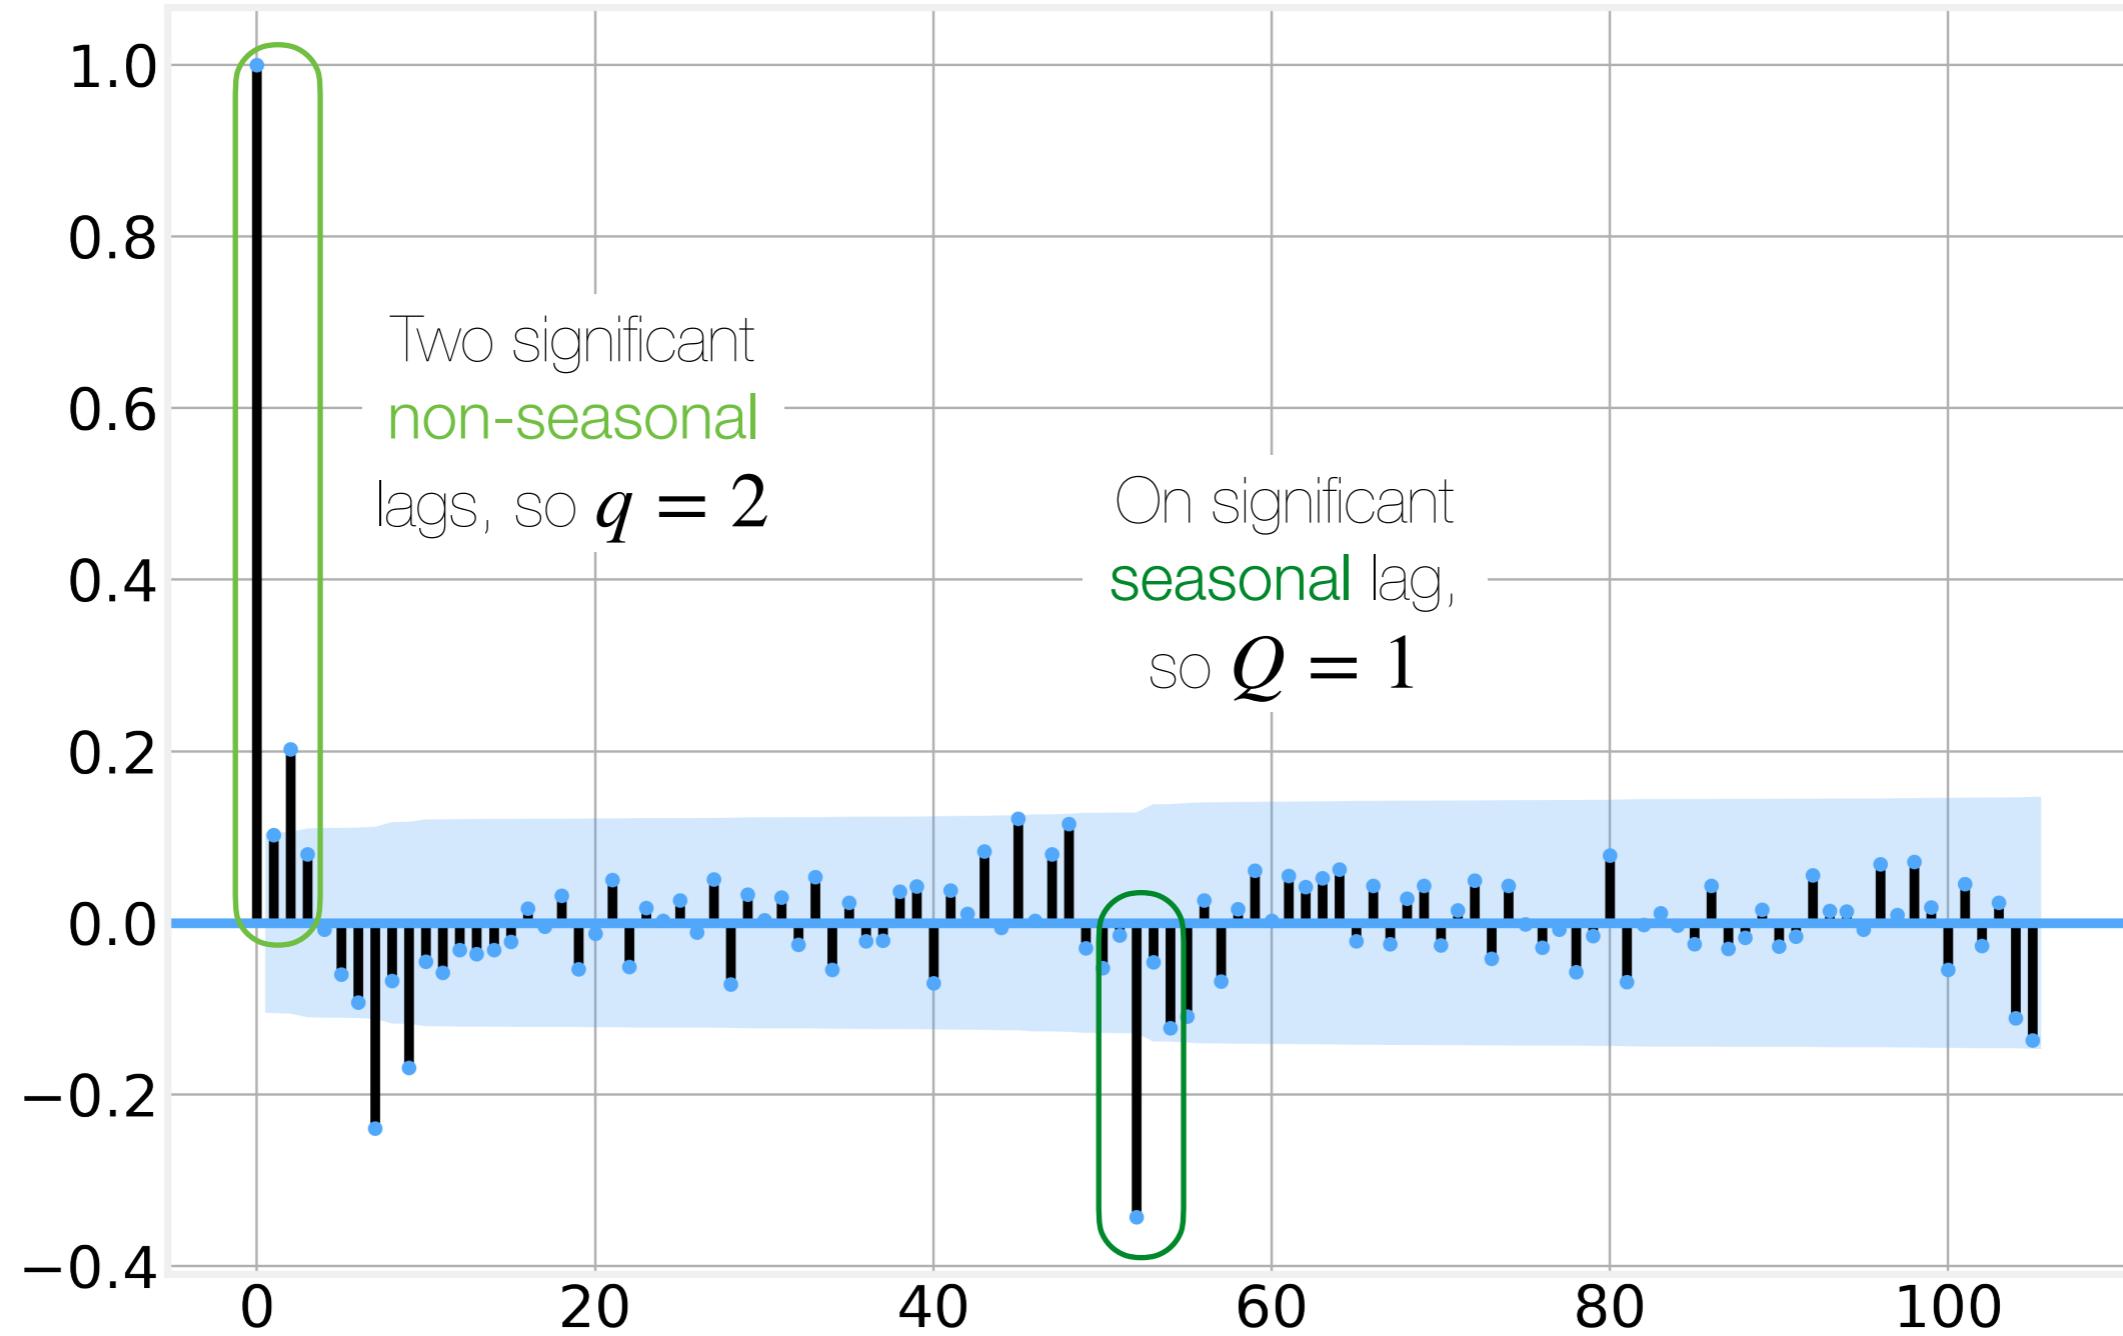
<https://www.statsmodels.org/devel/generated/statsmodels.tsa.statespace.sarimax.SARIMAX.html>

- **statsmodels** provides the **SARIMAX** model in the **tsa.statsmodels** package.
- The API is similar to that of the **ARIMA** class of models, but now we must provide both:
  - **order** -  $(p, d, q)$  - the usual **ARIMA** parameters
  - **seasonal\_order** -  $(P, D, Q, s)$  - the parameters specifying the seasonal component
- With this approach, we are now able to perform significantly better forecasts!
- The first step is to take the **seasonal** (52) followed but the **non-seasonal** (1) differences.



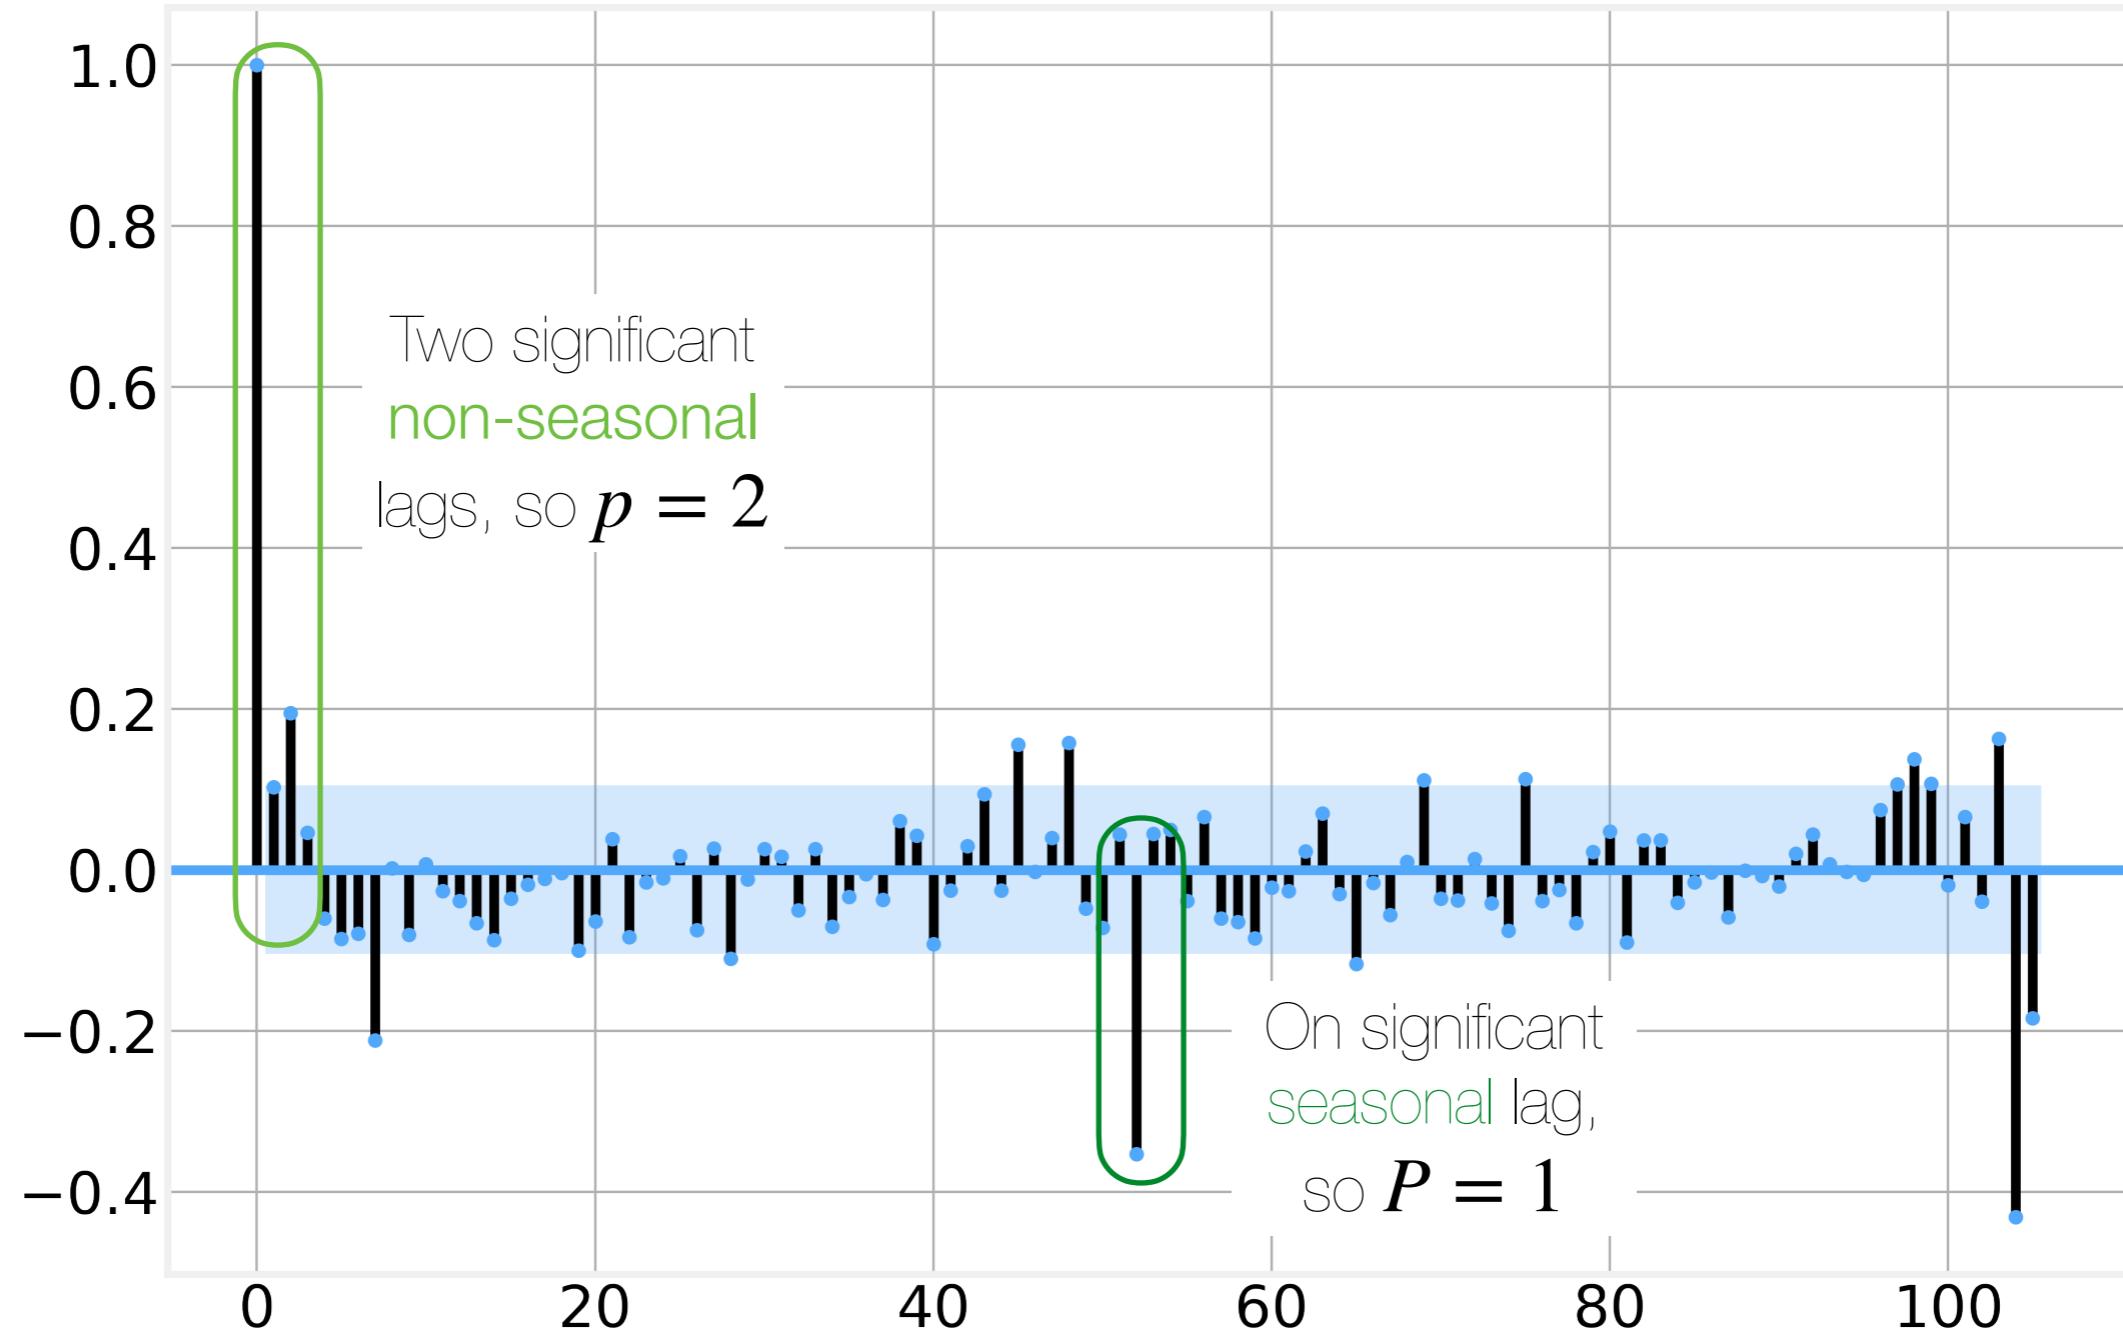
## Plot ACF

### Autocorrelation

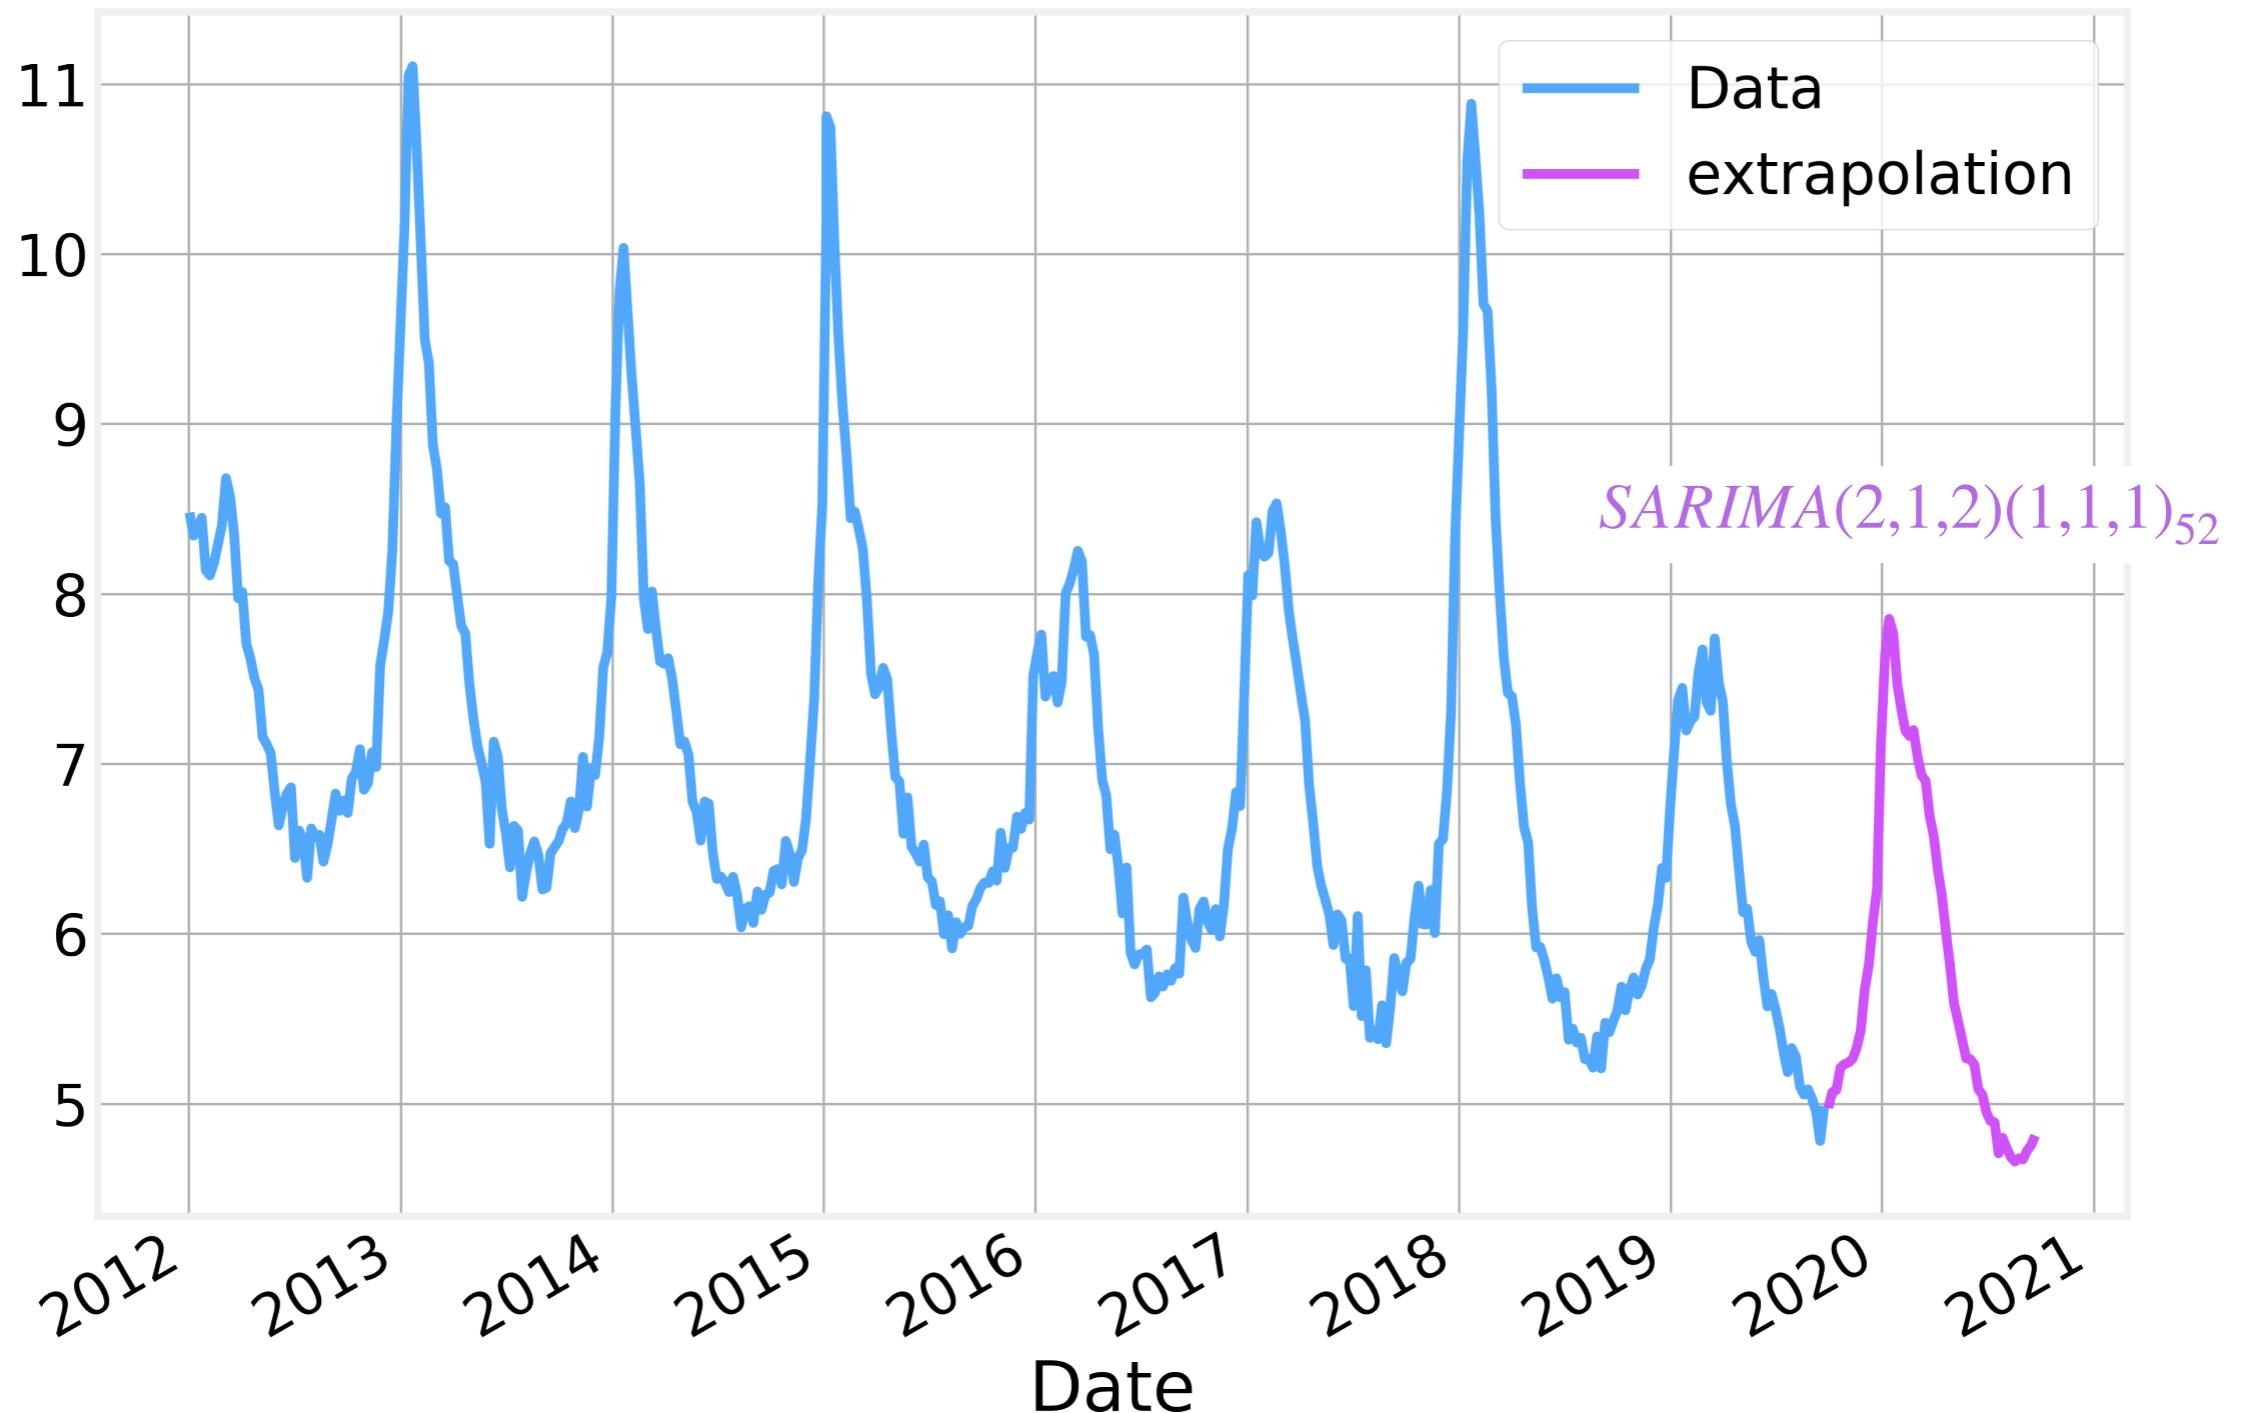


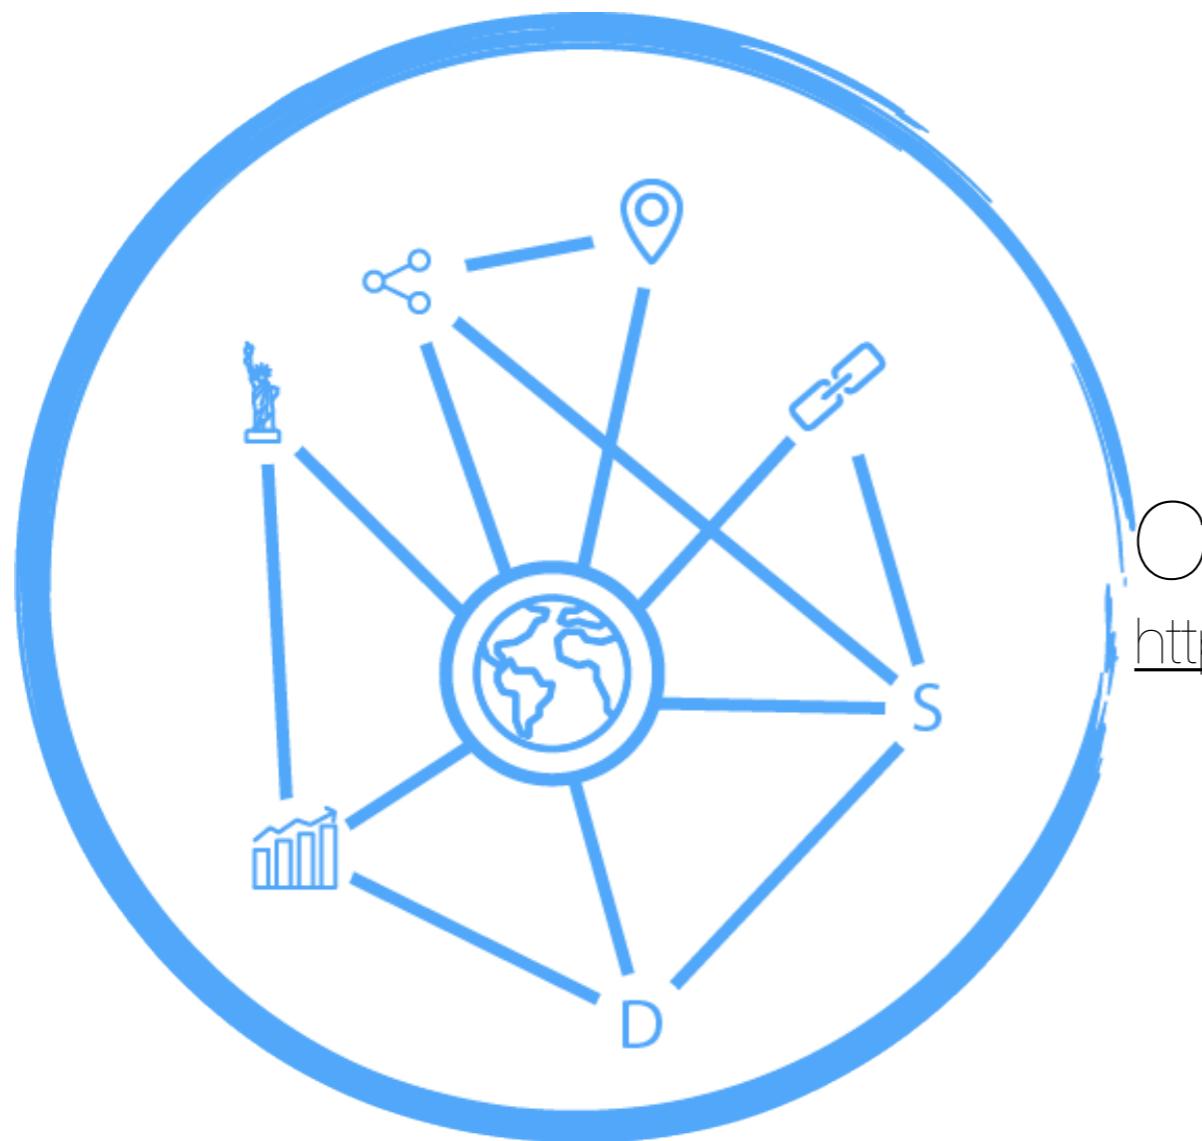
## Plot PACF

### Partial Autocorrelation



# Forecast





Code - ARIMA Models

<https://github.com/DataForScience/AdvancedTimeseries>

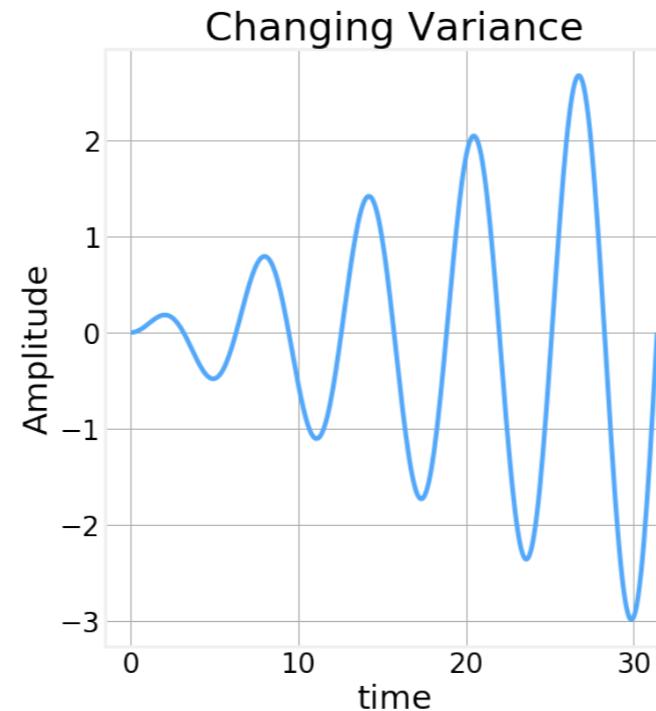


## Lesson IV: ARCH Models

# Heteroscedasticity

<https://en.wikipedia.org/wiki/Heteroscedasticity>

- Some time series display **variance that changes in time**:



The temporal behavior of **Variance** is not required to be monotonous.

- This phenomenon is called **Heteroscedasticity**
- This kind of time series required a specific class of models from the ones we have considered so far

# Heteroscedastic Models

- **ARIMA** class models can be written as:

$$x_t = f(x_{t-1}, \dots)$$

- **Heteroskedastic Models**, referred as **ARCH** models, can be interpreted as being a simple **AR** model applied to the variance of a series.

$$\text{Var}(x_t | x_{t-1}) \equiv \sigma_t^2 = f(x_{t-1}, \dots)$$

- In the simplest case of the **Autoregressive Model, AR(1)** we have:

$$x_t = \mu + \sigma_t \epsilon_t$$

- With

$$\sigma_t^2 = \alpha_0 + \alpha_1 x_{t-1}^2$$

- We obtain the order one **Autoregressive Conditionally Heteroscedastic Model, ARCH(1)** model.
- We further impose the constraints  $\alpha_0 > 0$ ,  $\alpha_1 > 0$  and  $\alpha_1^2 < 1/3$  to avoid negative variances.

# ARCH Model

- The **ARCH(1)** Model can be written:

$$x_t = \sigma_t \epsilon_t$$

$$\sigma_t^2 = \alpha_0 + \alpha_1 x_{t-1}^2$$

- Where we assume that the **stochastic variable**  $\epsilon_t$  is Normally distributed:

$$\epsilon_t \sim \mathcal{N}(0,1)$$

- Under these conditions, we have:

$$x_t^2 = \alpha_0 + \alpha_1 x_{t-1}^2 + \text{error}$$

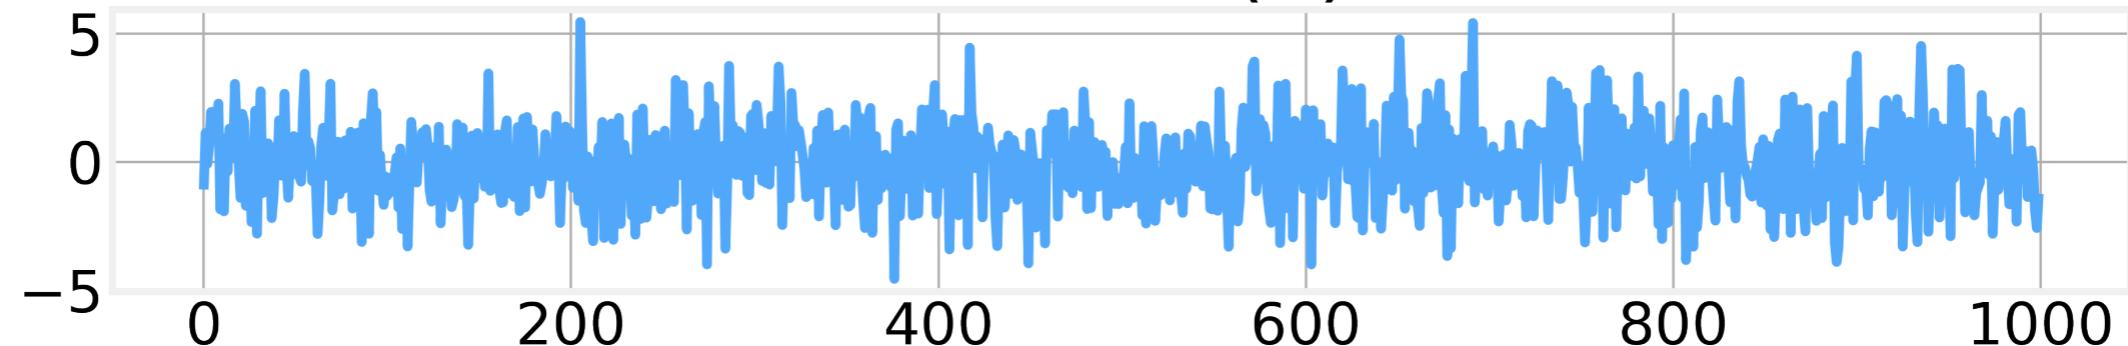
or, in other words, an **AR(1)** model for  $x_t^2$ ! But with a significant error term...

- And  $x_t$  is just White Noise!

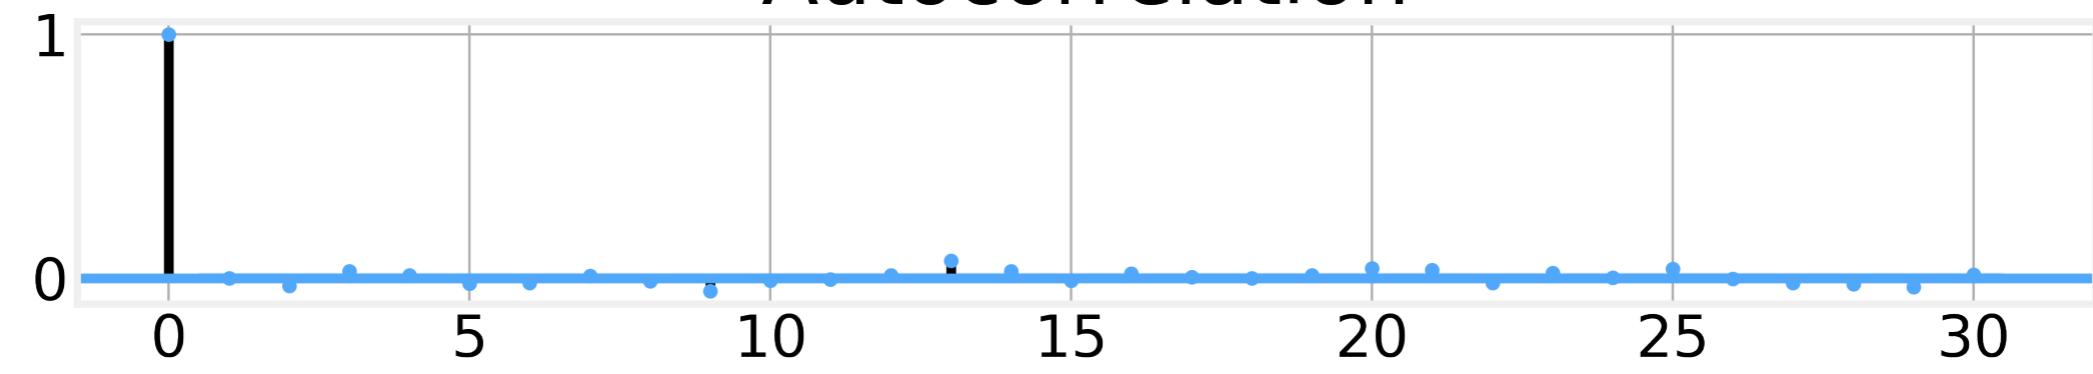
- Let's start by simulating this process in order to understand it better

# ARCH Process

Simulated ARCH(1) Process

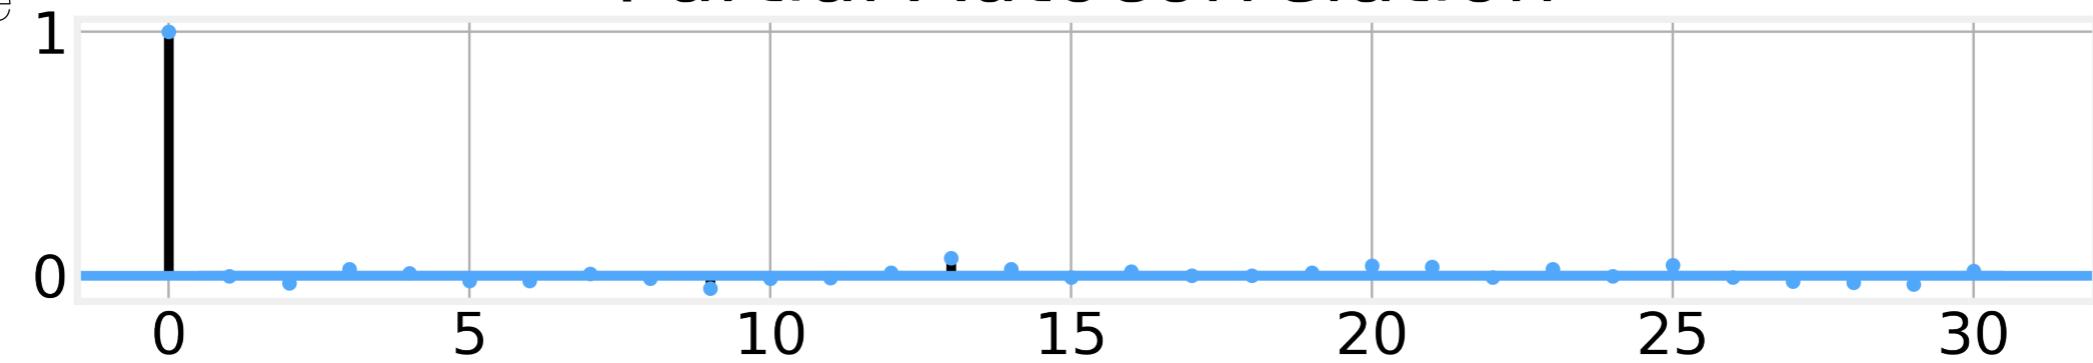


Autocorrelation



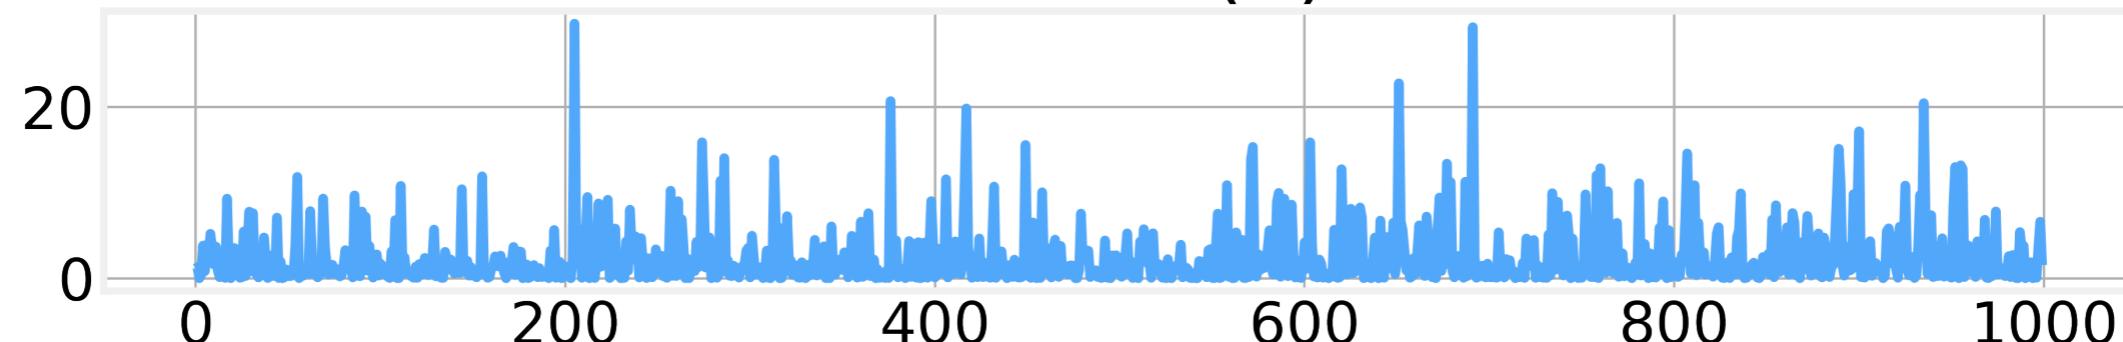
Apparently  
just  
White-Noise

Partial Autocorrelation

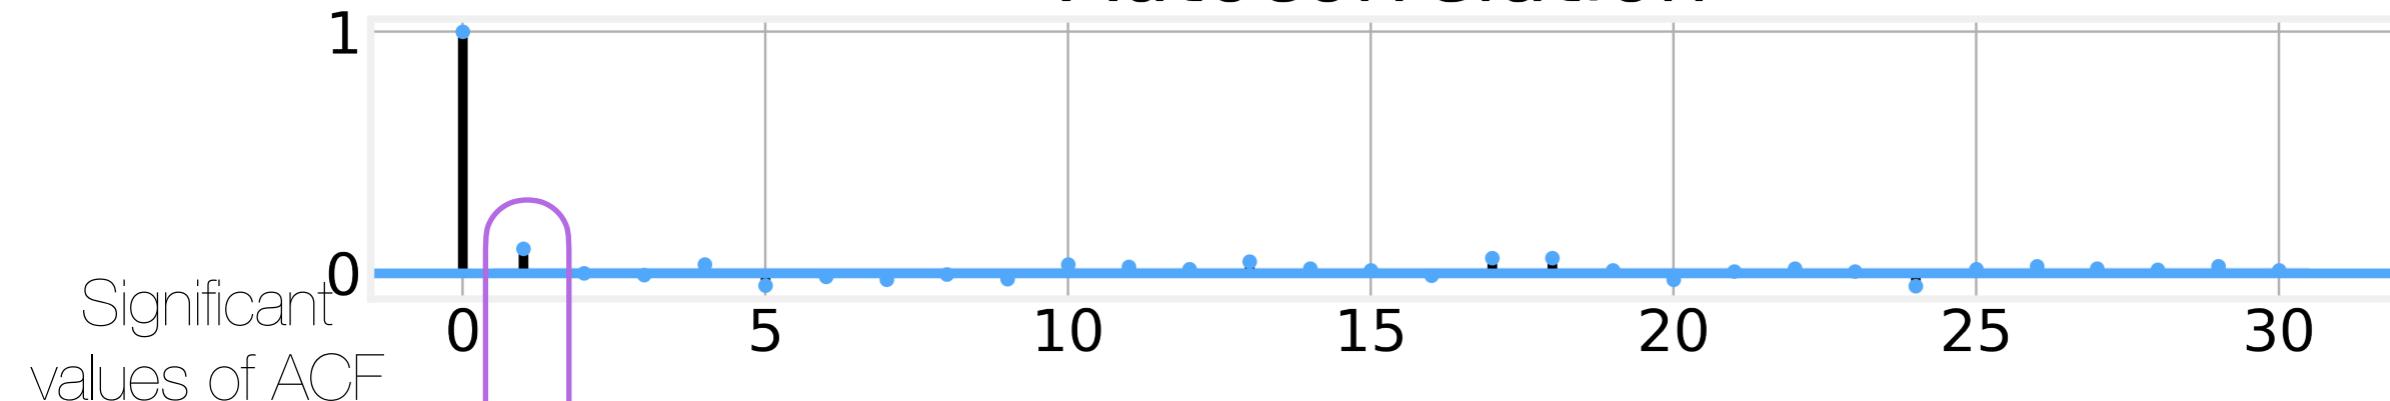


# ARCH Process

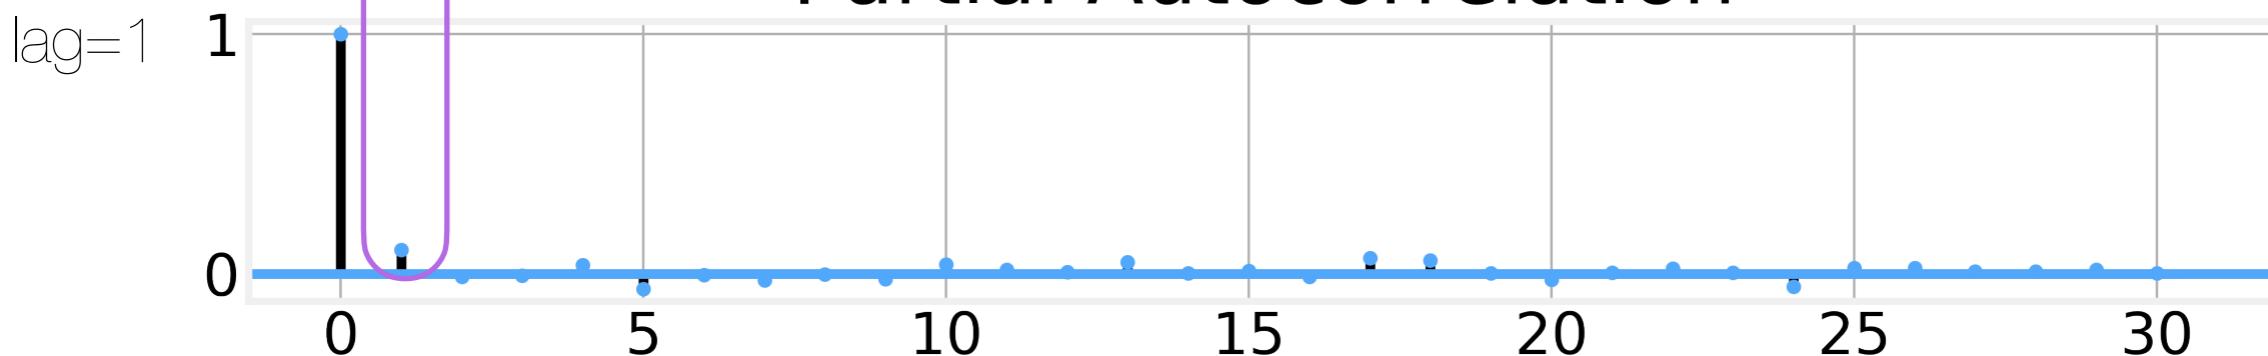
Simulated ARCH(1)<sup>2</sup> Process



Autocorrelation



Partial Autocorrelation



# GARCH Model

- There are many variants of the ARCH model, such as the **ARCH(m)**:

$$Var(x_t | x_{t-1}, \dots, x_{t-m}) = \sigma_t^2 = \alpha_0 + \sum_{l=1}^m \alpha_l x_{t-l}^2$$

- But perhaps the best known variant is the **Generalized Autoregressive Conditionally Heteroskedastic (GARCH)** Model.
- In **GARCH** models, we allow for both a **AR** and a **MA** components.
- The **GARCH** model can be defined as:

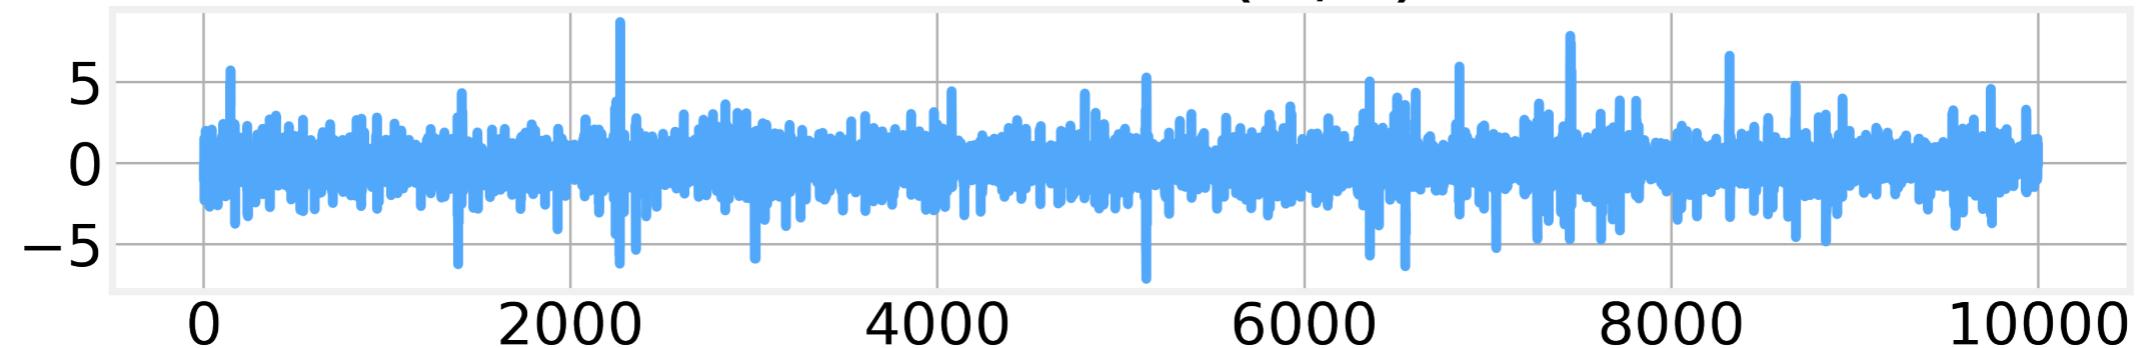
$$x_t = \sigma_t \epsilon_t$$

$$\sigma_t^2 = \alpha_0 + \alpha_1 x_{t-1}^2 + \beta_1 \sigma_{t-1}^2$$

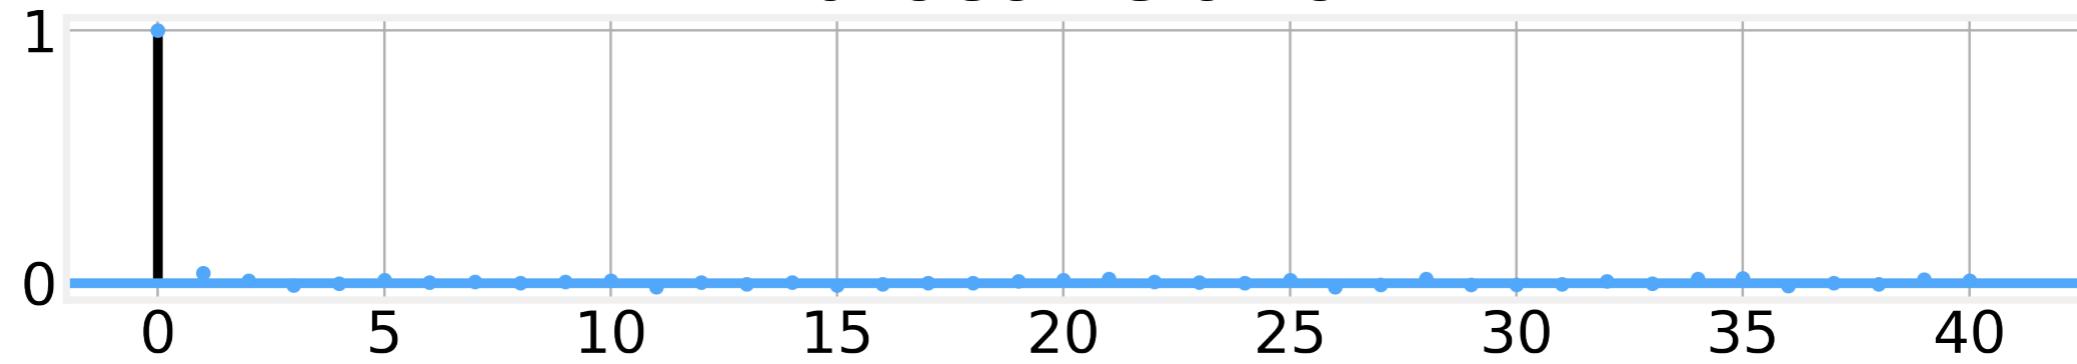
- where the  $\beta$  term represents the **MA** component.
- If we simulate this process, we obtain

# GARCH Model

Simulated GARCH(1,1) Process

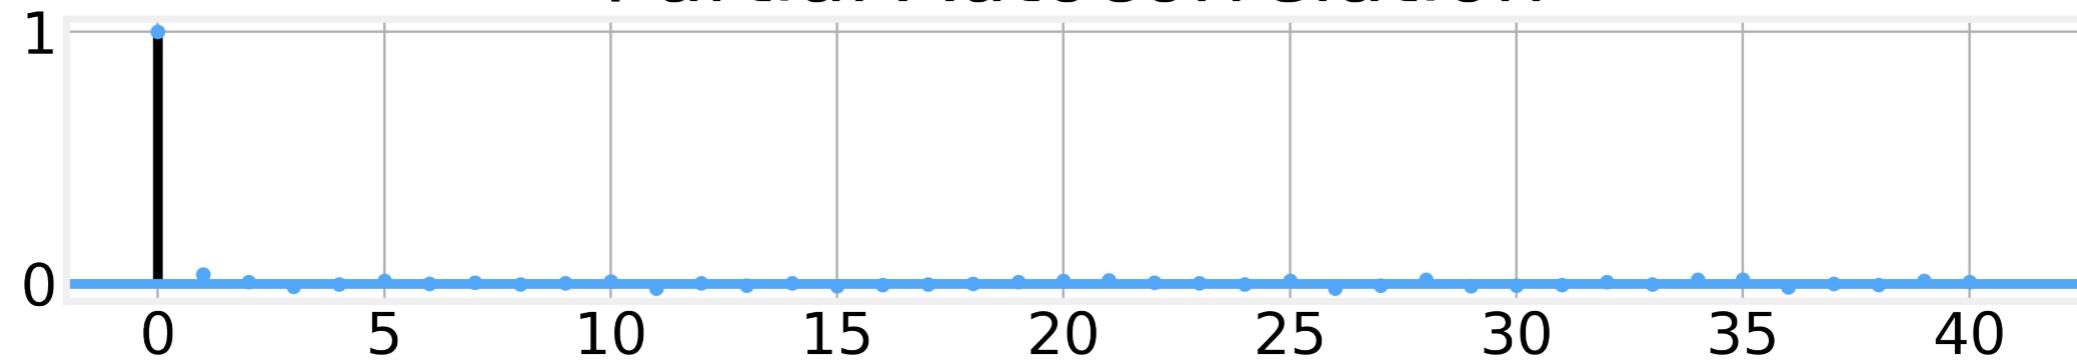


Autocorrelation



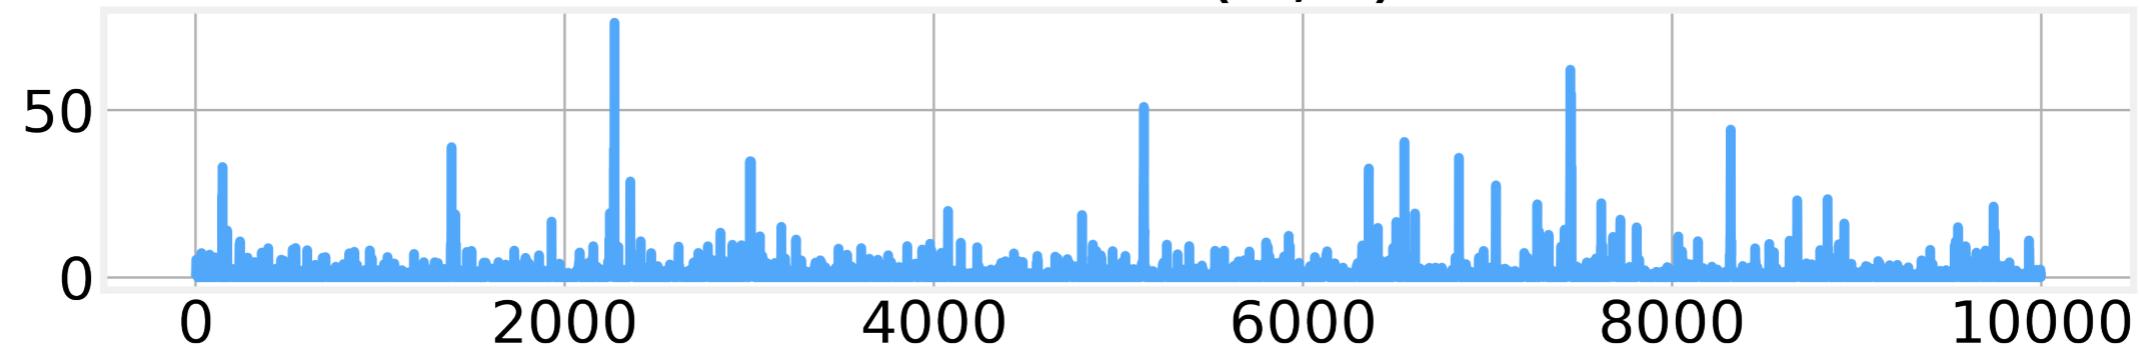
Again just  
White-Noise

Partial Autocorrelation

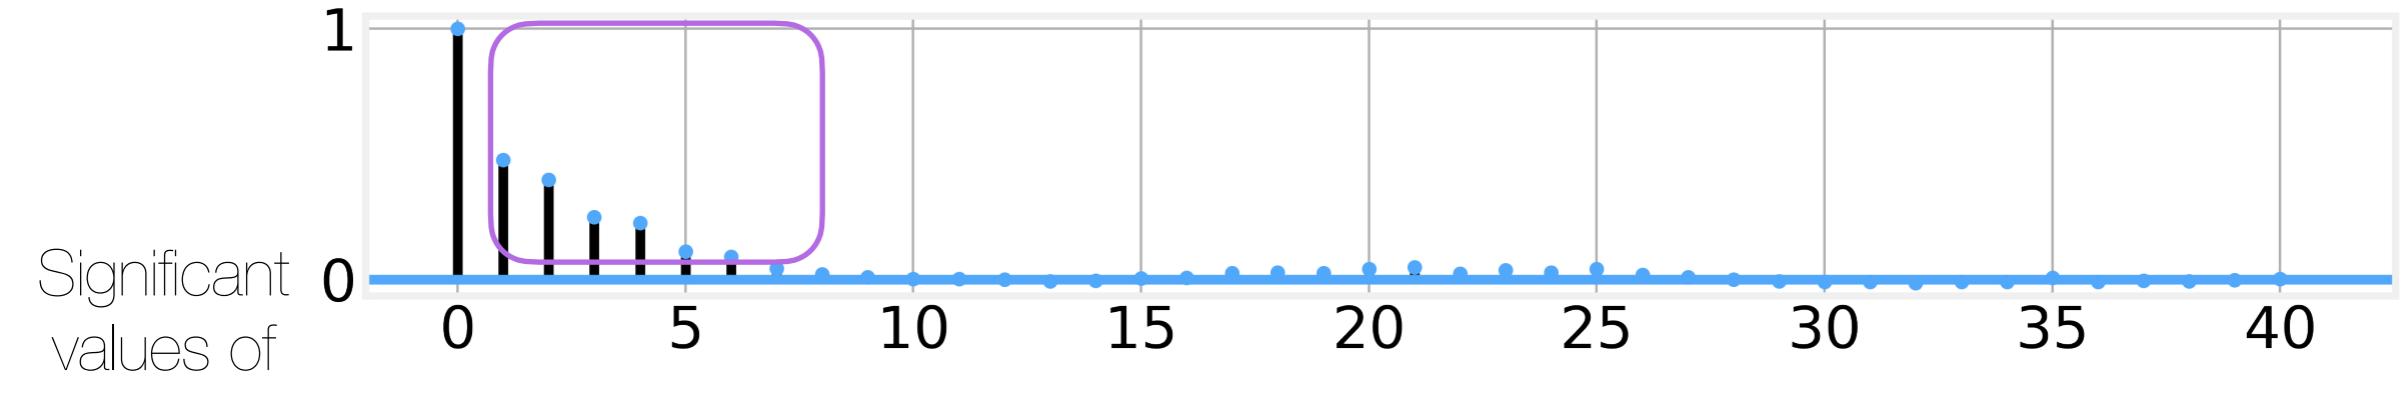


# GARCH Model

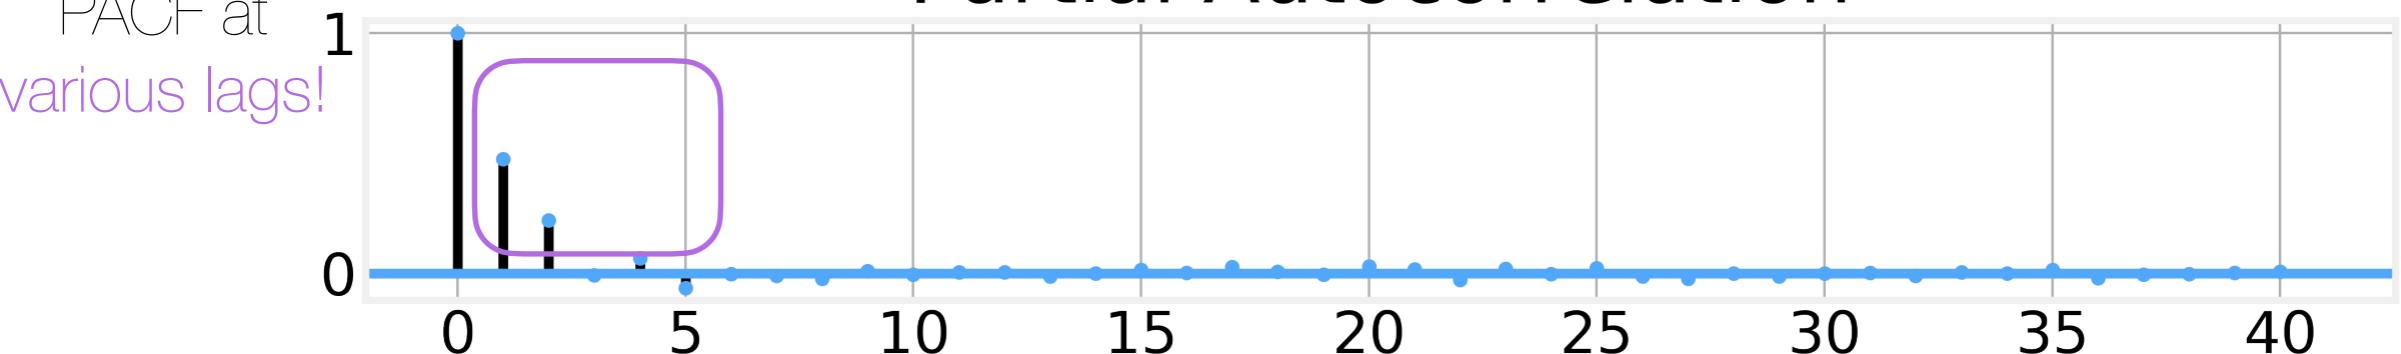
Simulated GARCH(1,1)<sup>2</sup> Process



Autocorrelation

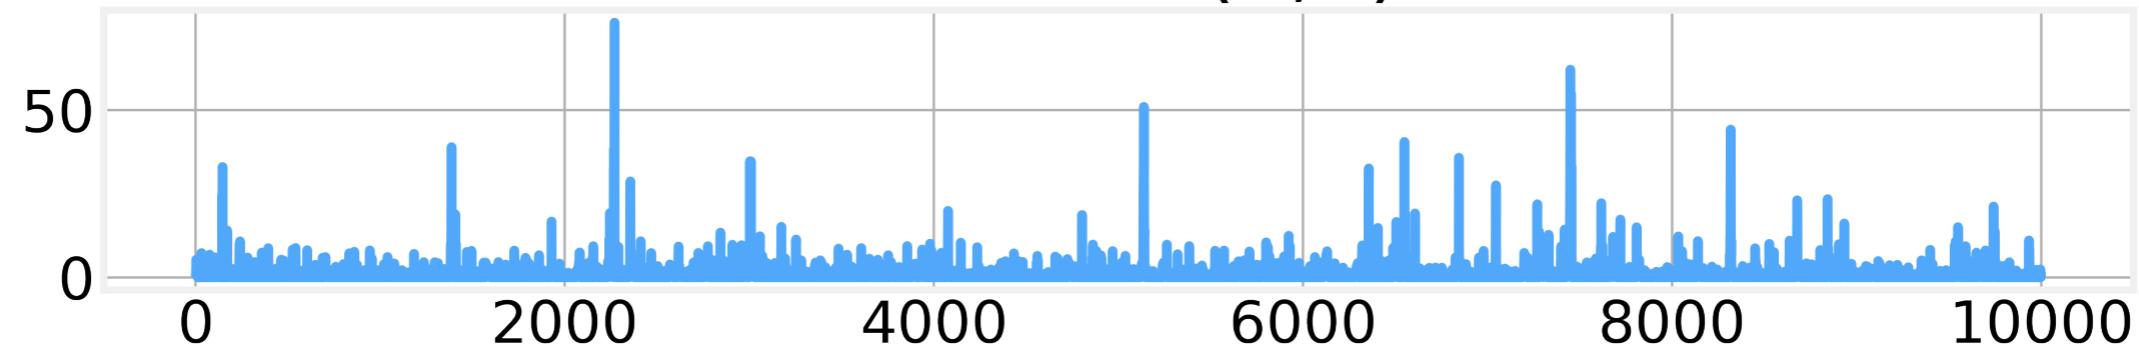


Partial Autocorrelation



# GARCH Model

Simulated GARCH(1,1)<sup>2</sup> Process

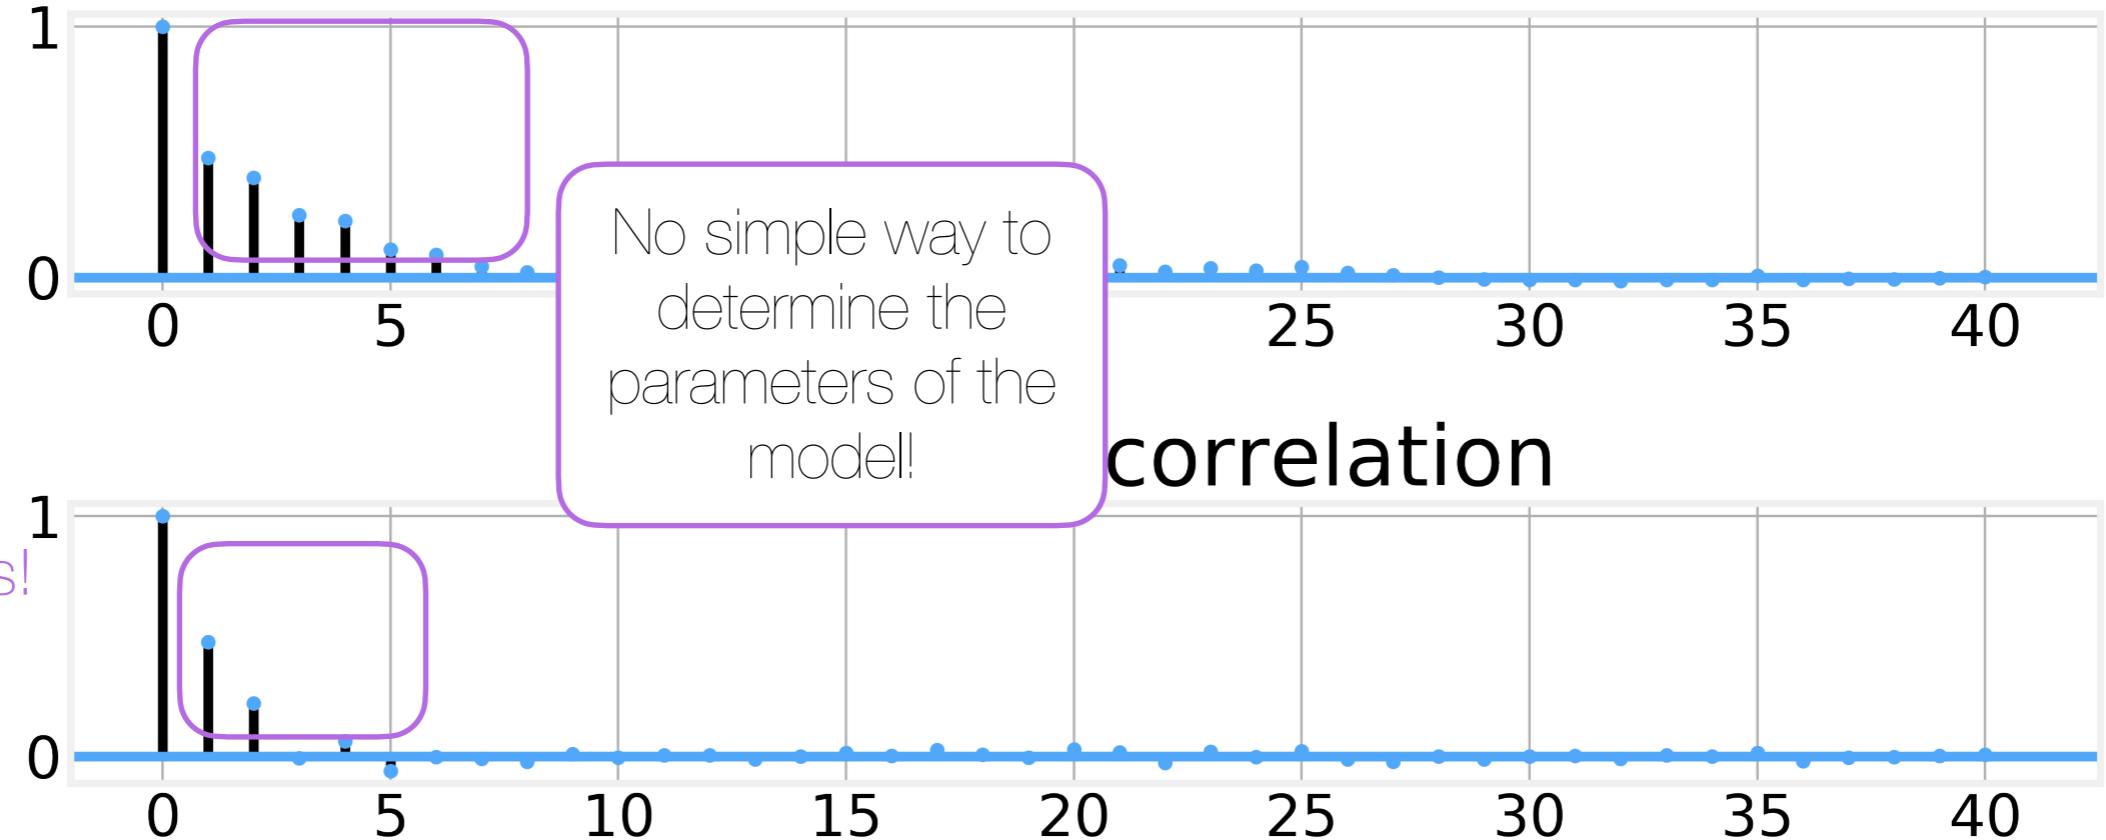


Autocorrelation

Significant values of ACF and PACF at various lags!

No simple way to determine the parameters of the model!

correlation



# Fitting G/ARCH models

Much more involved than fitting ARIMA class models!

Visualize Time Series

Plot ACF/PACF

Plot Squared ACF/PACF

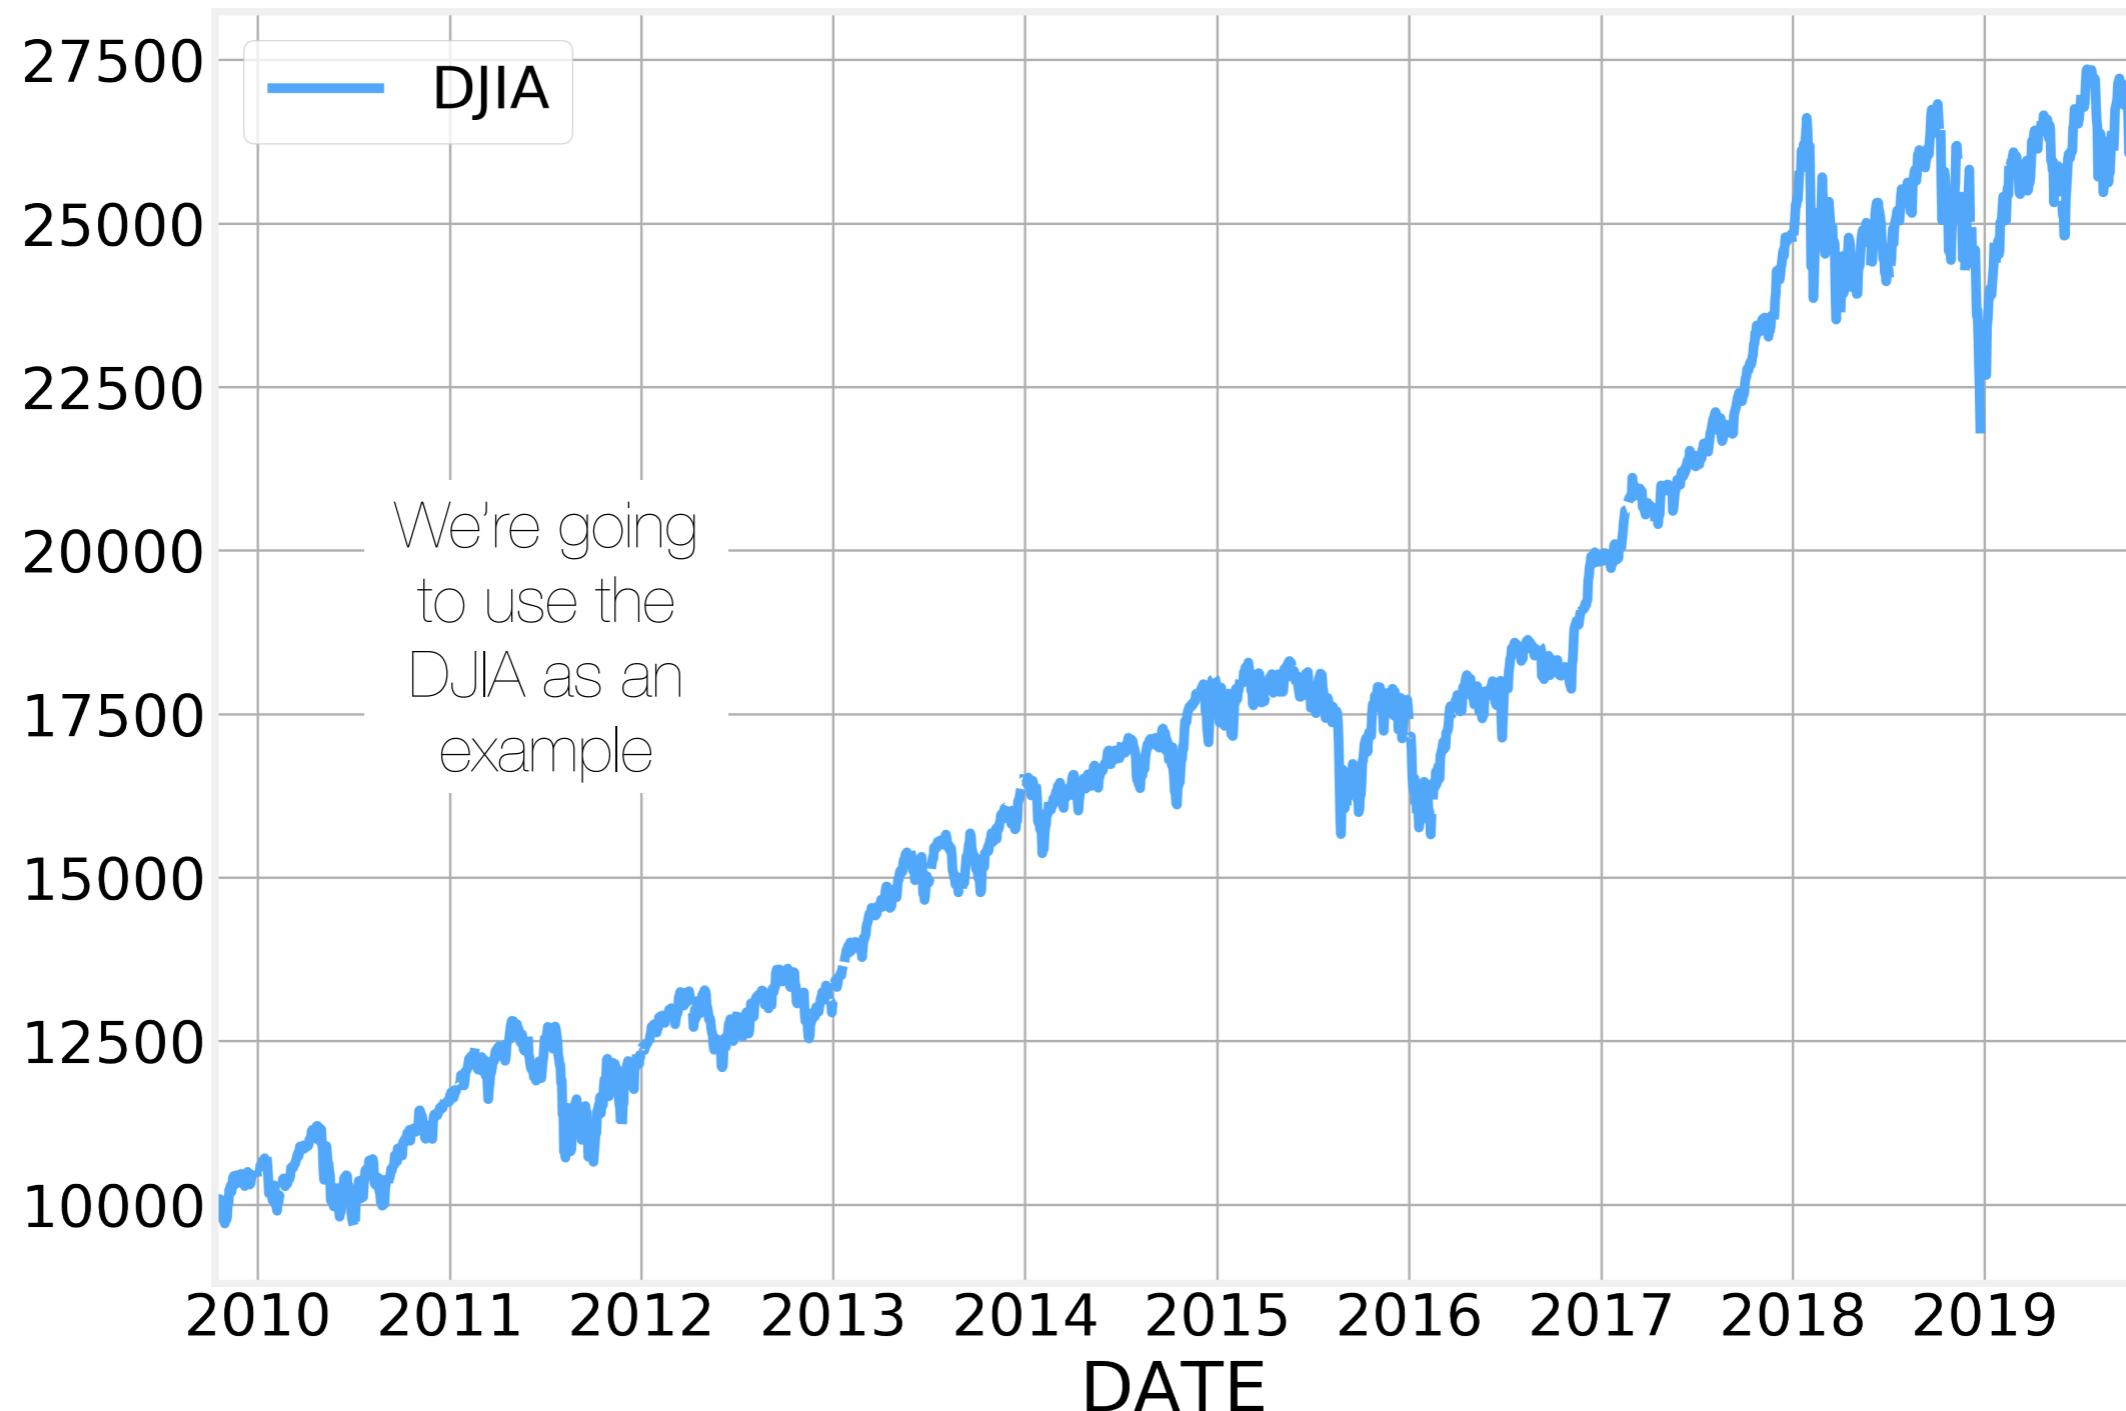
Determine G/ARCH

Build G/ARCH Model

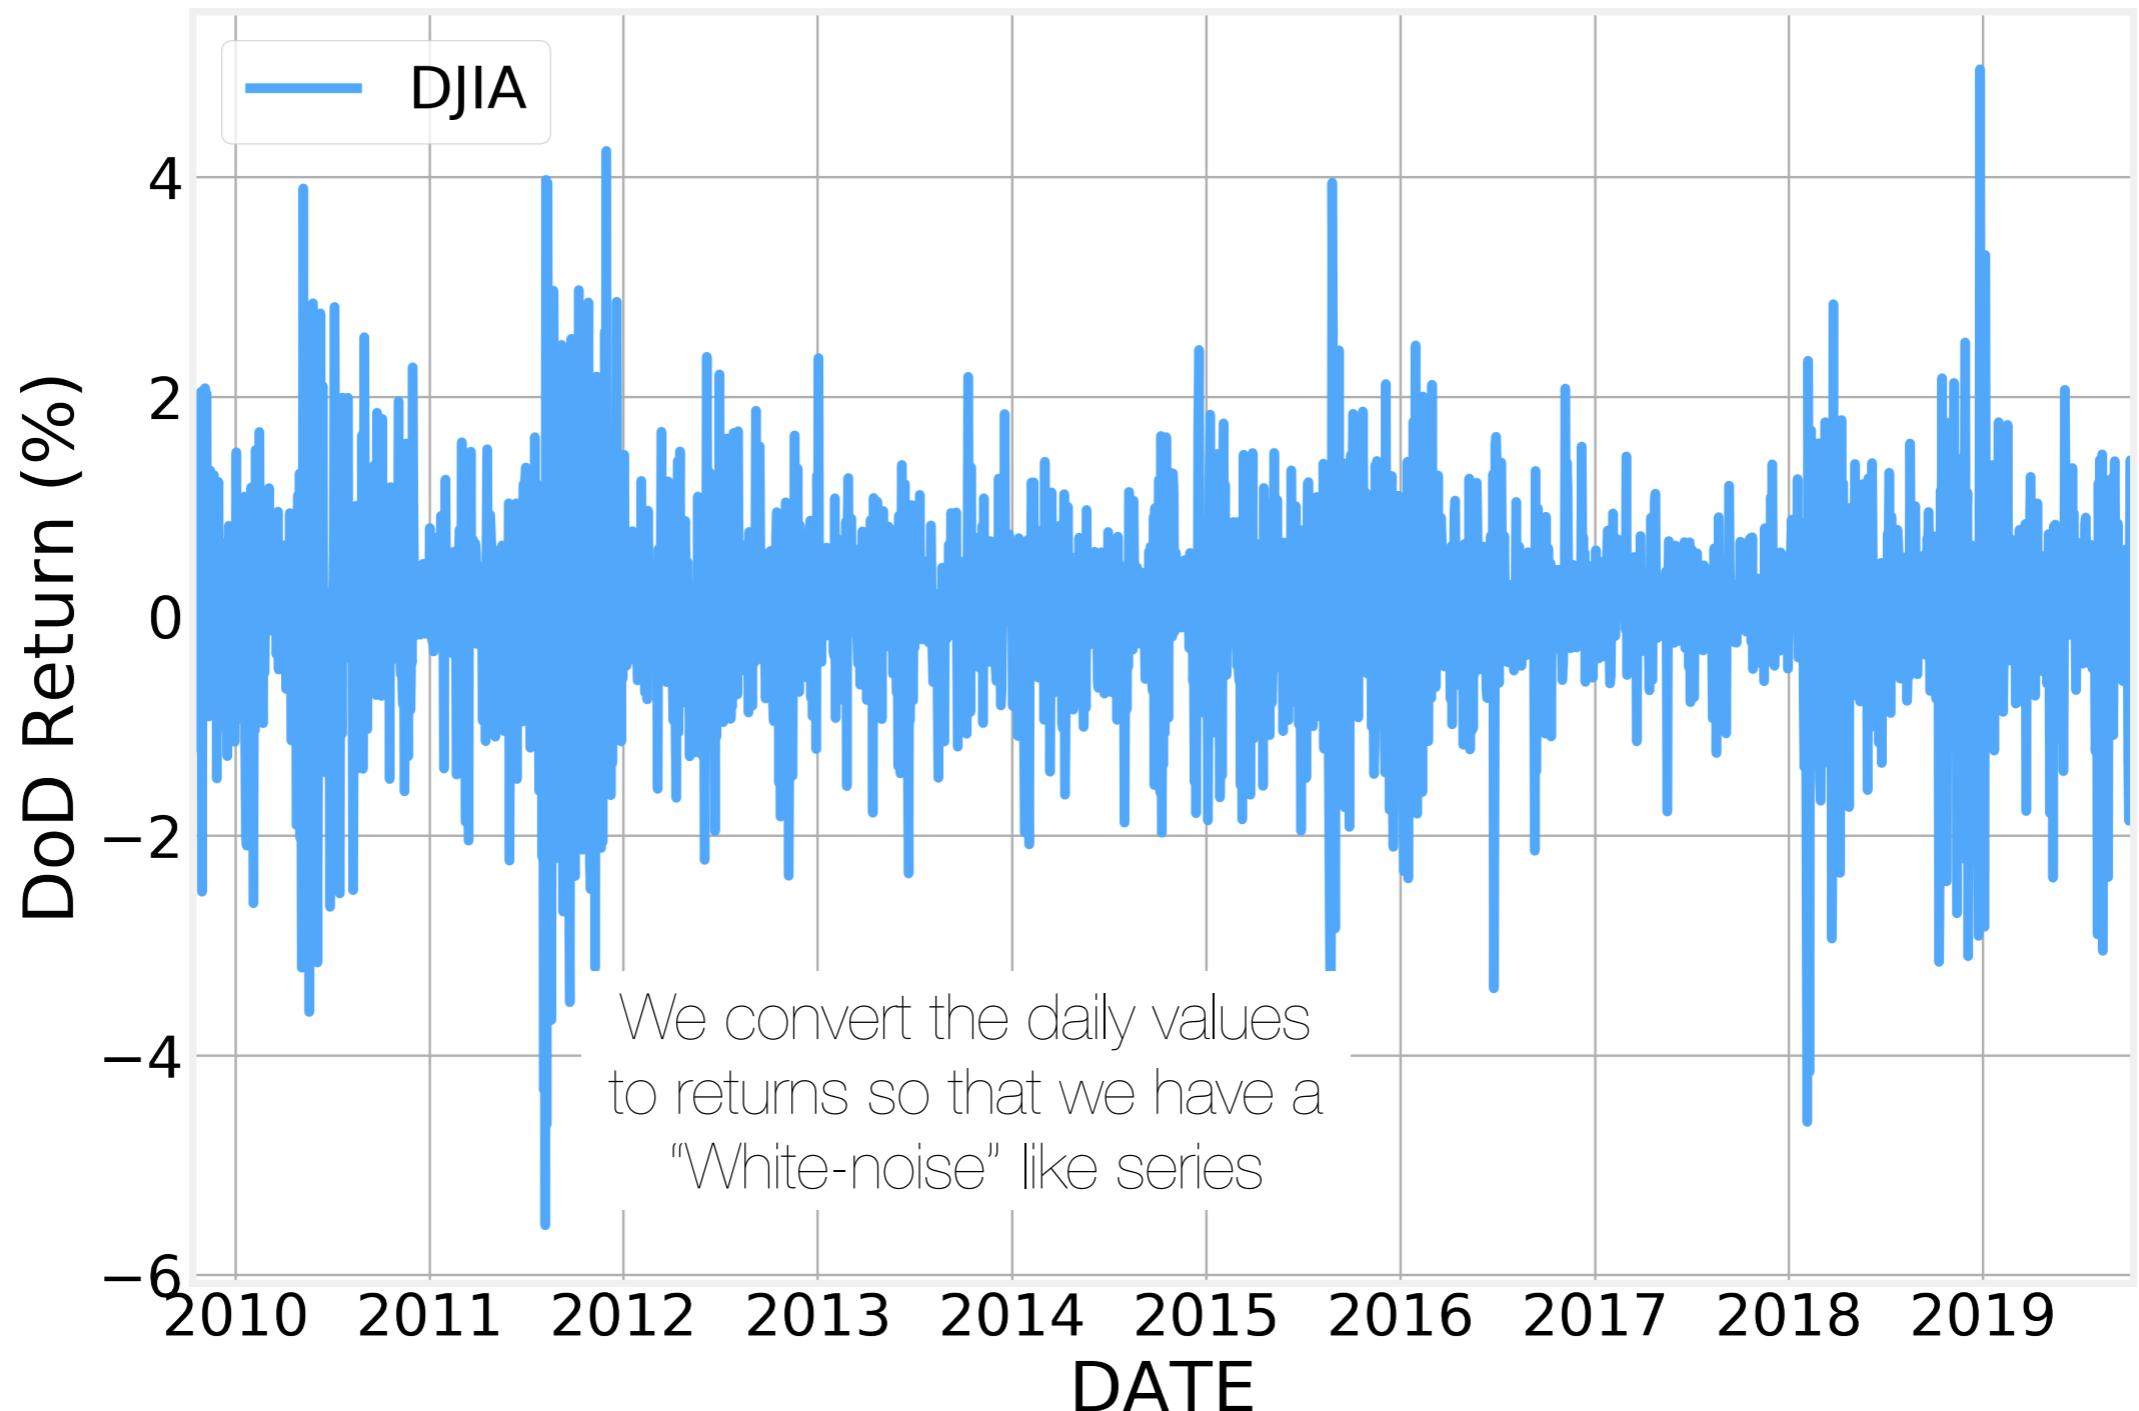
Forecast

Determine the correct order for model by comparing statistics metrics like **AIC!**

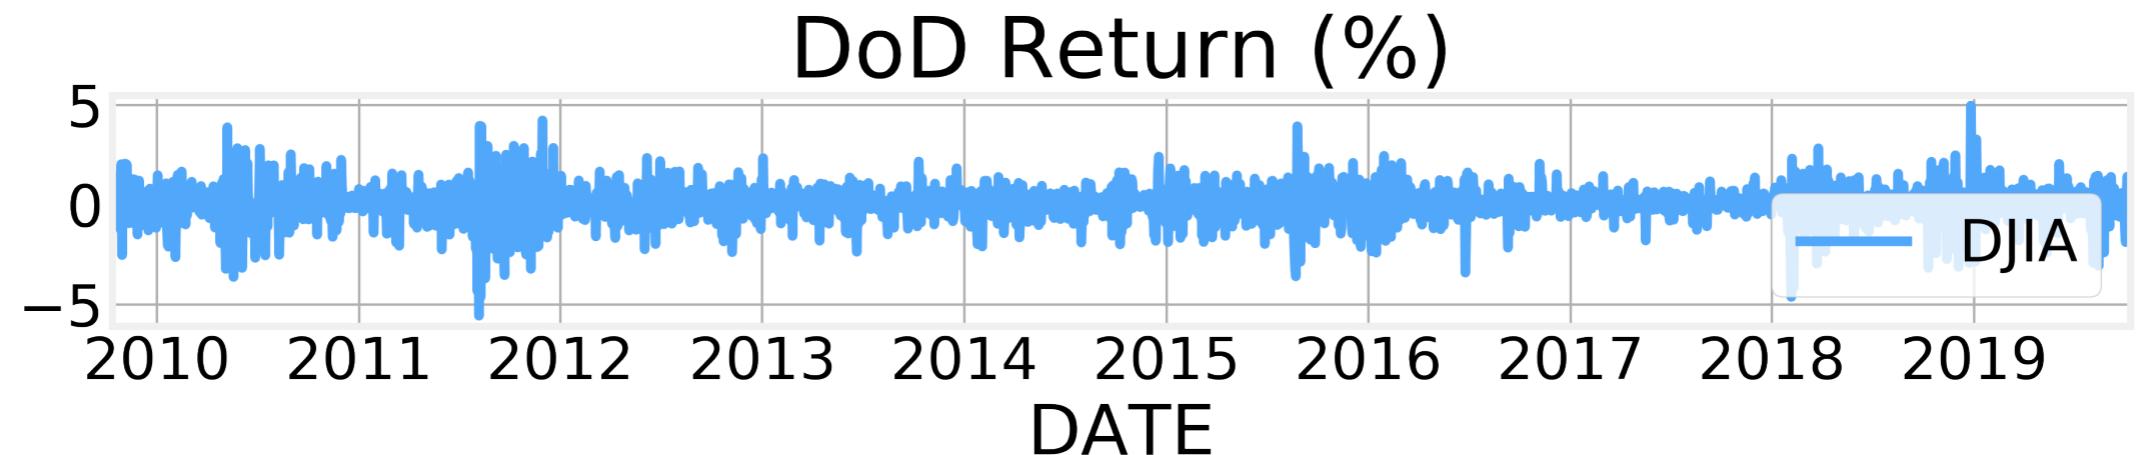
# Visualize Time Series



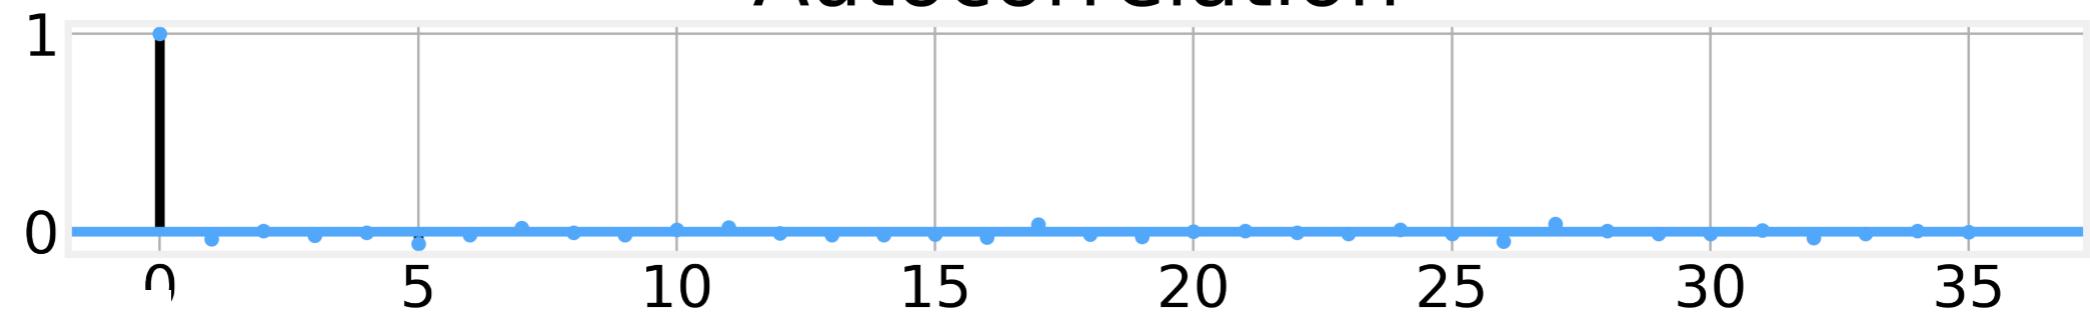
# Visualize Time Series



## Plot ACF/PACF

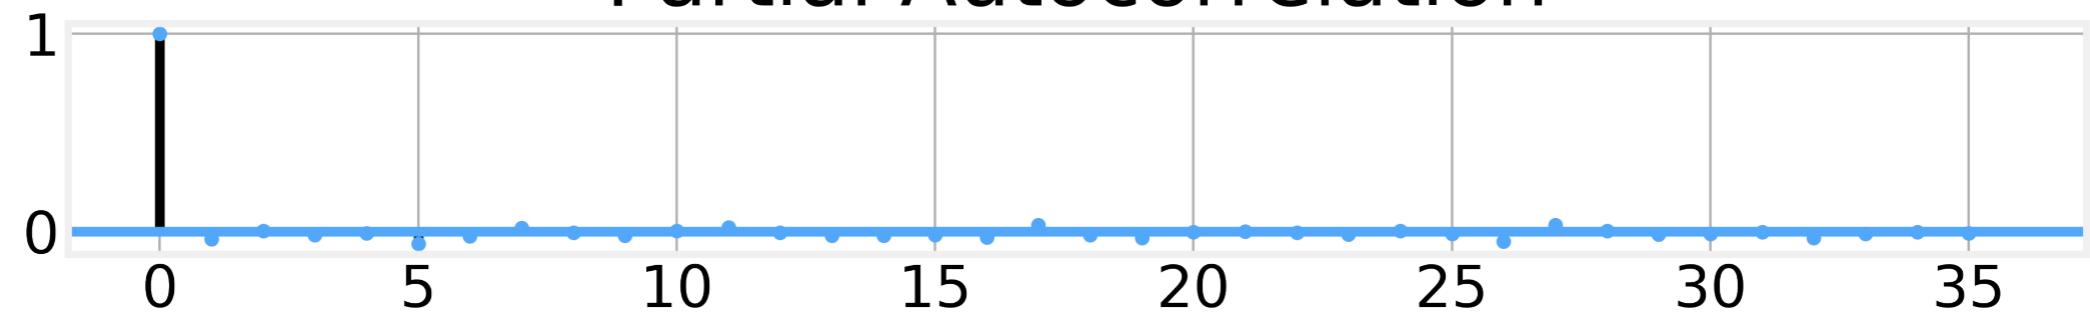


Autocorrelation

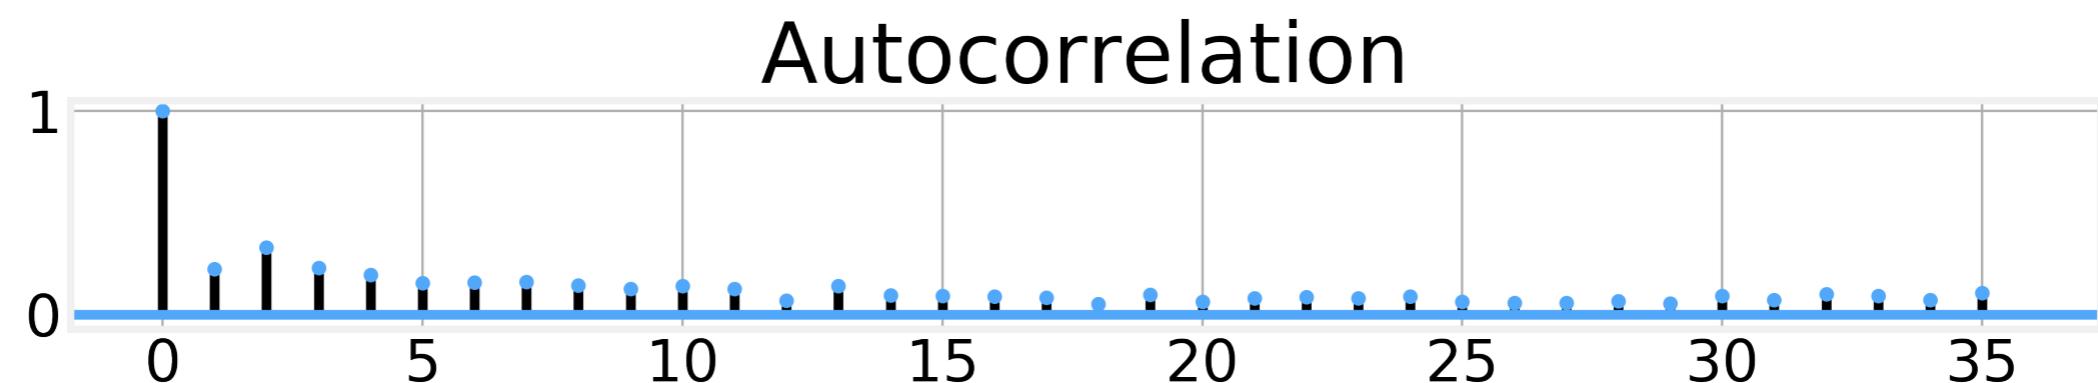
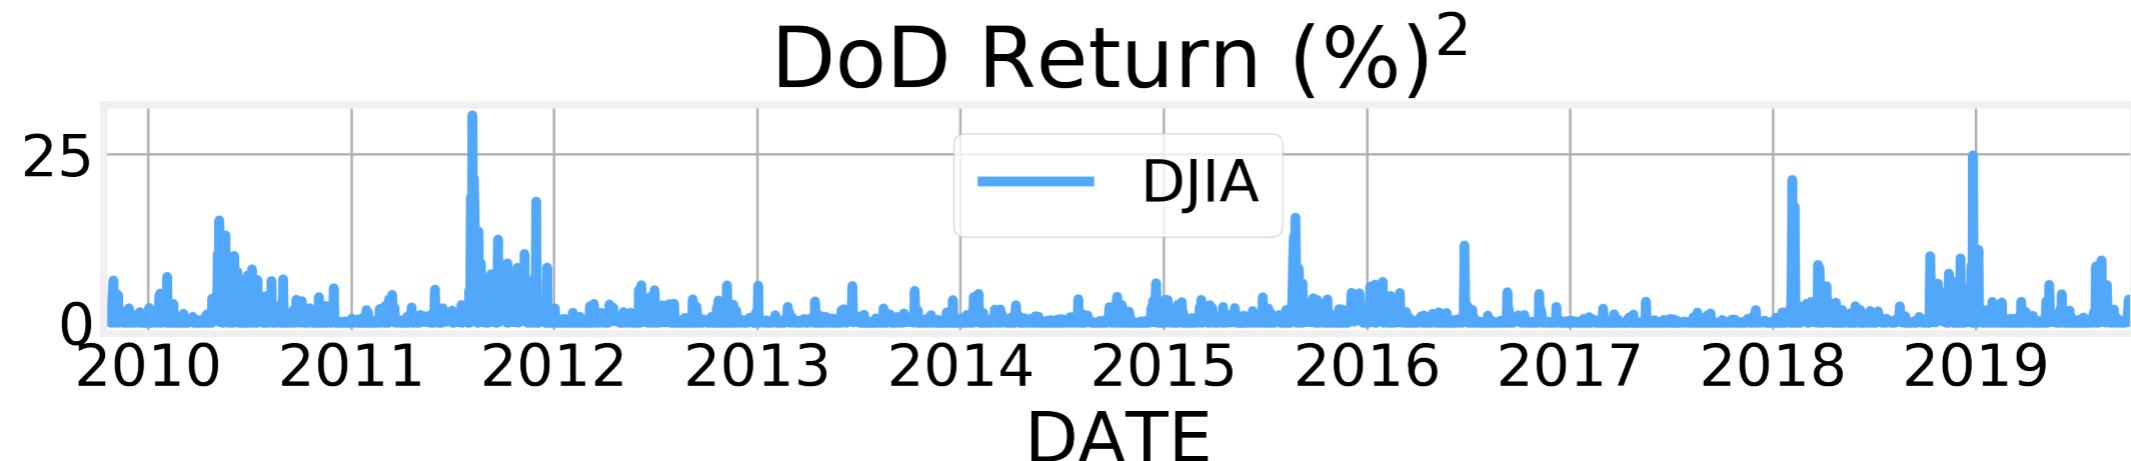


White Noise, as  
expected!

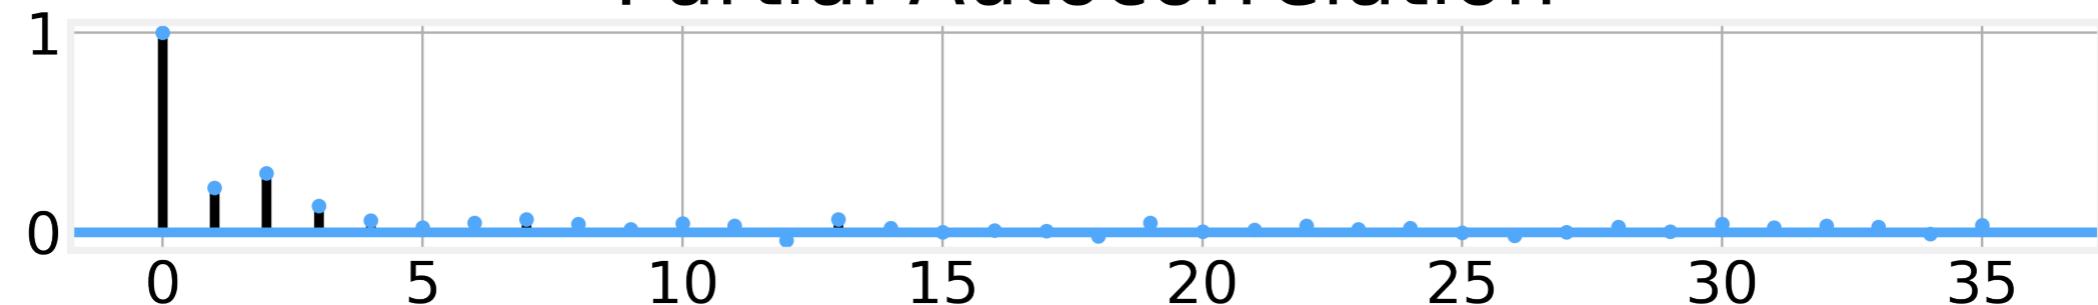
Partial Autocorrelation



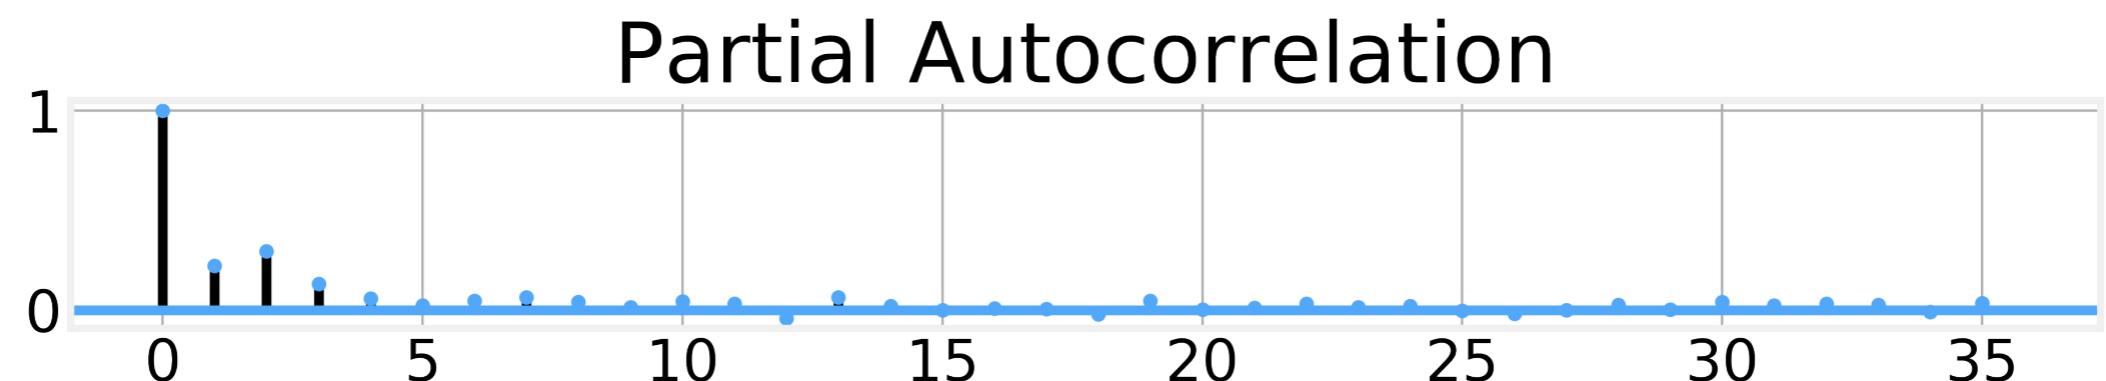
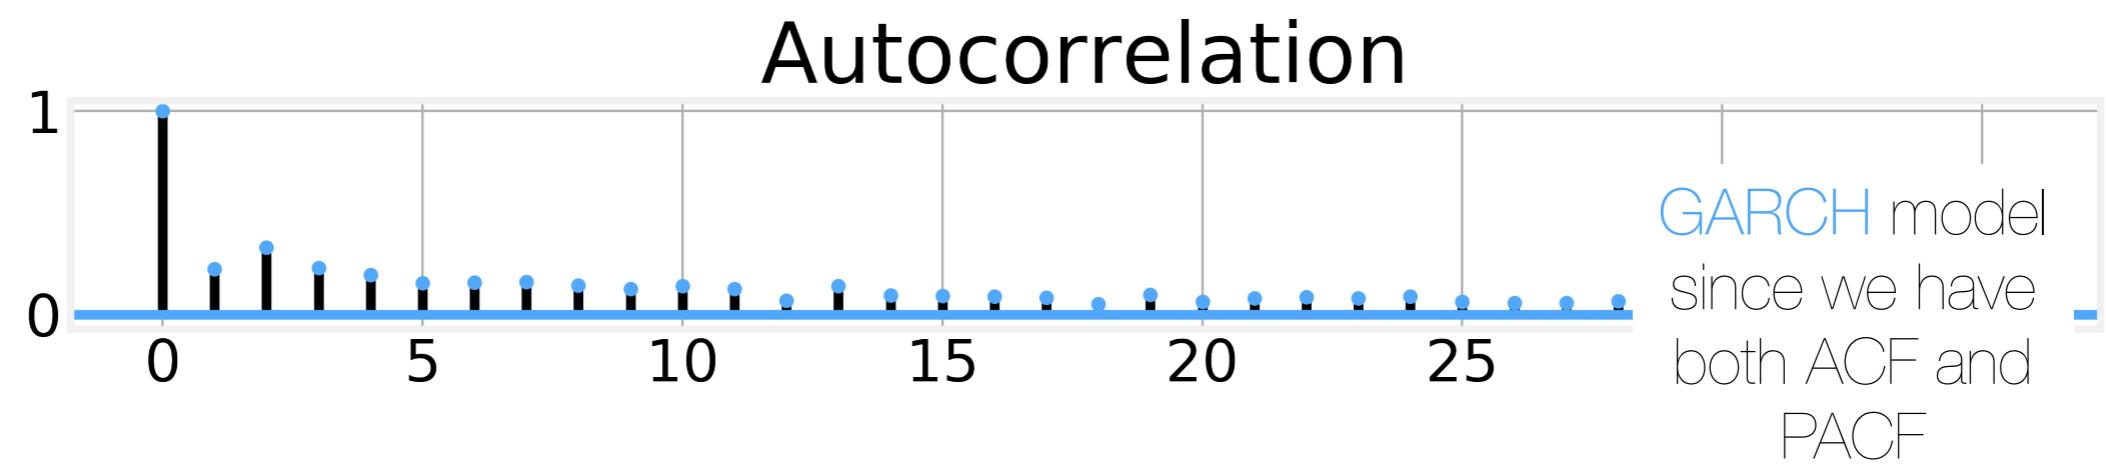
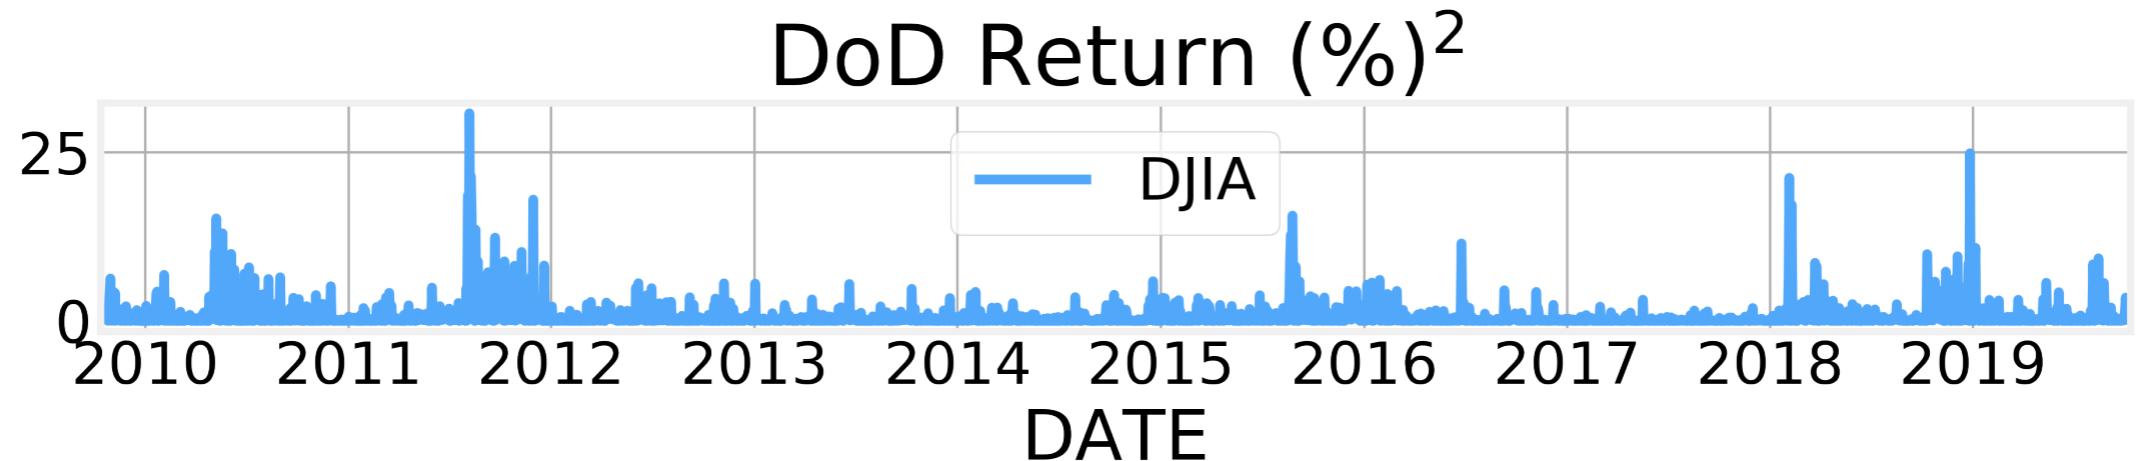
## Plot Squared ACF/PACF



Partial Autocorrelation



## Determine G/ARCH





[arch.readthedocs.io](https://arch.readthedocs.io)

- To fit and model **ARCH** and **GARCH** models, we will use the **arch** library
- **arch** includes a large amount of functionality, but we will focus on the **univariate** sub-package.
- **arch.univariate** provides the **arch\_model()** function that allows us to define a wide variety of **ARCH** or **GARCH** models.
- While **arch\_model()** can take many arguments, there's only a few that are relevant for our purposes:
  - **y** - the (dependent) variable data to train on
  - **vol** - the name of the volatility model, typically **“ARCH”** or **“GARCH”**
  - **p, o, q** - the model parameters



<https://arch.readthedocs.io/en/latest/univariate/generated/arch.univariate.base.ARCHModel.html>

- When a `arch_model()` returns a `ARCHModel` object that represents our model. This object provides a great deal of functionality, but we're particularly interested in:
  - `fit()` - fit the parameters of the model (the values of  $\alpha_i$ ,  $\beta_j$ , etc.)
  - `simulate()` - generate new values from the fitted model.
  - `forecast()` - forecast new values of the data.

And continuing with our example...

# Build G/ARCH Model

- We use **arch** to generate a **GARCH(1,1)** model and train it on the DJIA returns dataset.
- **arch** provides us with a nice summary of the model parameters, similarly to **statsmodels**

Constant Mean - GARCH Model Results			
<b>Dep. Variable:</b>	DJIA	<b>R-squared:</b>	-0.001
<b>Mean Model:</b>	Constant Mean	<b>Adj. R-squared:</b>	-0.001
<b>Vol Model:</b>	GARCH	<b>Log-Likelihood:</b>	-3000.92
<b>Distribution:</b>	Normal	<b>AIC:</b>	6009.85
<b>Method:</b>	Maximum Likelihood	<b>BIC:</b>	6033.31
		<b>No. Observations:</b>	2608
<b>Date:</b>	Mon, Aug 10 2020	<b>Df Residuals:</b>	2604
<b>Time:</b>	23:46:42	<b>Df Model:</b>	4

Can use AIC/BIC to compare various models and pick the best (smallest) one

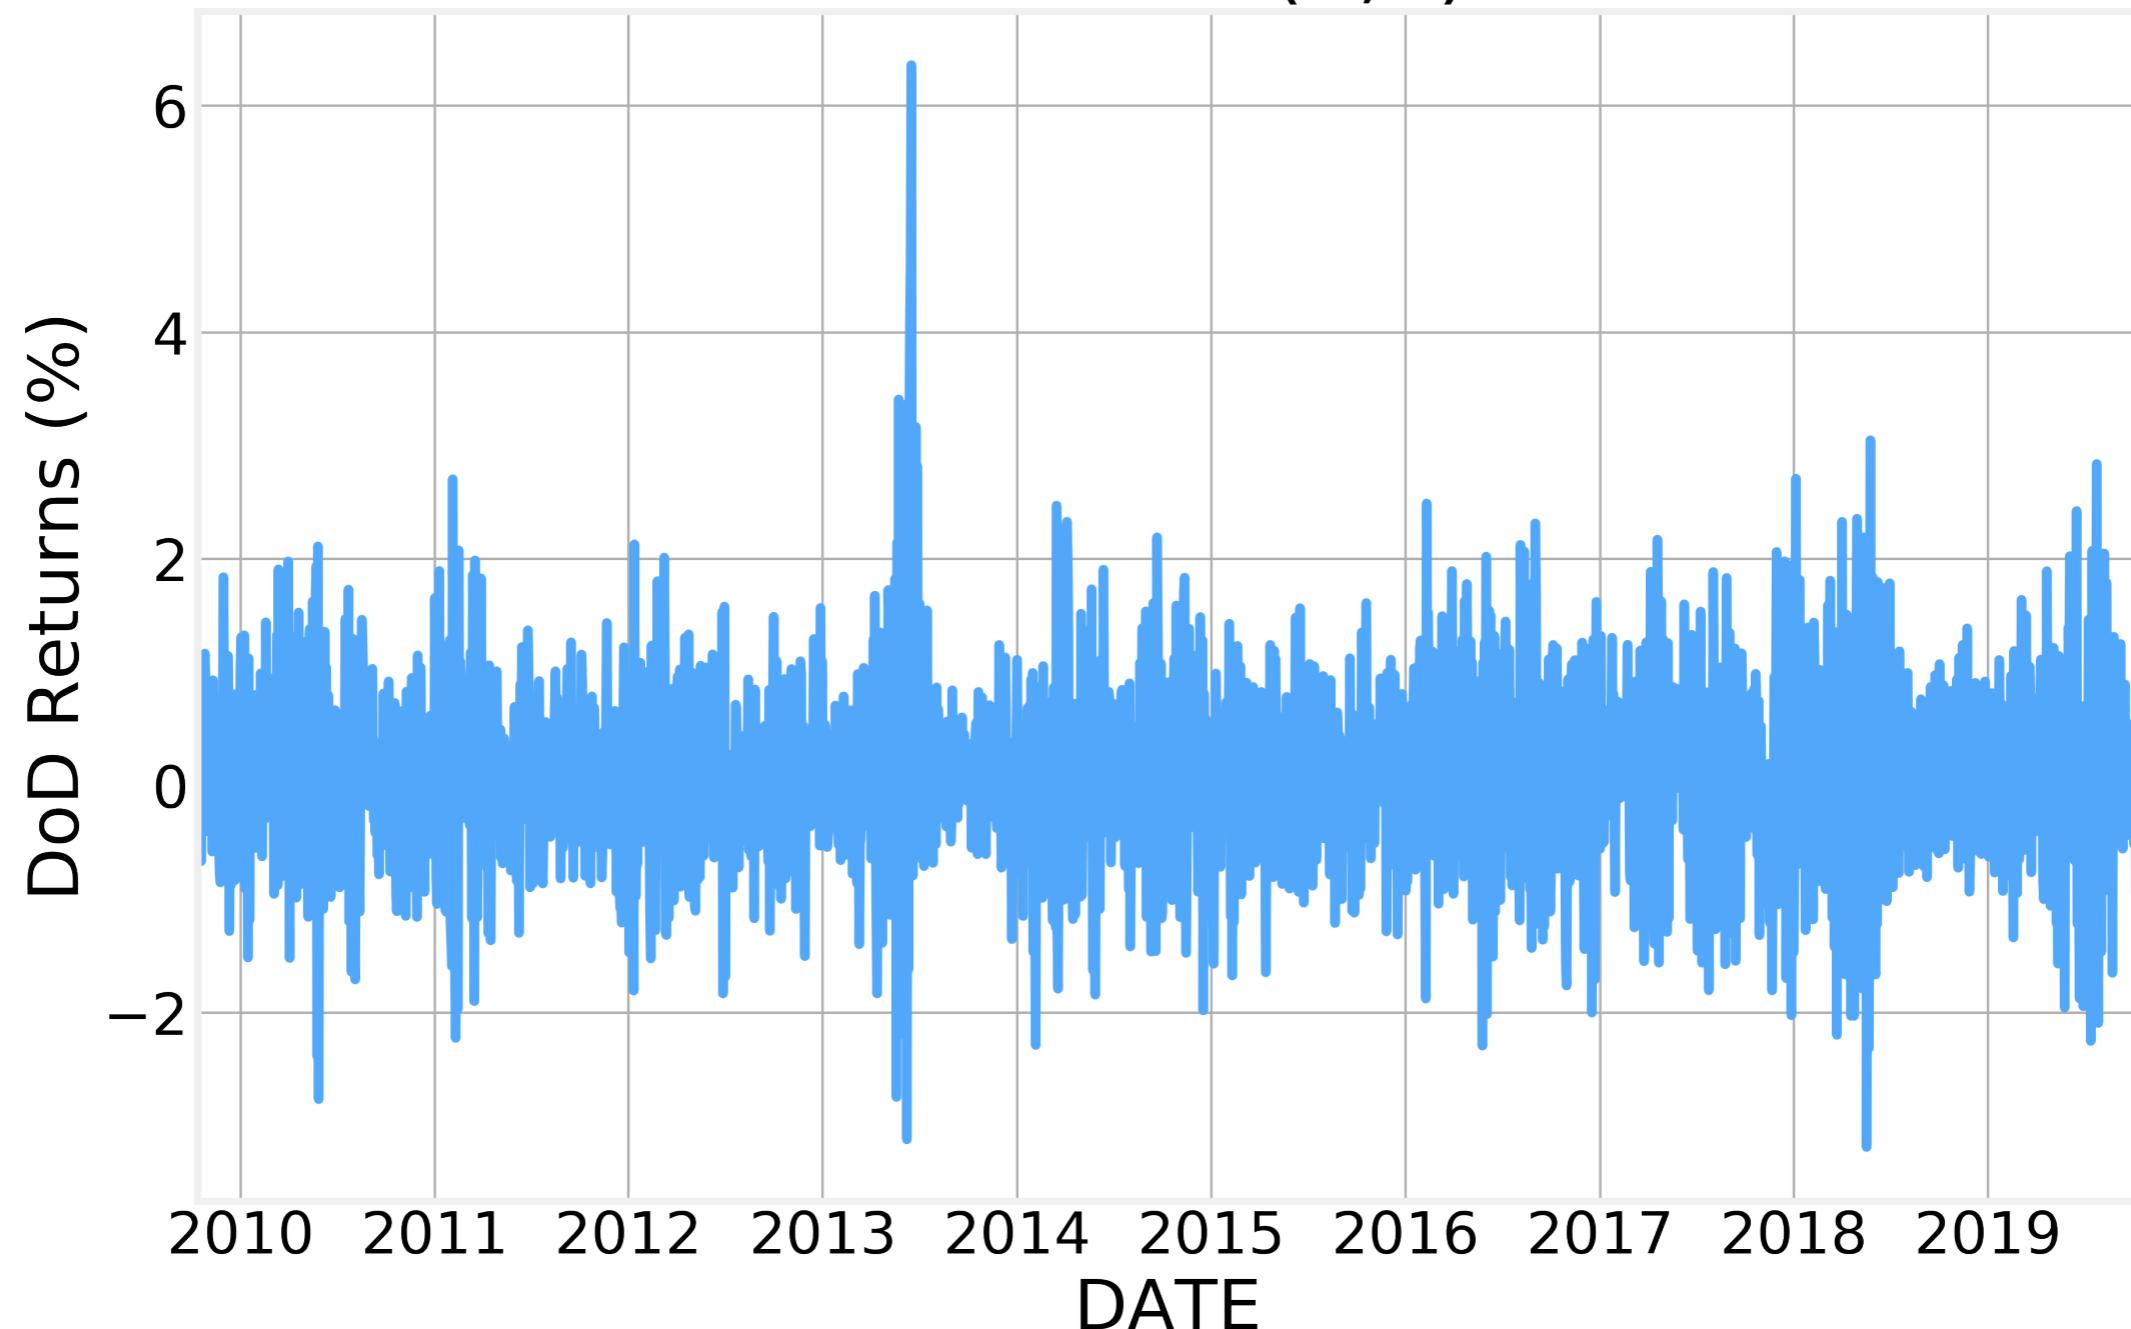
Mean Model					
	<b>coef</b>	<b>std err</b>	<b>t</b>	<b>P&gt; t </b>	<b>95.0% Conf. Int.</b>
<b>mu</b>	0.0744	1.327e-02	5.608	2.043e-08	[4.840e-02, 0.100]

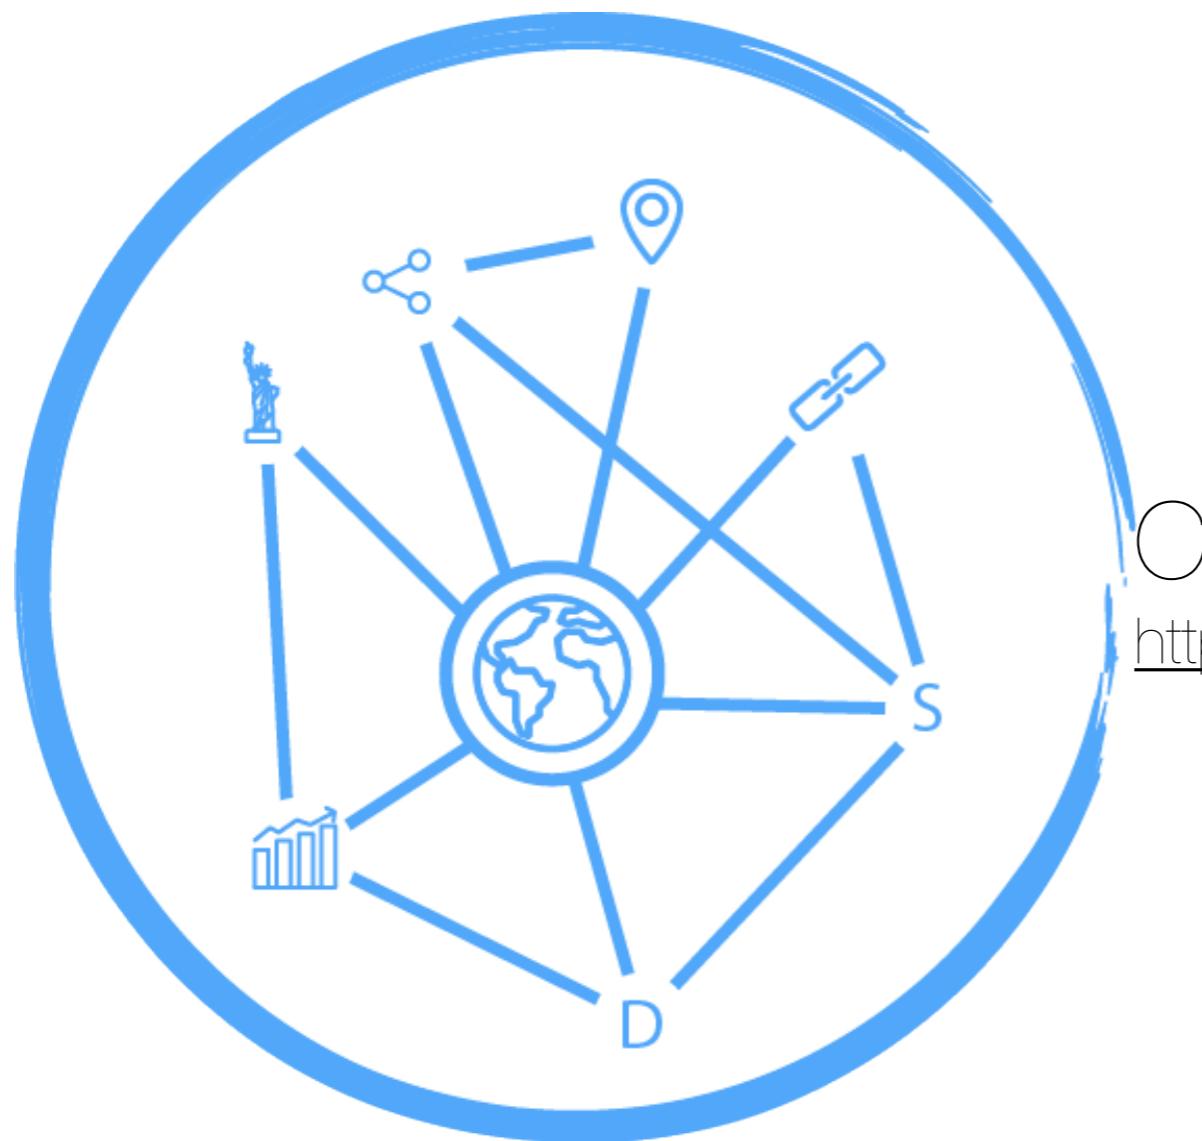
Volatility Model					
	<b>coef</b>	<b>std err</b>	<b>t</b>	<b>P&gt; t </b>	<b>95.0% Conf. Int.</b>
<b>omega</b>	0.0318	6.646e-03	4.780	1.749e-06	[1.875e-02, 4.480e-02]
<b>alpha[1]</b>	0.1587	2.316e-02	6.851	7.310e-12	[0.113, 0.204]
<b>beta[1]</b>	0.8033	2.280e-02	35.232	6.527e-272	[0.759, 0.848]

parameter coefficients

## Forecast

Simulated GARCH(1,1) Process

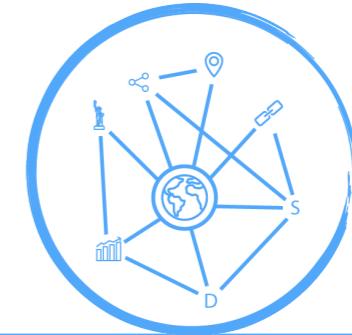




Code - ARCH Models

<https://github.com/DataForScience/AdvancedTimeseries>

# Events



[www.data4sci.com/newsletter](http://www.data4sci.com/newsletter)



<http://paypal.me/data4sci>



**Natural Language Processing (NLP) from Scratch**

<http://bit.ly/LiveLessonNLP> - On Demand